

3. Instruction Set Architecture – The LC2k and ARM architectures

EECS 370 – Introduction to Computer Organization - Winter 2016

Profs. Valeria Bertacco & Reetu Das


EECS Department
University of Michigan in Ann Arbor, USA

© Bertacco-Das, 2016

The material in this presentation cannot be
copied in any form without our written permission

**Get started on project 1
Assembler due 9/21!**

**Grab coffee and donuts
in the meantime**



PROCRASTINATION

HARD WORK OFTEN PAYS OFF AFTER TIME,
BUT LAZINESS ALWAYS PAYS OFF NOW.

Announcements

❑ Project 1

- Project 1.a due 6:00pm, Thursday, 1/21
- Project 1.m & 1.s due 6:00pm, Thursday, 2/4

❑ Homework 1 assigned

- Due Thursday, 1/21
- Submit electronically to www.gradescope.com, any electronic or scanned format is fine as long as the graders can find/read your answers, see homework sheet for submission details

❑ Discussions only on Friday 1/15, no classes on Monday 1/18

- Please attend any Friday discussion, or review discussion materials on the class website and/or in office hours

Recap

- ❑ ISAs, instruction encoding
- ❑ Storage: registers, status register, memory, (accumulator)
- ❑ Addressing modes: direct, indirect, reg.indirect, base+displacement, pc-relative

Instruction Set Architecture (ISA) Design Lectures

- ❑ Lecture 2: Storage types and addressing modes
- ❑ **Lecture 3 : LC-2K and ARM architecture**
- ❑ Lecture 4 : Converting C to assembly – basic blocks
- ❑ Lecture 5 : Converting C to assembly – functions
- ❑ Lecture 6 : Translation software; libraries, VMs
- ❑ Lecture 7 : Memory layout

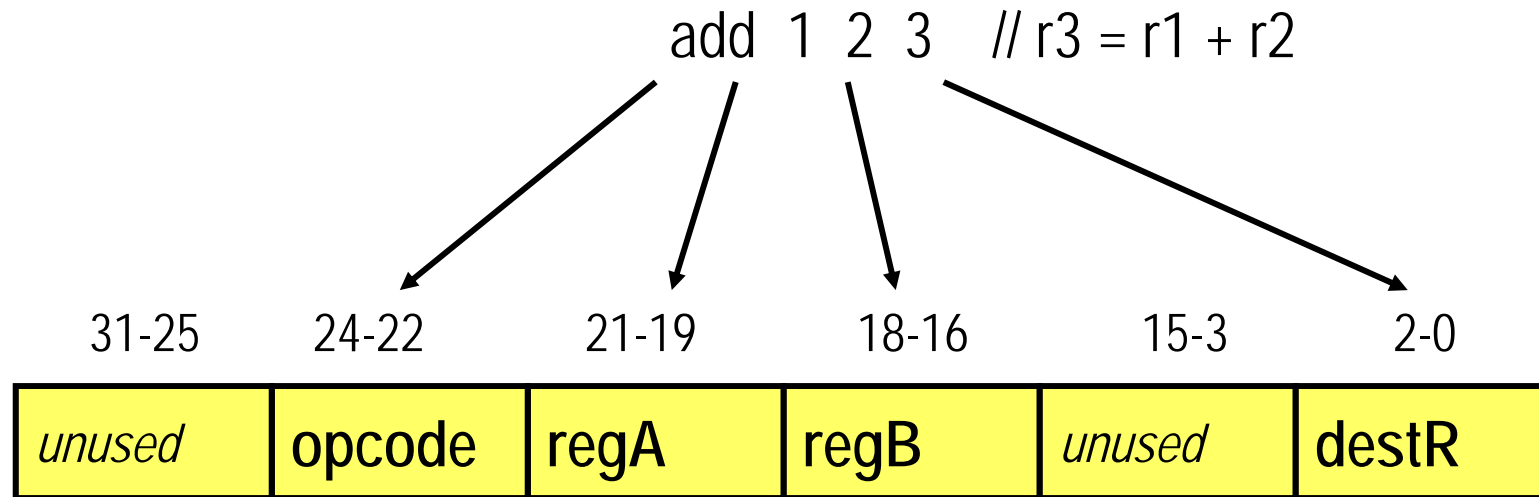
LC2K Processor

- ❑ 32-bit processor
 - Instructions are 32 bits
 - Integer registers are 32 bits
- ❑ 8 registers
- ❑ supports 65536 words of memory (addressable space)

- ❑ 8 instructions
 - add, nor, lw, sw, beq, jalr, halt, noop

Instruction Encoding


- Instruction set architecture defines the mapping of assembly instructions to machine code




Instruction Formats

- Positional organization of bits (**Implies nothing about bit values!!!**)

- R type instructions (add, nor)



- I type instructions (lw, sw, beq)



Bit Encodings

❑ Opcode encodings

- add (000), nor (001), lw (010), sw (011), beq (100), jalr (101), halt (110), noop (111)

❑ Register values

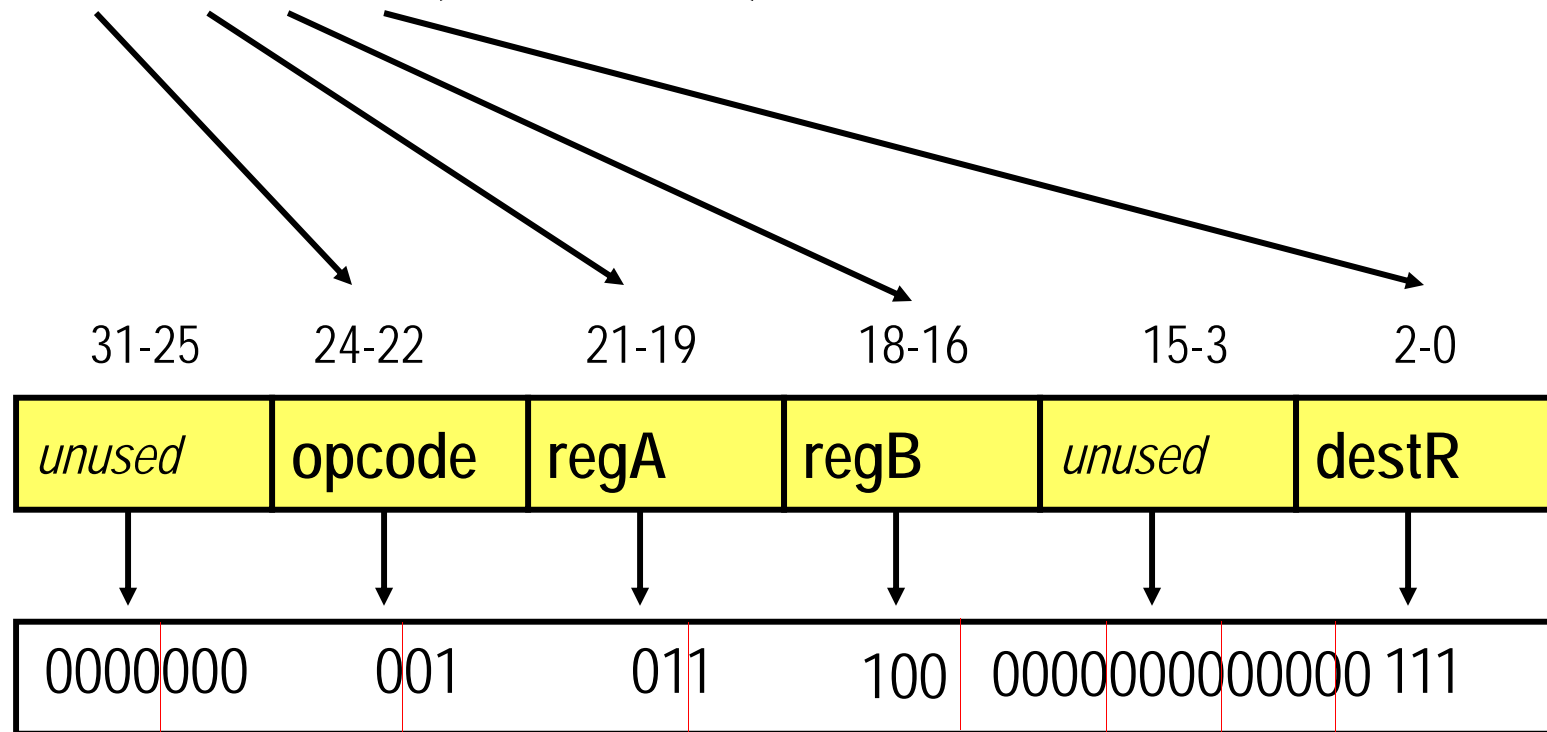
- Just encode the register number (r2 = 010)

❑ Immediate values

- Just encode the values (Remember to give all the available bits a value!!)

Example Encoding - nor

❑ `nor 3 4 7 (r7 = r3 nor r4)`

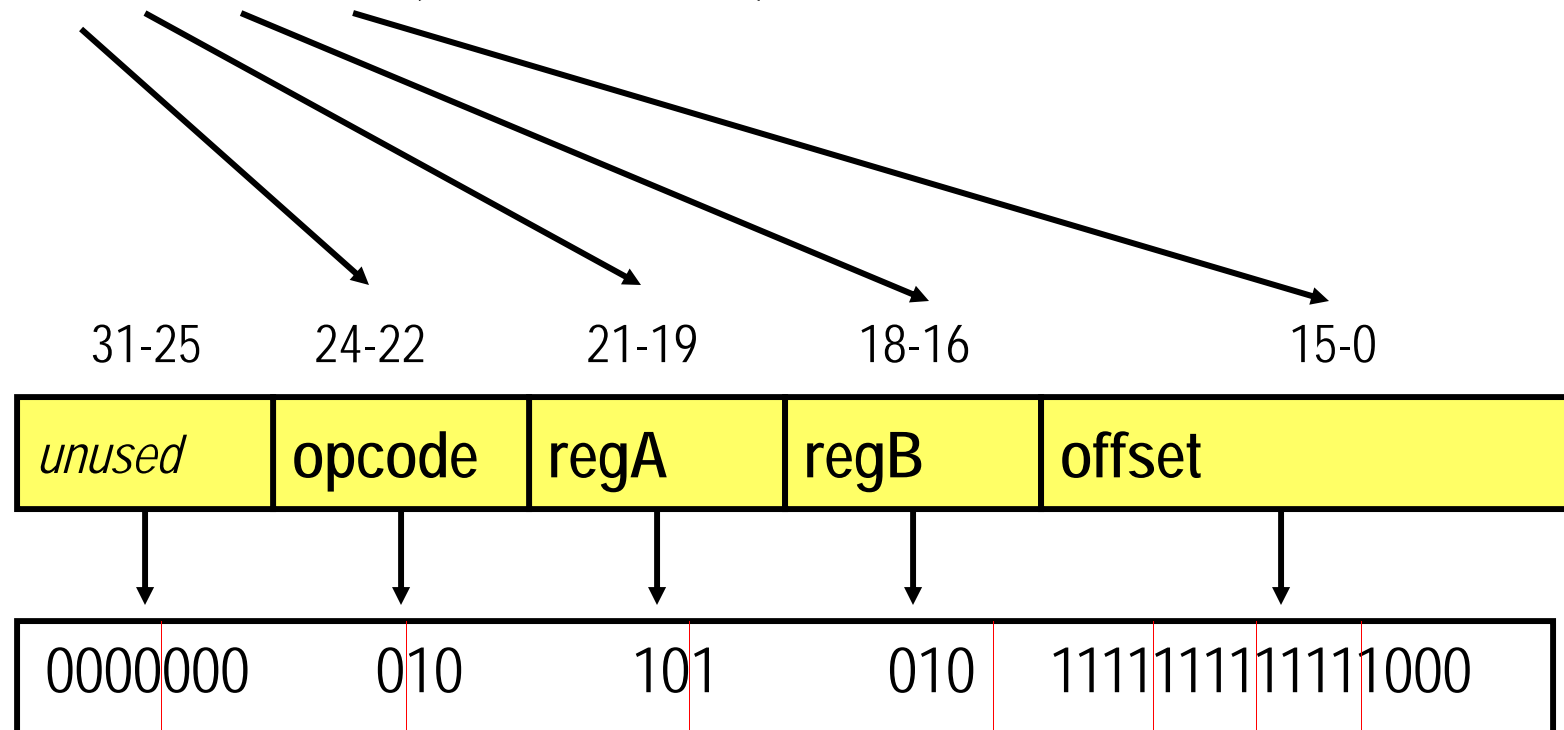


Convert to Hex → 0x005C0007

Convert to Dec → 6029319

Example Encoding - lw

❑ lw 5 2 -8 (r2 = M[r5 + -8])




Convert to Hex → 0x00AAFF8

Convert to Dec → 11206648

Class Problem 1

- Compute the encoding in Hex for:
- add 3 7 3 ($r3 = r3 + r7$) (add = 000)
 - sw 1 5 67 ($M[r1+67] = r5$) (sw = 011)



ARM Instruction Set

□ Three main types of instructions:

1. Arithmetic

- Add, subtract, multiply, multiply-accumulate
- Logical: and, or, xor, shift, rotate, etc.
- Compare : eq, ne, lt, gt, le, mi, pl, etc.
- Data movement: mov, mvn

2. Memory access

- ldr, ldm (load multiple), str, stm (store multiple)

3. Sequencing / control flow


- b (branch), bl (branch and link)

ARM Arithmetic Instructions

- Format: three operand fields
 - Dest. register usually the first one – check instruction format
 - e.g., `add r3, r4, r7`
 - Both SUB and RSB available

r3→g
r4→h
r5→i
r6→j
r7→f

- C-code example: `f = (g + h) - (i + j)`



ARM Arithmetic Instructions – special 2nd operand

Second source operand can be

1. A **register** (see previous)
2. An **immediate** (8 bit) – optionally rotated by an immediate (4bits)
 - e.g., RSB r3, r4, #10
 - bit 25 of the instruction indicates whether reg or imm

C-code example: $f = 10 - g$

```
rsb    r7, r3, #10
```

3. A **register to which a shift/rotate is applied first**
 - the shift amount can be specified in a register, or a 5-bit immediate
- e.g., MVN r5, r3 LSL #2 $r5 = \text{NOT}(r3 \ll 2)$

Opcode and mnemonic do not change

Mini-review: What is a *shift* ?

❑ C/C++

- `a = b >> 2;`
- `c = d << 4;`

b = 01101110

00110111

❑ ARM

- `mov r3, r2, LSR #2`
- `mov r5, r4, LSL #4`

❑ Rotate also available in ARM

- `mov r3, r2, ROR #3`

r2 = 01101110

11001101

What About immediates Larger Than 8 Bits??

- ❑ Use rotate right 0, 2, ..30 option on immediate operands
 - e.g., `add r3, r4, #FF ROR 8` (equals `#FF000000`)
 - If this is not possible, load constant from memory into a register, then use a register operand
- ❑ Pseudo-instruction `LDR rd, =K`
- ❑ C-code example: `f = f + 0x104`

```
add    r7, r7, 0x41 ROR 30
```

(Note: ROR 30 = ROL 2)

0x041 (0b000001000001)




0x104 (0b000100000100)

ARM Arithmetic Instructions - recap

- `add` `dest, src1, src2imm` `// add`
- `adc` `dest, src1, src2imm` `// add with carry`
- `mul` `dest, src1, src2` `// multiply`
- `sub` `dest, src1, src2imm` `// sub src2imm from src1`
- `sbc` `dest, src1, src2imm` `// sub src2imm from src1 with carry`
- `rsb` `dest, src1, src2imm` `// sub src1 from src2imm`
- `rsc` `dest, src1, src2imm` `// sub src1 from src2imm with carry`
- `and` `dest, src1, src2imm` `// logical AND of bits`
- `orr` `dest, src1, src2imm` `// logical OR of bits`
- `eor` `dest, src1, src2imm` `// exclusive OR of bits`
- `mov` `dest, src1imm` `// mov src1imm into dest (can be shifted)`
- `mvn` `dest, src1imm` `// mov NOT(src1imm) to dest (shift ok)`
- `cmp` `src1, src2imm` `// compare src1 to src2imm & set PSR`

Class Problem 2

- Show the C and ARM assembly for extracting the value in bits 15:10 from a 32-bit integer variable



Want these bits

Assume the variable is in r1

Remember each hex digit is 4 bits

ARM Memory Instructions

- ❑ Supports base + displacement and base + register modes
 - Base is a register
 - Displacement is a 12-bit immediate, positive or negative
 - Register displacement can be shifted/rotated (5-bits), immediate cannot
 - Pre-indexed and post-indexed offset modes allow to modify the base register before or after the memory access
- ❑ Examples:
 - `ldr r3, [r4, #1000] // load word: r3= M[r4+1000]`
 - `ldr r3, [r4, -r1, LSL #2] // load word: r3= M[r4-(r1<<2)]`
 - `str r3, [r4, # -34]! // load word pre-ndx: M[r4-34]= r3; r4-=34`
 - `ldr r3, [r4], # 20 // load word post-ndx: r3= M[r4]; r4+=20`

Load Instruction Sizes

How much data is retrieved from memory at the given address?

- ❑ `ldr r3, [r4, #1000]`
 - retrieve a word (32 bits) from address (r4+1000)

- ❑ `ldrh r3, [r4, #1000]`
 - retrieve a halfword (16 bits) from address (r4+1000)

- ❑ `ldrb r3, [r4, #1000]`
 - retrieve a byte (8 bits) from address (r4+1000)

Sign/Zero Extension

- ❑ Registers in ARM are 32 bits!
- ❑ So what happens when you load 8 or 16 bits?

- Sign extend if the load is signed



- Zero extend if the load is unsigned (default)



ARM Memory Instructions - recap

❑ Load instructions

- `ldrsb` $\backslash\backslash$ load byte signed (load 8 bits, sign extend)
- `ldrb` $\backslash\backslash$ load byte unsigned (load 8 bits, zero extend)
- `ldrsh` $\backslash\backslash$ load halfword (load 16 bits, sign extend)
- `ldrh` $\backslash\backslash$ load halfword unsigned (load 16 bits, zero extend)
- `ldr` $\backslash\backslash$ load word (load 32 bits, no extension)

❑ Store instructions (No sign/zero extension for stores)


- `strb r3, [r4, #1000]` $\backslash\backslash$ store 8 LSBs of r3 to M[r4+1000]
- `strh r3, [r4, #1000]` $\backslash\backslash$ store 16 LSBs of r3 to M[r4+1000]
- `str r3, [r4, #1000]` $\backslash\backslash$ store all 32 bits of r3 to M[r4+1000]

❑ Addressing: `base+displ`; `displ` = {reg, imm, reg w. shift/rot}

❑ Pre- and post-indexing


Load Instruction in Action

- ❑ `ldrsb r3, [r4, #1000]` // load signed byte
Retrieves 8-bit value from memory location ($r4+1000$) and puts the result into r3 (sign extended)



Load Instruction in Action – other example

- ❑ `ldrsb r3, [r4, #1000]` // load signed byte
Retrieves 8-bit value from memory location ($r4+1000$) and puts the result into r3 (sign extended)



Example Code Sequence


What is the final state of memory once you execute the following instruction sequence? (assume r0 stores the value 0)

```
ldr    r4, [r0, #100]
ldrsb  r3, [r0, #102]
str     r3, [r0, #100]
strb    r4, [r0, #102]
```

register file

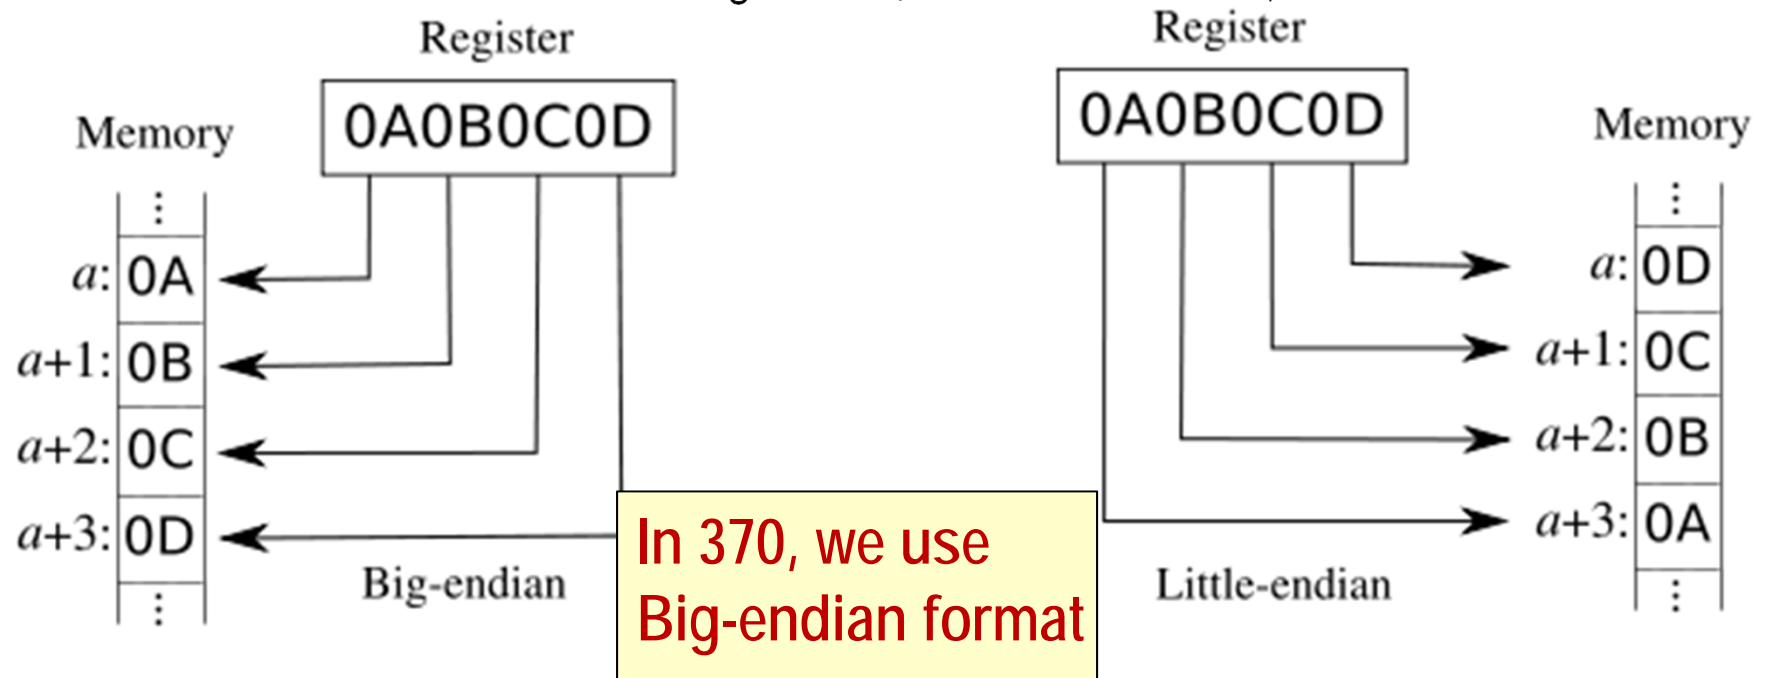


Memory
(each location is 1 byte)



Important note on addressing: Big-endian vs. Little-endian

- ❑ Endian-ness: ordering of bytes within a word
 - Little - increasing numeric significance with increasing memory addresses
 - Big – The opposite, most significant byte first
 - Internet and SPARC are big endian, x86 is little endian, ARM is bi-endian




Example Code Sequence – insn 1


What is the final state of memory once you execute the following instruction sequence? (assume r0 stores the value 0)

```
ldr    r4, [r0, #100]
ldrsb  r3, [r0, #102]
str     r3, [r0, #100]
strb    r4, [r0, #102]
```

register file



Memory
(each location is 1 byte)



Example Code Sequence – insn 2

What is the final state of memory once you execute the following instruction sequence? (assume r0 stores the value 0)

```
ldr    r4, [r0, #100]
ldrsb  r3, [r0, #102]
str     r3, [r0, #100]
strb    r4, [r0, #102]
```

register file

r3	0xFFFF FFFF
r4	0x0203 FF05

Memory
(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103

Example Code Sequence – insn 3

What is the final state of memory once you execute the following instruction sequence? (assume r0 stores the value 0)

```
ldr    r4, [r0, #100]
ldrsb  r3, [r0, #102]
str     r3, [r0, #100]
strb    r4, [r0, #102]
```

register file

r3	0xFFFF FFFF
r4	0x0203 FF05

Memory
(each location is 1 byte)

0xFF	100
0xFF	101
0xFF	102
0xFF	103

Example Code Sequence – insn 4

What is the final state of memory once you execute the following instruction sequence? (assume r0 stores the value 0)

```
ldr    r4, [r0, #100]
ldrsb  r3, [r0, #102]
str     r3, [r0, #100]
strb   r4, [r0, #102]
```

register file

r3	0xFFFF FFFF
r4	0x0203 FF05

Memory
(each location is 1 byte)

0xFF	100
0xFF	101
0x05	102
0xFF	103

Class Problem 3


What is the final state of memory once you execute the following instruction sequence? (assume r0 stores the value 0)

```
ldrh    r3, [r0, #100]
ldrb    r4, [r0, #102]
str      r3, [r0, #100]
strh     r4, [r0, #102]
```

register file



Memory
(each location is 1 byte)



ARM Sequencing Instructions

- ❑ Sequencing instructions change the flow of instructions that are executed
 - This is achieved by modifying the program counter (r15)
- ❑ Conditional branches
 - If (*condition_test*) goto *target_address*
 - *condition_test* examines flags from the processor status word (PSR)
 - *target_address* is a 24-bit word displacement on current PC+8

- `cmp r1, r2`
`beq label`
- if ($r1 == r2$) then $PC = \text{label}$ else $PC = PC + 4$

ARM Condition Codes Determine Branch Direction

❑ Most arithmetic/logic instructions can set condition codes in PSR

- add, sub, cmp, and, eor, etc...
- You need to add "S" to the instruction: adds, subs, ands,...

❑ Four primary condition codes evaluated:

- N – set if the result is **negative** (i.e., bit 31 is non-zero)
- Z – set if the result is **zero** (i.e., all 32 bits are zero)
- C – set if last addition/subtraction had a **carry**/borrow out of bit 31
- V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produced a positive result)

❑ Branch conditions:

- eq – (Z == 1) gt – (Z == 0 && N == V) mi - ?
- ne – (Z == 0) le – (Z == 1 || N != V) pl – ?
- ge – (N == V) lt – (N != V) al – 1 (can use shorthand "b label")

Setting the Branch Displacement Field

- ❑ Target address is a 24-bit signed aligned displacement on current PC+8
 - $\text{Target} = \text{PC} + 8 + 4 * \text{24_bit_signed_displacement}$
 - `beq 1` // branch **3** instructions ahead if flag Z == 1
 - `beq -3` // branch **1** instruction back if flag Z == 1
 - `beq -2` // Infinite loop if flag Z == 1

Example code sequence

PC+8



add
sub
mul
beq
ldr
str
mov

Other Branching Instructions

- `cmp r2, r3` `// branch to 4*offset+PC+8 if r2 \neq r3`
`bne offset`
- `cmp r2, r3` `// branch to 4*offset+PC+8 if r2 < r3`
`blt offset`
- `cmp r2, r3` `// branch to 4*offset+PC+8 if r2 \geq r3`
`bge offset`
- `b offset` `// jump to 4*offset+PC+8`
 `// unconditional jump, is taken regardless of PSR flags`
- `mov r15, r3` `// jump to address in r3` -- when is this useful?
- `bl offset` `// put PC+4 into register r14 (LR) and jump to 4*offset+PC+8`

Branch - Example

Convert the following C code into ARM assembly (assume x is in r1, y in r2):

```
int x, y;  
if (x == y)  
    x++;  
(L1) else  
    y++;  
(L2) ...
```

Using Labels

```
cmp r1, r2  
bne L1  
add r1, r1, #1  
b L2  
L1: add r2, r2, #1  
L2: ...
```

Without Labels

```
cmp r1, r2  
bne 1  
add r1, r1, #1  
b 0  
add r2, r2, #1
```

Assemblers should deal with labels and assign displacements – why?