

5. Instruction Set Architecture – -from C to Assembly – function calls

EECS 370 – Introduction to Computer Organization - Winter 2016

Profs. Todd Austin & Don Winsor

EECS Department
University of Michigan in Ann Arbor, USA

© Bertacco-Das, 2016

The material in this presentation cannot be
copied in any form without our written permission

Announcements

- ❑ Due today
 - Project 1.a
 - Homework 1
- ❑ Reading assignments available in LECTURES tab of the website

❑ Correction: **C: $a = b + \text{names}[i]$;**

~~mul r5, r3, #4~~ // calc mov r5, r3, LSL #2 // r5 = r3 << 2

ldr r4, [r5, #1000] // load names[i]

add r1, r2, r4 // calculate b + names[i]

Recap

❑ ARM – branch instructions

❑ C –to assembly

- Basic statements
- Memory alignment
- If-then-else

Loop – Example

```
// assume all variables are integers
// i is in r1, start of a is at address 500, sum is in r2
for (i=0 ; i < 100 ; i++) {
    if (a[i] > 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $3 \times 100 + 1 = 301$

a.k.a. while-do template

```
    mov r1, #0
    mov r4, #400
Loop1: cmp r1, r4
       bge endLoop
       ldr r5, [r1, #500]
       cmp r5, #0
       ble endIf
       add r2, r2, r5
endIf: add r1, r1, #4
       b   Loop1
endLoop:
```

Same Loop, Different Assembly

```
// assume all variables are integers
// i is in r1, start of a is at address 500, sum is in r2
for (i=0 ; i < 100 ; i++) {
    if (a[i] > 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $2 \times 100 = 200$

a.k.a. do-while template

```
mov r1, #0
mov r4, #400
Loop1: ldr r5, [r1, #500]
      cmp r5, #0
      ble endIf
      add r2, r2, r5
endIf: add r1, r1, #4
      cmp r1, r4
      blt Loop1
endLoop:
```

Class Problem 2

Write the ARM assembly code to implement the following C code:

```
// assume ptr is in r1
// struct {int val; struct node *next;} node;
// struct node *ptr;

if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

ARM conditional assembly

- ❑ All ARM instructions can be executed conditionally, not just branches
- ❑ Conditional assembly programs may boost performance because execution penalty is < branch misprediction

Example: while (r0 != r1) {
 if (r0>r1) r0 = r0 - r1;
 else r1 =r1 – r0;

 }
standard assembly

```
TOP      cmp r0, r1
        beq END
        blt LESS
        sub r0, r0, r1
        b TOP
LESS     sub r1, r1, r0
        b TOP
END     ...
```

conditional assembly

```
TOP      cmp r0, r1
        subgt r0, r0, r1
        sublt r1, r1, r0
        bne TOP
        ...
```

Instruction Set Architecture (ISA) Design Lectures

- ❑ Lecture 2: Storage types and addressing modes
- ❑ Lecture 3 : LC-2K and ARM architecture
- ❑ Lecture 4 : Converting C to assembly – basic blocks
- ❑ **Lecture 5 : Converting C to assembly – functions**
- ❑ Lecture 6 : Translation software; libraries, memory layout

Converting function calls to assembly code

C: `printf("hello world\n");`

- Need to pass parameters to the called function (`printf`)
- Need to save return address of caller
- Need to save register values
- Need to jump to `printf`



Execute instructions for `printf()`
Jump to to return address



- Need to get return value (if used)
- Restore register values

Task 1: Passing parameters

- ❑ Where should you put all of the parameters?
 - Registers?
 - Fast access but few in number and wrong size for some objects
 - Memory?
 - Good general solution but where?

- ❑ ARM answer:
 - Registers and memory
 - Put the first few parameters in registers (if they fit) (r0 – r3)
 - Put the rest in memory on the call stack

 - Example:
 `mov r0, #1000 // put address of char array “hello world” in r0`

Call stack

- ❑ ARM conventions (and most other processors) allocate a region of memory for the call stack
 - This memory is used to manage all the storage requirements to simulate function call semantics
 - Parameters (that were not passed through registers)
 - Local variables
 - Temporary storage (when you run out of registers and need somewhere to save a value)
 - Return address
 - Etc.
- ❑ Sections of memory on the call stack [a.k.a. **stack frames**] are allocated when you make a function call, and de-allocated when you return from a function.

ARM (Linux) Memory Map

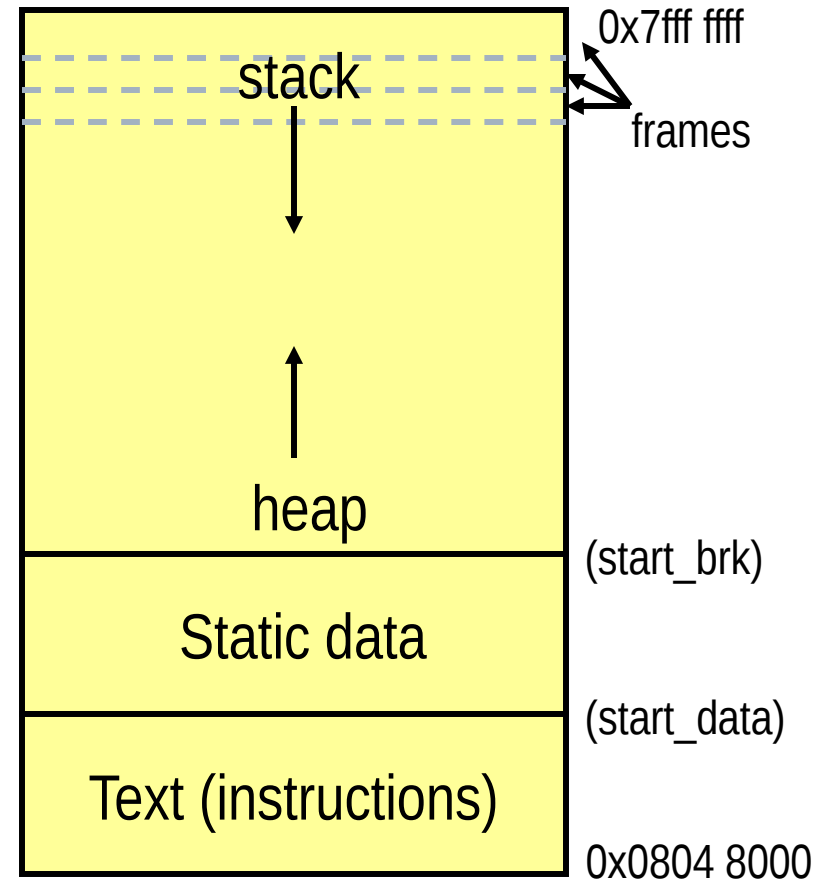
FUNCTION CALLS

Stack: starts at 0x7fff ffff and grows down to lower addresses. Bottom of the stack resides in the SP register

Heap: starts above static (page aligned) and grows up to higher addresses. Allocation done explicitly with malloc(). Deallocation with free(). Runtime error if no free memory before running into SP address.


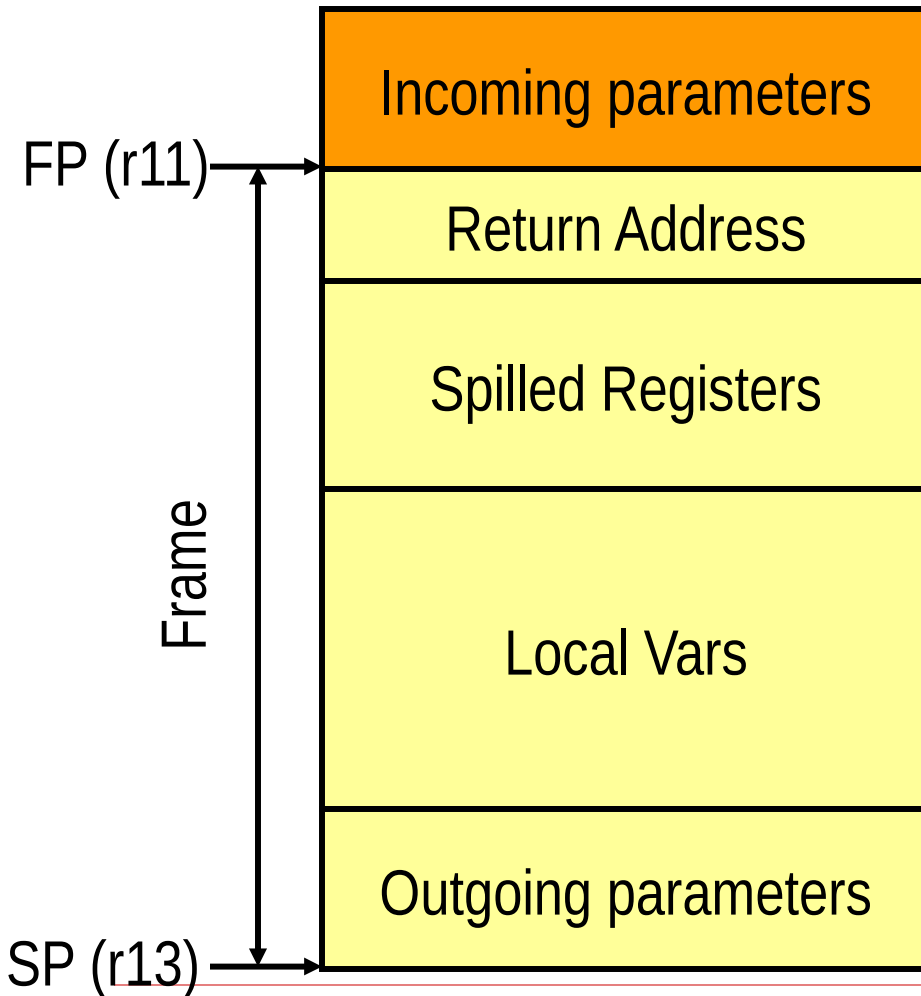
Static: starts above text (page aligned). Holds all global variables and those locals explicitly declared as “static”.

Text: starts at 0x08048000. Holds all instructions in the program (except for Dynamically linked library routines DLLs)



The ARM Stack Frame

FUNCTION CALLS



Allocating space to local variables

- ❑ Local variables (by default) are created when you enter a function, and disappear when you leave
 - Technical terminology: local variables are placed in the automatic storage class (as opposed to the static storage class used for globals).
- ❑ Automatics are allocated on the call stack
 - How?
by incrementing (or decrementing) the pointer to the top of the call stack
 - ```
sub r13, r13, #12 // SP = SP - 12, allocate space for 3 integer locals
add r13, r13, #12 // SP = SP + 12, de-allocate space for locals
```


# The stack grows as functions are called

FUNCTION CALLS

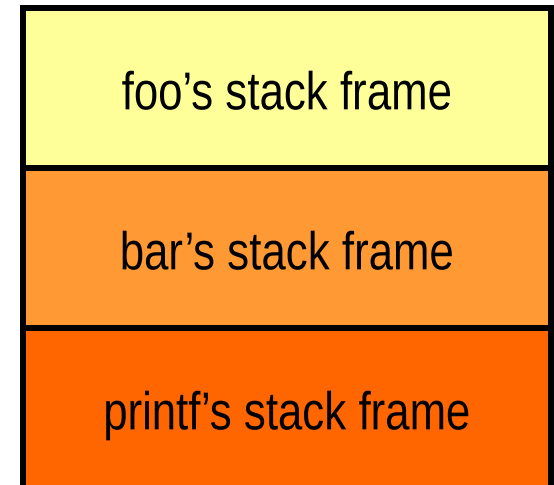
```
void foo()
```

```
{
 int x, y[2];
 bar(x);
}
```

inside foo




bar calls printf



```
void bar(int x)
```

```
{
 int a[3];
 printf();
}
```


foo calls bar



# The stack shrinks as functions return


```
void foo()
{
 int x, y[2];
 bar(x);
}
```

printf returns



```
void bar(int x)
{
 int a[3];
 printf();
}
```

bar returns





## Stack frame contents

```
void foo()
{
 int x, y[2];
 bar(x);
}
```

```
void bar(int x)
{
 int a[3];
 printf();
}
```

inside foo – foo's stack frame

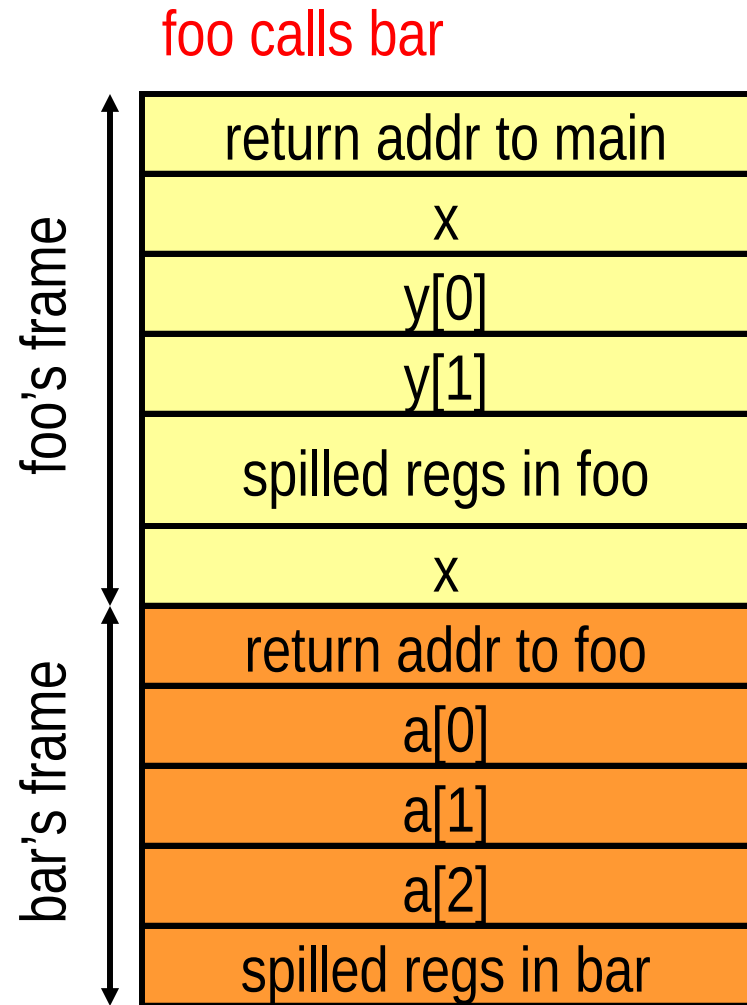
|                     |
|---------------------|
| return addr to main |
| x                   |
| y[0]                |
| y[1]                |
| spilled regs in foo |

## Stack frame contents (2)

FUNCTION CALLS

```
void foo()
{
 int x, y[2];
 bar(x);
}
```

```
void bar(int x)
{
 int a[3];
 printf();
}
```



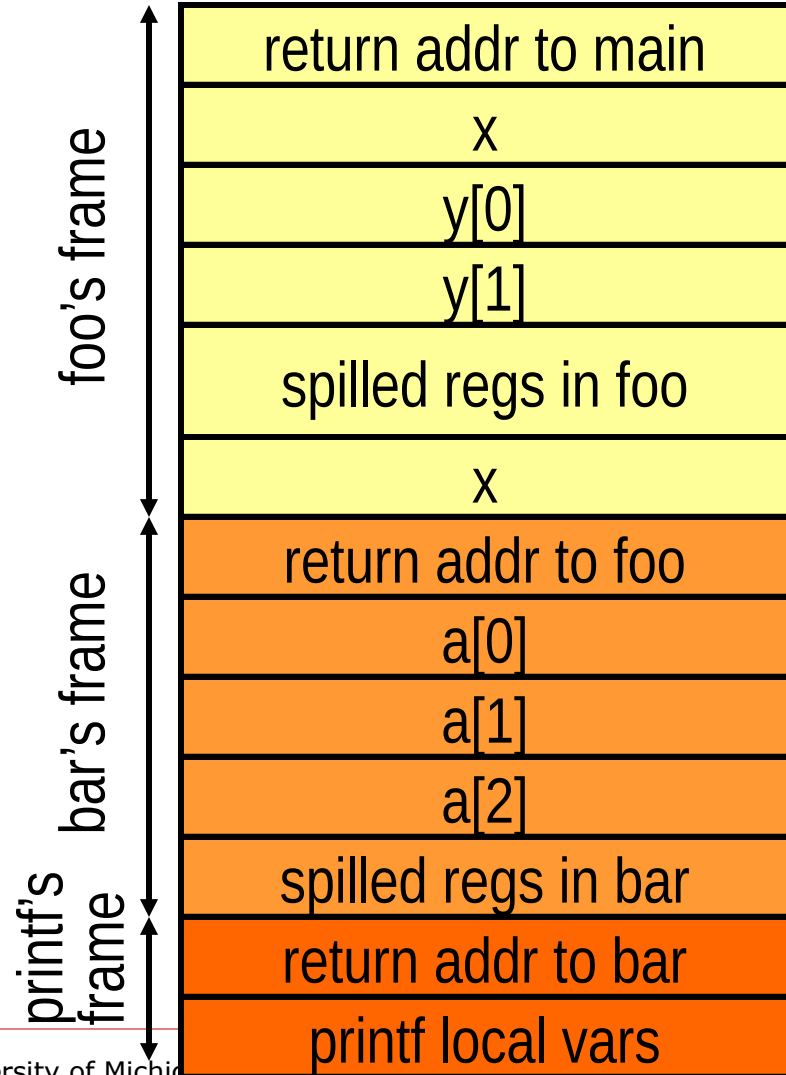
## Stack frame contents (3)

bar calls printf

FUNCTION CALLS

```
void foo()
{
 int x, y[2];
 bar(x);
}
```

```
void bar(int x)
{
 int a[3];
 printf();
}
```



# Recursive function example

FUNCTION CALLS

```
main()
```

```
{
 foo(2);
}
```

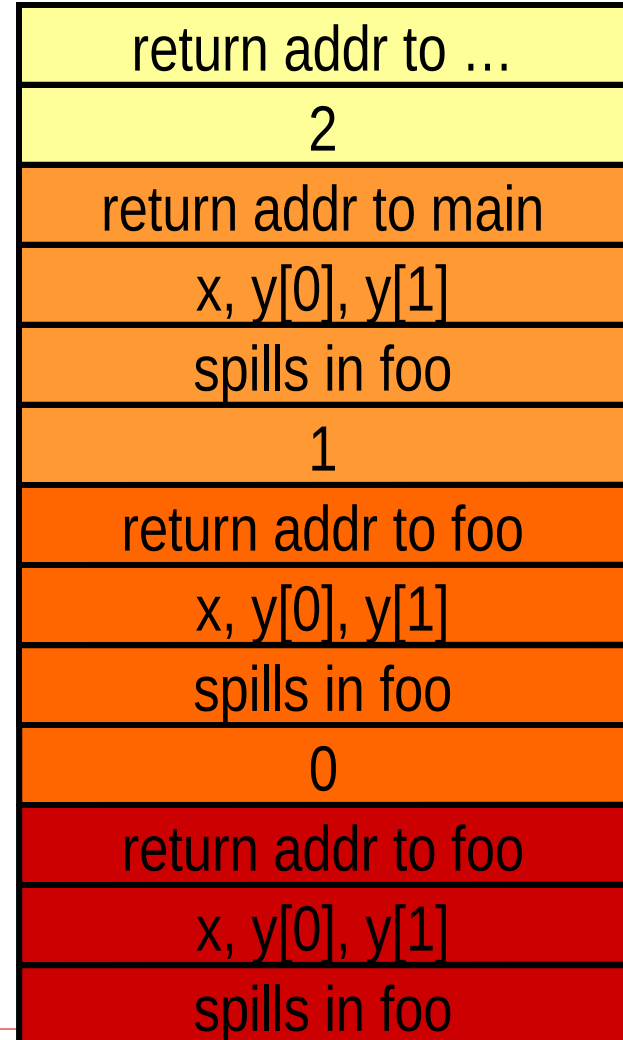
main calls foo

```
void foo(int a)
```

```
{
 int x, y[2];
 if (a > 0)
 foo(a-1);
}
```

foo calls foo

foo calls foo



# Virtual functions

---

- ❑ Call stack is identical
- ❑ key difference: call is implemented as table lookup (i.e., indirect call versus direct call)
  - “class->method()” is translated to:
    - “switch (class.type)
    - case A: r1 = &class.method\_typeA
    - case B: r1 = &class.method\_typeB
    - etc.
    - call r1”

# Assigning variables to memory spaces

FUNCTION CALLS

```
int w;
void foo(int x)
```

w goes in static, as it's a global

x goes on the stack, as it's a parameter


```
{
 static int y[4];
 char *p;
 p = malloc(10);
 ...
 printf("%s\n", p);
}
```

y goes in static, 1 copy of this!!

p goes on the stack

allocate 10 bytes on heap, ptr  
set to the address

string goes in static, pointer  
to string on stack, p goes on  
stack



# Need for Saving registers during a call

- What happens to the values we have in registers when we make a function call? Assume variables `x` in `foo()` and `y` in `bar()` happen to be allocated to the same register `r1`.

```
void foo()
{
 int x = 1;
 bar(x);
 x = x + 1;
}
```

```
void bar(int k)
{
 int y = 2;
 y++
}
```

```
void foo()
{
 r1 = 1;
 bar(r1);
 r1 = r1 + 1;
}
```

```
void bar(int k)
{
 r1 = 2;
 r1 = r1 + 1
}
```

# Saving registers during a call

- What happens to the values we have in registers when we make a function call?  
$$a = b * c + \text{sqrt}(d);$$

Options:

1. You can save your registers **before** you make the function call and restore the registers when you return (**caller-save register**). Where?

The stack frame is used to store anything required to support function calls

What if the function you are calling doesn't use that register? No harm done, but wasted work!!!


2. You can save your registers **after** you make the function call and restore the registers before you return (**callee-save register**). Where?

What if the caller function doesn't use that register? No harm done, but wasted work!!!

**Most assembly programming conventions use a mix of both**



# Caller-Callee save/restore



**Caller save:** Callee may change, so caller responsible for saving immediately before call and restoring immediately after call

**Callee save:** Callee may not change, so callee (called function) must leave these unchanged. Can be ensured by inserting saves at the start of the function and restores at the end

# Caller/callee example

CALLER-CALLEE

```
foo()
{
 r0 = 5;
 r4 = -1;
 bar();
 r3 = r0 + r4;
}
```

If r0 is caller-save  
and r4 is callee-  
save



```
bar()
{
 r0 = 10;
 r4 = 5;
}
```

```
foo()
{
 r0 = 5;
 r4 = -1;
 save r0; i.e., str r0, [r13, #20]
 bar();
 restore r0; i.e., ldr r0, [r13, #20]
 r3 = r0 + r4;
}
```

```
bar()
{
 save r4; i.e., str r4, [r13, #8]
 r0 = 10;
 r4 = 5;
 restore r4; i.e., ldr r4, [r13, #8]
}
```

# Saving/Restoring Optimizations

---

## ❑ Caller-saved

- Only needs saving if it is “live” across a function call
- **Live** = contains a useful value: Assign value before function call, use that value after the function call
- In a leaf function, caller saves can be used without saving/restoring

## ❑ Callee-saved

- Only needs saving at beginning of function (generally infrequent as outside of loops) and restoring at end of function
- Only save/restore it if function overwrites the register

## ❑ Each has its advantages. Neither is always better.

# ARM caller-callee convention

- Have some registers that are caller-saved, and some that are callee-saved

Example: ARM

- 5 caller saved
- 10 callee saved

## ARM register conventions

|        |                                       |
|--------|---------------------------------------|
| r0     | parameter, return value, caller saved |
| r1-r3  | parameters, caller saved              |
| r4-r10 | callee saved                          |
| r11    | frame pointer, callee saved           |
| r12    | caller saved                          |
| r13    | stack pointer, callee saved           |
| r14    | link register, callee saved           |
| r15    | program counter, not saved            |

- Choose registers for variables so to minimize the number of dynamic saves/restores

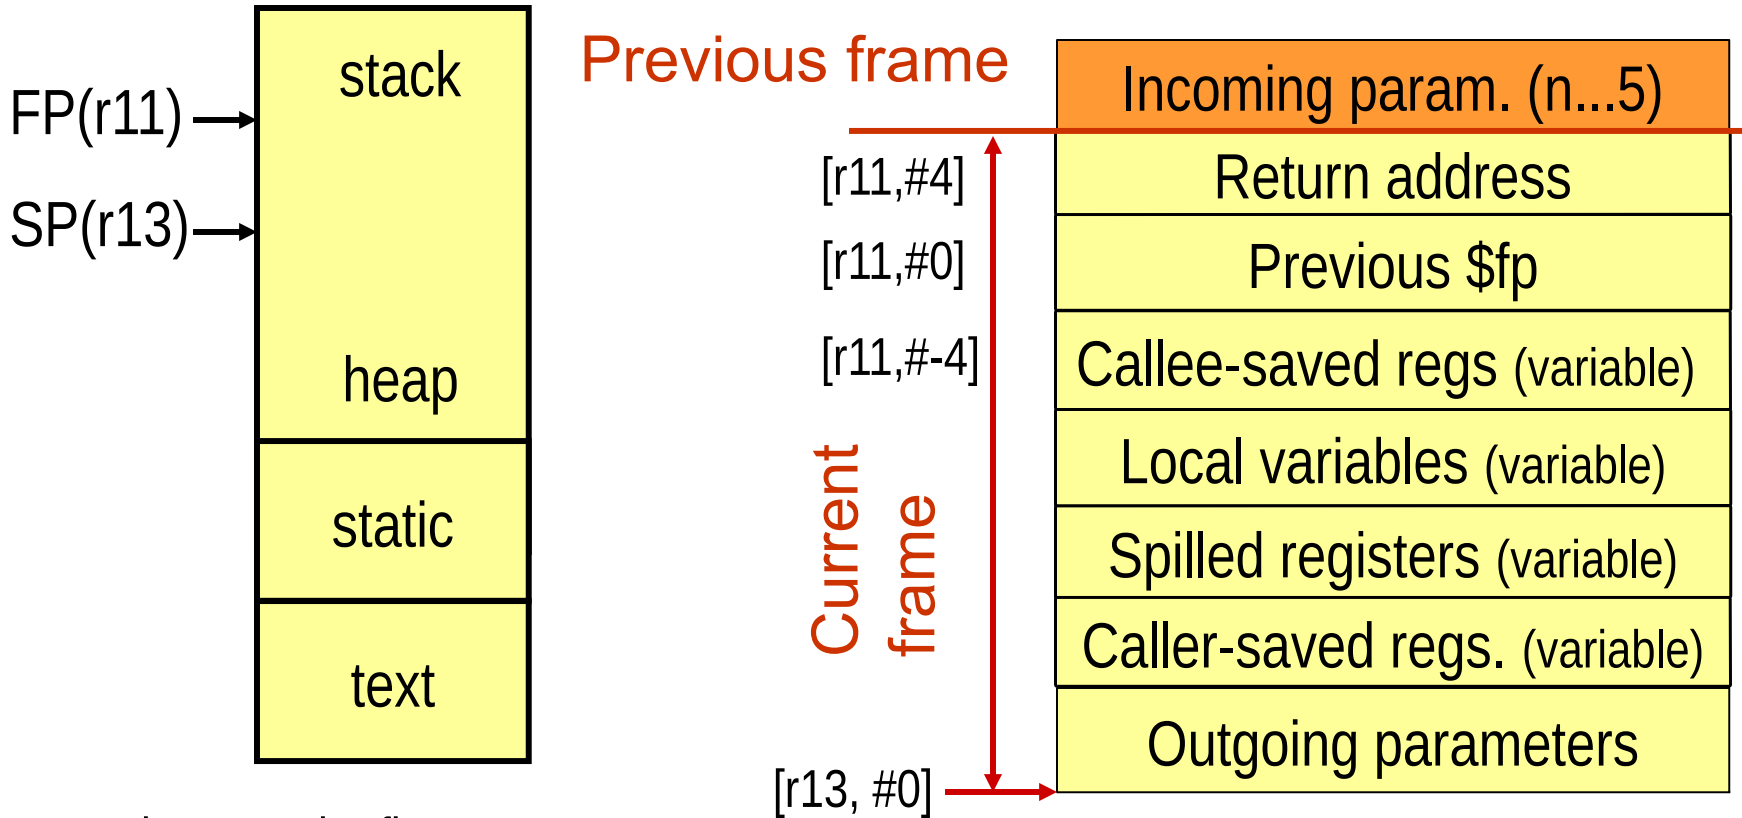
# Calling convention

---

- ❑ This is a **convention**: calling convention
  - There is no difference in H/W between caller and callee save registers
  
- ❑ Passing parameters in registers is also a convention
- ❑ Allows assembly code written by different people to work together
  - Need conventions about who saves regs and where args are passed.
- ❑ These conventions collectively make up the ABI or “application binary interface”
  
- ❑ Why are these conventions important?
  - What happens if a programmer/compiler violates them?

# ARM Stack Frame (typical organization)

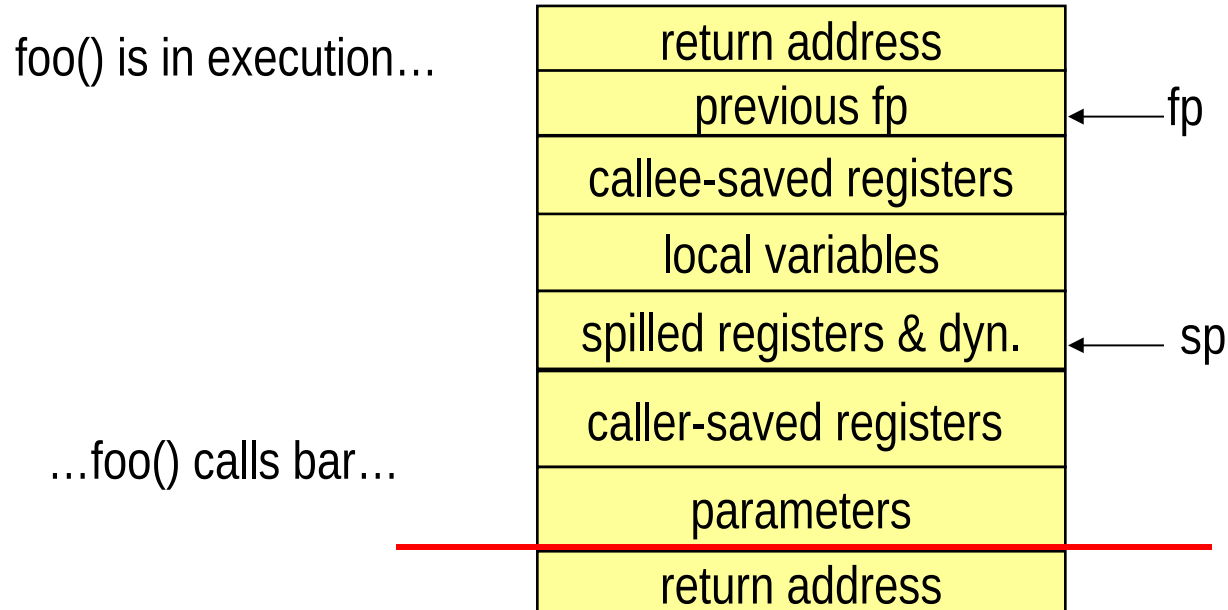
CALLER-CALLEE



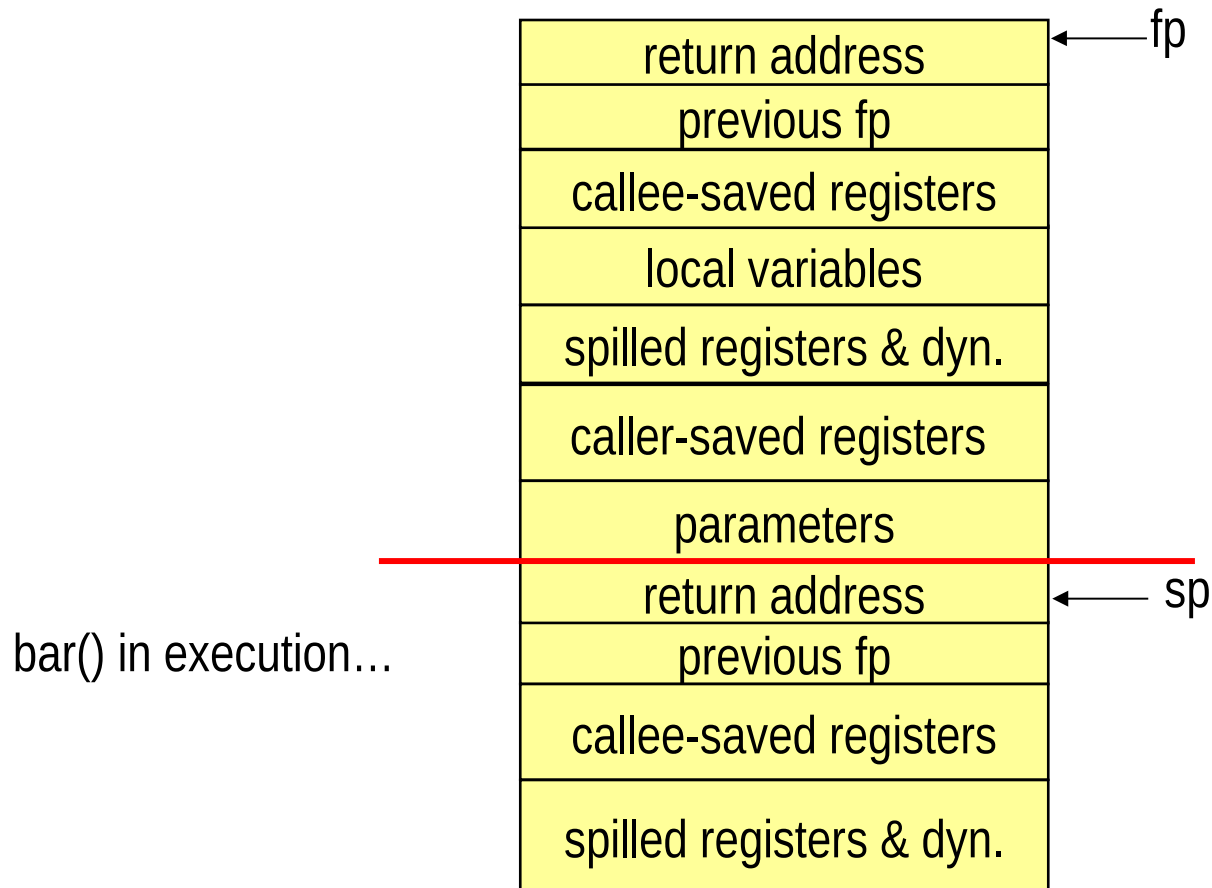
Note 1: in ARM, the first 4 parameters are passed via registers. Other ISAs have  $\neq$  conventions

Note 2: why is the last parameter first on the stack ?

# The stack during program execution...

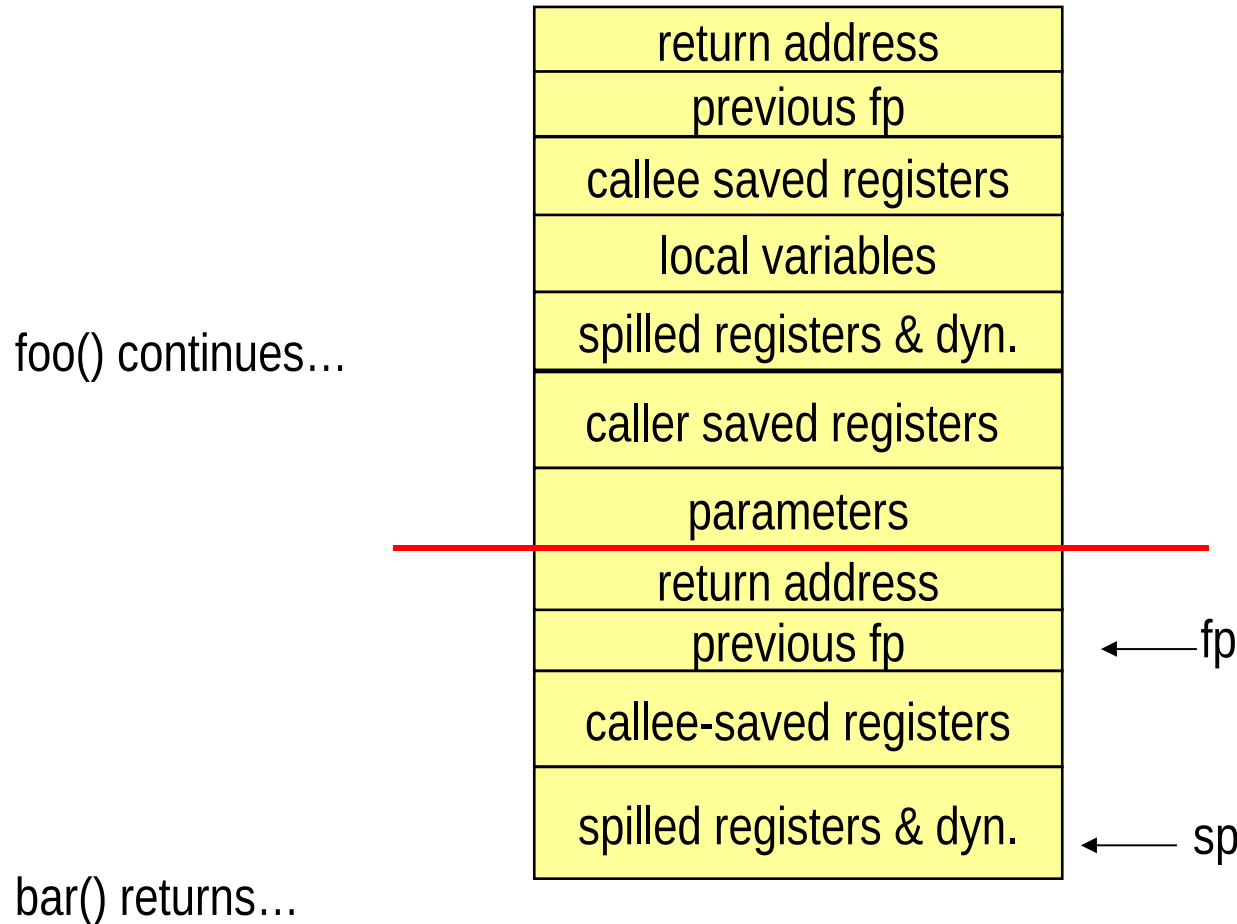


# The stack





# The stack



## Putting it all together (using activation records)

---

```
mov r0, #1000 // param "hello world\n"
str r12, [r13, #24] // caller save r12
bl _printf // call printf()
 // r0 holds return value (ignored)
ldr r12, [r13, #24] // restore caller-saved r12 value
```

```
sub r13, r13, #16 // allocate space for 2 locals +
str r4, [r13, #8] // callee save r4
str r14, [r13, #12] // save return address
... // function body
mov r0, #0 // return value 0
ldr r4, [r13, #8] // restore callee-saved r4
ldr r14, [r13, #12] // restore return address
add r13, r13, #16 // deallocate call frame
mov r15, r14 // return to calling function
```

# Calculating Caller/Callee Costs

---

Consider the cost of placing each variable  $v$  from function  $f$  in a callee register and a caller register:

**Cost = number of store/load instructions required to accomplish the required saving/restoring**

Callee\_cost  $\sqsubseteq$  save at the start of the function, restore at end  
 $= 2 * \text{number of invocations of } f$

Caller\_cost  $\sqsubseteq$  potentially save/restore across each funct. call in  $f$   
Caller cost = 0

For each function call in  $f$ , call <sub>$i$</sub>

if ( $v$  is live) caller\_cost  $+= 2 * \text{number of times call}_i$  is executed

# Caller/Callee Selection

---

- ❑ Select assignment of variables to registers such that the sum of caller/callee costs is minimized
  - Execute fewest save/restores
- ❑ Each function greedily picks its own assignment ignoring the assignments in other functions
  - Calling convention assures all necessary registers will be saved
- ❑ 2 types of problems
  1. Given a single function  $\Rightarrow$  Assume it is called 1 time
  2. Set of functions or program  $\Rightarrow$  Compute number of times each function is called if it is obvious (i.e., loops with known trip counts or you are told)

# Assumptions

---

- ❑ A function can be invoked by many different call sites in different functions.
- ❑ Assume no inter-procedural analysis (hard problem)
  - A function has no knowledge about which registers are used in either its caller or callee
  - Assume main() is not invoked by another function
- ❑ Implication
  - Any register allocation optimization is done using function local information

## Class Problem 3

```
foo() {
 a = ...
 b = ...
 bar();
 ... = a;
 ... = b;
 for (1 to 15) {
 c = ...
 d = ...
 ... = c;
 printf();
 ... = d;
 }
}
```

Assume that you have 2 caller and 2 callee save registers. Pick the best assignment for a, b, c, d. Assume each requires its own register.

# Caller-saved vs. callee saved – Multiple function case

```
void main() {
 int a,b,c,d;
 .
 c = 5; d = 6;
 a = 2; b = 3;
 foo();
 d = a+b+c+d;
 .
 .
 .
}
```

```
void foo() {
 int a,b;
 .
 .
 a = 2; b = 3;
 bar();
 a = a + b;
 .
 .
 .
}
```

```
void bar() {
 int a,b,c,d;
 .
 c = 0; d = 1;
 a = 2; b = 3;
 final();
 a = a+b+c+d;
 .
 .
 .
}
```

```
void final() {
 int a,b,c;
 .
 .
 a = 2; b = 3;
 .
 c = a+b;
 .
 .
 .
}
```

Note: assume main does not have to save any callee reg. (that is really the case for start)

# Caller-saved vs. callee saved – Multiple function case

---

## Questions:

1. In assembly code, how many regs. need to be stored/loaded in total if we use a **caller-save** convention ?
2. In assembly code, how many regs. need to be stored/loaded in total if we use a **callee-save** convention ?
3. In assembly code, how many regs. need to be stored/loaded in total if we use a mixed **caller/callee**-save convention with 3 callee-s. and 3 caller-s. registers ?
4. Assume bar() is in a loop inside foo() and the loop is iterated 10 times ? When the program is executed, how many regs. need to be stored/loaded in total for each of the above three scenarios?



# Question 1: Caller-save

---

```
void main() {
 .
 .
 .
 [4 str]
 foo();
 [4 ldr]
 .
 .
 .
}
```

```
void foo() {
 .
 .
 .
 [2 str]
 bar();
 [2 ldr]
 .
 .
 .
}
```

```
void bar() {
 .
 .
 .
 [4 str]
 final();
 [4 ldr]
 .
 .
 .
}
```

```
void final() {
 .
 .
 .
 .
 .
 .
 .
 .
 .
}
```

Total: 10 str / 10 ldr

## Question 2: Callee-save

---

```
void main() {
 .
 .
 .
 .
 foo();
 .
 .
 .
 .
}
```

```
void foo() {
 [2 str]
 .
 .
 .
 bar();
 .
 .
 .
 [2 ldr]
}
```

```
void bar() {
 [4 str]
 .
 .
 .
 final();
 .
 .
 .
 [4 ldr]
}
```

```
void final() {
 [3 str]
 .
 .
 .
 .
 .
 .
 .
 [3 ldr]
}
```

Total: 9 str / 9 ldr

## Question 3: Mixed 3 caller / 3 callee

---

```
void main() {
 .
 .
 .
 [1 str]
 foo();
 [1 ldr]
 .
 .
 .
}
```

1 caller r.  
3 callee r.

```
void foo() {
 [2 str]
 .
 .
 .
 bar();
 .
 .
 .
 [2 ldr]
}
```

```
void bar() {
 [4 str]
 .
 .
 .
 .
 final();
 .
 .
 .
 .
 [4 ldr]
}
```

```
void final() {
 .
 .
 .
 .
 .
 .
 .
 .
 .
 .
}
```

3 caller r.

Total: 7 str / 7 ldr

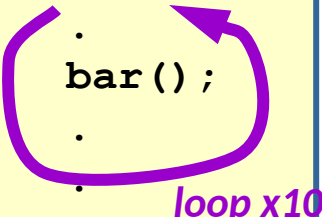
# Caller-saved vs. callee saved – Question 4

❑ Mixed 3 caller / 3 callee

```
void main() {
 .
 .
 .
 [1 str]
 foo();
 [1 ldr]
 .
 .
 .
}
```

1 caller r.  
3 callee r.

```
void foo() {
 [2 str]
 .
 .
 .
 bar();
 .
 .
 .
 [2 ldr]
}
```



2 callee r.

```
void bar() {
 [4 str]
 .
 .
 .
 .
 final();
 .
 .
 .
 [4 ldr]
}
```

*x10*

```
void final() {
 .
 .
 .
 .
 .
 .
 .
 .
 .
 .
}
```

3 caller r.  
*x10*

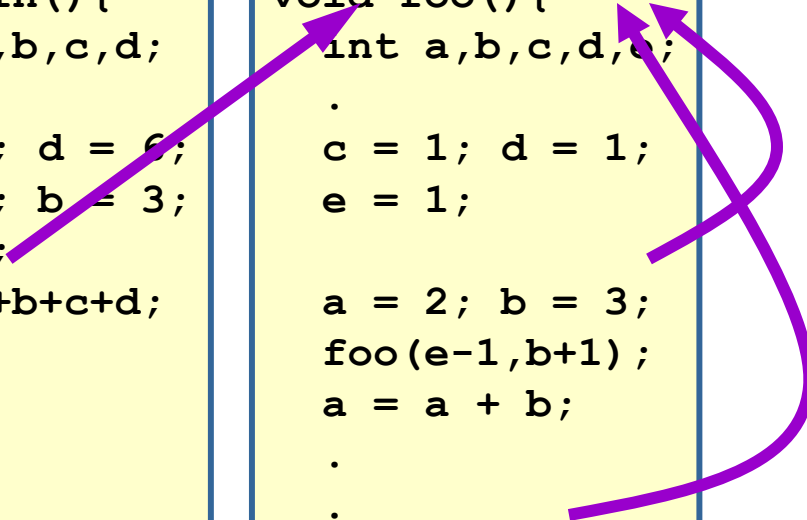
Total: 43 str / 43 ldr

Pure caller: (4+20+40+0) str / ldr - Pure callee (0+2+40+30) str / ldr

# Caller-saved vs. callee saved – A more interesting case

```
void main(){
 int a,b,c,d;
 .
 c = 5; d = 6;
 a = 2; b = 3;
 foo();
 d = a+b+c+d;
 .
 .
 .
}
```

```
void foo(){
 int a,b,c,d,e;
 .
 c = 1; d = 1;
 e = 1;
 .
 a = 2; b = 3;
 foo(e-1,b+1);
 a = a + b;
 .
 .
 a = 5, b = 4;
 foo(b,9);
 b = a - b;
 .
 c++; d++; e++;
}
```



The diagram illustrates the control flow between the `main` and `foo` functions. A purple arrow originates from the `foo();` call in `main` and points to the start of the `foo` function. Another purple arrow originates from the closing brace of the `foo` function and points back to the line immediately following the `foo();` call in `main`, representing the return path.

# Caller-saved vs. callee saved – the interesting case

---

- ❑ Assume the function foo() is called recursively 15 times in total
- ❑ When the program is executed, how many regs need to be stored/loaded in total for the following scenarios:
  - Use a **caller-save** convention ?
  - Use a **callee-save** convention ?
  - Use a mixed **caller/callee**-save convention with 3 callee-s. and 3 caller-s. registers ?