

# 15. Cache organization: The basics

---

EECS 370 – Introduction to Computer Organization - Winter 2016

**Profs. Valeria Bertacco & Reetu Das**

EECS Department  
University of Michigan in Ann Arbor, USA

© Bertacco-Das, 2016

The material in this presentation cannot be  
copied in any form without our written permission

# Memory

---

- ❑ So far, we have discussed two structures that hold data:
  - Register file (little array of words)
  - Memory (bigger array of words)
  
- ❑ We have discussed several methods of implementing storage devices:
  - Static memory (made with logic gates)
  - Dynamic memory (transistor and capacitor)
  - ROM, and other ROM-like storage, e.g., flash (floating gate transistors)

# Memory Hierarchy

---

- ❑ We want a lot of memory
  - LC2 can handle  $2^{18}$  bytes of memory
  - MIPS can handle  $2^{32}$  bytes of memory
  - Athlon-64 or EM64T can handle  $2^{64}$  bytes of memory
  
- ❑ What are our choices?
  - SRAM, DRAM, ROM, disk, tape, DVD?

## Option 1: SRAM

---

- ❑ Fast:  $\sim 2\text{ns}$  access time or faster
  - Decoders are big
  - Array is big
    - Why?
  
- ❑ Expensive, high area requirement
  - SRAM: \$5.0 per megabyte
    - \$0.13 for LC2
    - \$20,000 for MIPS
    - \$88,000,000,000,000 for Athlon-64

## Option 2: DRAM

---

- ❑ Slower: ~60ns access time
  - Hurry up and wait design philosophy doesn't work
  - Must stall for dozens of cycles on each memory load
  
- ❑ Less expensive than SRAM.
  - DRAM costs \$0.012 per megabyte
    - \$0.00 for LC2
    - \$50 for MIPS/Pentium-IV/Athlon-XP
    - \$210,000,000,000 for Alpha/G5/x86\_64

## Option 3: Flash

---

- ❑ Slower still: ~250ns access time
  - Hurry up and wait design philosophy doesn't work
  - Must stall for dozens of cycles on each memory load
- ❑ Less expensive than SRAM
  - Flash costs \$0.0012 per megabyte
    - \$0.00 for LC2
    - \$4.9 for MIPS/Pentium-IV/Athlon-XP
    - \$21,000,000,000 for Alpha/G5/x86\_64
- ❑ Non-volatile

## Option 4: Disks

---

- ❑ Obnoxiously slow: 3,000,000ns access time
  - We could have stopped with the Intel 4004
- ❑ Cheap
  - Disk storage costs \$0.000043 per megabyte
    - \$0.00 LC2
    - \$0.18 for MIPS
    - \$760,000,000 for Athlon-64
- ❑ Non-volatile

# Memory Hierarchy Goals

---

- ❑ **Fast:** Ideally run at processor clock speed
  - 1 ns access
- ❑ **Cheap:** Ideally free
  - Not more expensive than rest of system
- ❑ Options DRAM, flash, disks are too slow
- ❑ Option SRAM is too expensive
- ❑ How to get best properties of multiple memory technologies?




# Memory Hierarchy Goals

move what you use here

With good locality of reference, memory appears as fast as and as large as

backup everything here




# Memory Hierarchy

---

- ❑ Use a small array of SRAM
  - Small so fast and cheap
  - For the [cache](#) (hopefully covers most loads and stores)
- ❑ Use a larger amount of DRAM
  - Cheaper than SRAM, faster than flash/disk
  - For the [main memory](#)
- ❑ Use a lot of flash and/or disk
  - Non-volatile. Cheap. Big
  - For [Virtual memory](#)
- ❑ Don't try to buy  $2^{64}$  bytes of anything
  - Use “virtual memory” to make it look like the entire address range is available
  - A few TB is enough for most desktop machines today, or a smartphone in a few years

# Memory hierarchy

---



# Definitions

---

- ❑ The architectural view of memory is
  - What the machine language (or programmer) sees
  - Just a big array
- ❑ Breaking up the memory system into different pieces – cache, main memory (made up of DRAM) and Disk – is not architectural
  - The machine language doesn't know about it
  - A new implementation may not break it up in the same way

# Function of the Cache

---

- ❑ The cache will hold the data that we think is **most likely** to be **referenced**
  - Because we want to maximize the number of references that are serviced by the cache to minimize the average memory access latency
  - How do we decide what the most likely accessed memory locations are?

# Cache Analogy

---

## ❑ Studying books in library

- **Option 1: Every time you switch to another book, return current book to shelf and get new book from shelf**
  - Latency = 5 minutes
- **Option 2: Keep 10 commonly-used books on shelf above desk**
  - Latency = 1 minute
- **Option 3: Keep three books open to appropriate locations on desk**
  - Latency = 10 seconds

## Example Problem

---

Given the following:

Cache has 1 cycle access time

Main memory has 100 cycle access time

Disk has 10,000 cycles access time

What is the average access time for 100 memory references if 90% of the cache accesses are hits and 80% of the accesses to main memory are hits?

## Basic Cache Design

---

- ❑ Cache memory can copy data from any part of main memory. It has 2 parts:
  - The **TAG** (CAM) holds the memory address
  - The **BLOCK** (SRAM) holds the memory data

addr	data
addr	data

**TAG    BLOCK**

- ❑ Accessing the cache: compare reference address with the tag
  - If they match, get the data from the cache block
  - If they don't match, get the data from main memory



# CAMs: content addressable memories

---

- ❑ Instead of thinking of memory as an array of data indexed by a memory address,
- ❑ Think of memory as a set of data matching a query
  - Instead of an address, we send a key to the memory, asking whether the key exists and, if so, what value it is associated with
  - Memory answers: yes/no (hit/miss for caches) and gives associated value (if there is one)

# Operations on CAMs

---

## ❑ **Search:** the primary way to access a CAM


- Send data to CAM memory
- Return “found” or “not found”
- If found, return location of where it was found or associated value

## ❑ **Write:**

- Send data for CAM to remember
  - Where should it be stored if CAM is full?
    - Replacement policy
      - Replace oldest data in the CAM
      - Replace least recently searched data

## CAM example

---




**5 storage element CAM array of 9 bits each**

## Previous use of CAMs

---

- ☐ You have seen a simple CAM used before. When?

# Fetch Stage with Branch Prediction



# Cache Organization

---

addr	data
addr	data

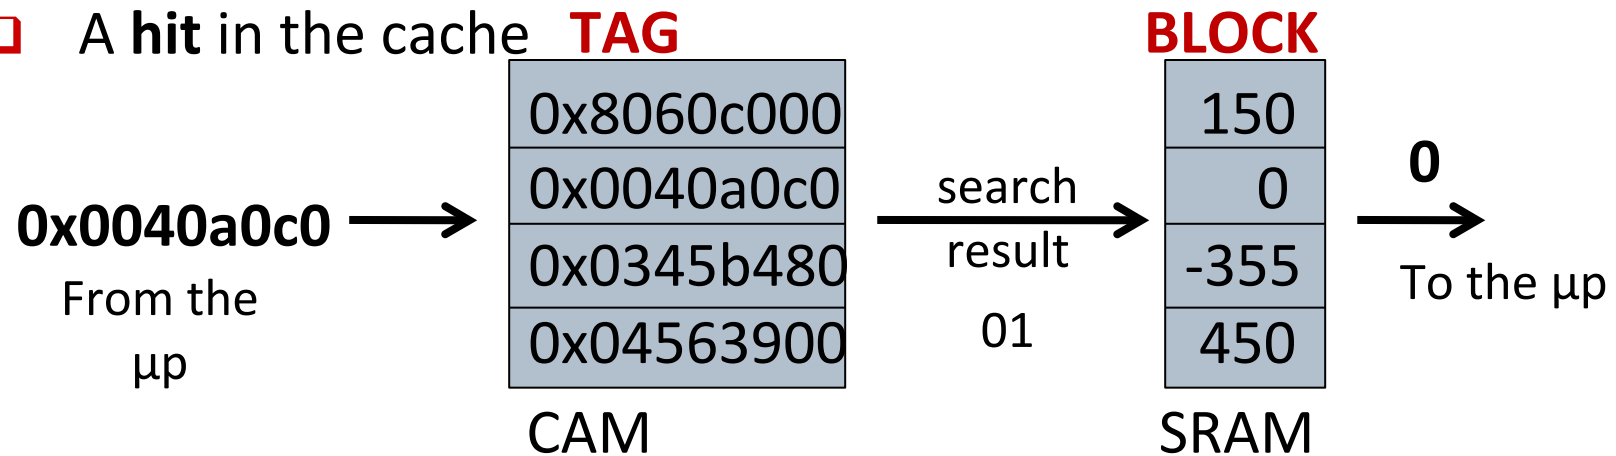
**TAG    BLOCK**

- ❑ A cache memory consists of multiple tag/block pairs (called **cache lines**)
  - Searches can be done in parallel (within reason)
  - At most one tag will match
- ❑ If there is a tag match, it is a cache **HIT**
- ❑ If there is no tag match, it is a cache **MISS**

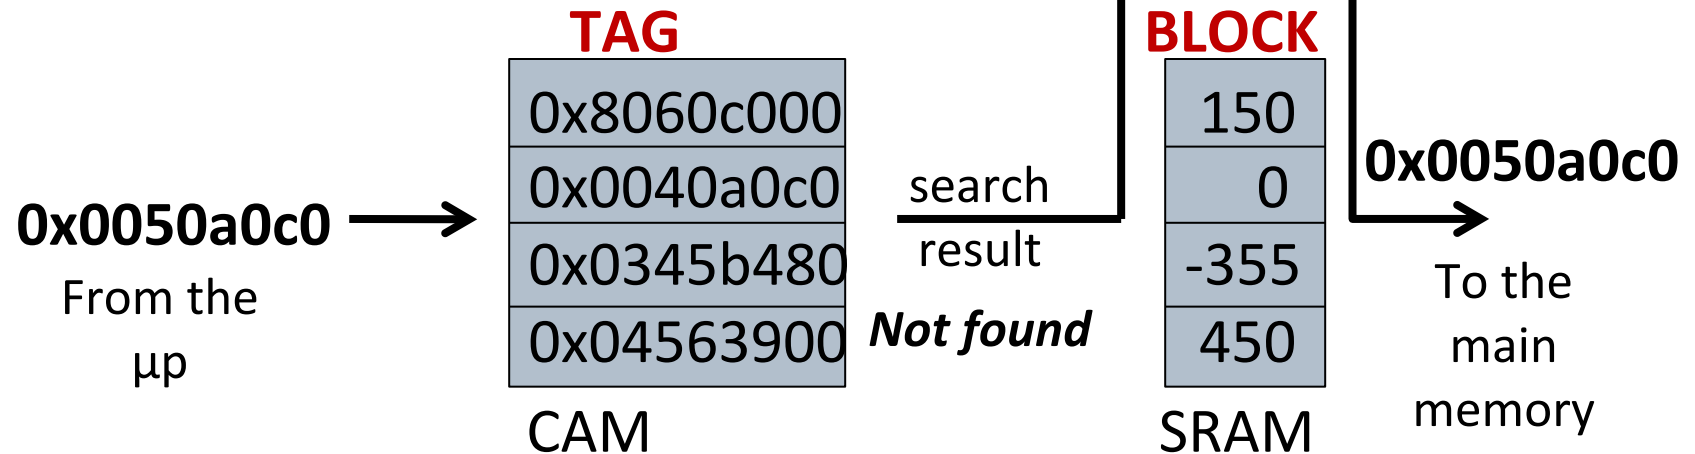
Our goal is to keep the data we think will be accessed in the near future in the cache

# Caches: the hardware view


- ❑ A **hit** in the cache



- ❑ A **miss** in the cache



# CAM = content addressable memory



When used in caches, all tags are fully specified (no X)



# Cache Operation

---

- ❑ Every cache **miss** will get the data from memory and **ALLOCATE** a cache line to put the data in
  - Just like any CAM write
- ❑ Which line should be allocated?
  - Random? OK, but hard to grade test questions
  - Better than random? How?

# Something To Think About

---

- ❑ Does an optimal replacement policy exist?
  - That is, given a choice of cache lines to replace, which one will result in the fewest total misses during program execution
  - Hint: a crystal ball will come in handy in solving this problem...
- ❑ Why would we care?

# Picking the Most Likely Addresses

---

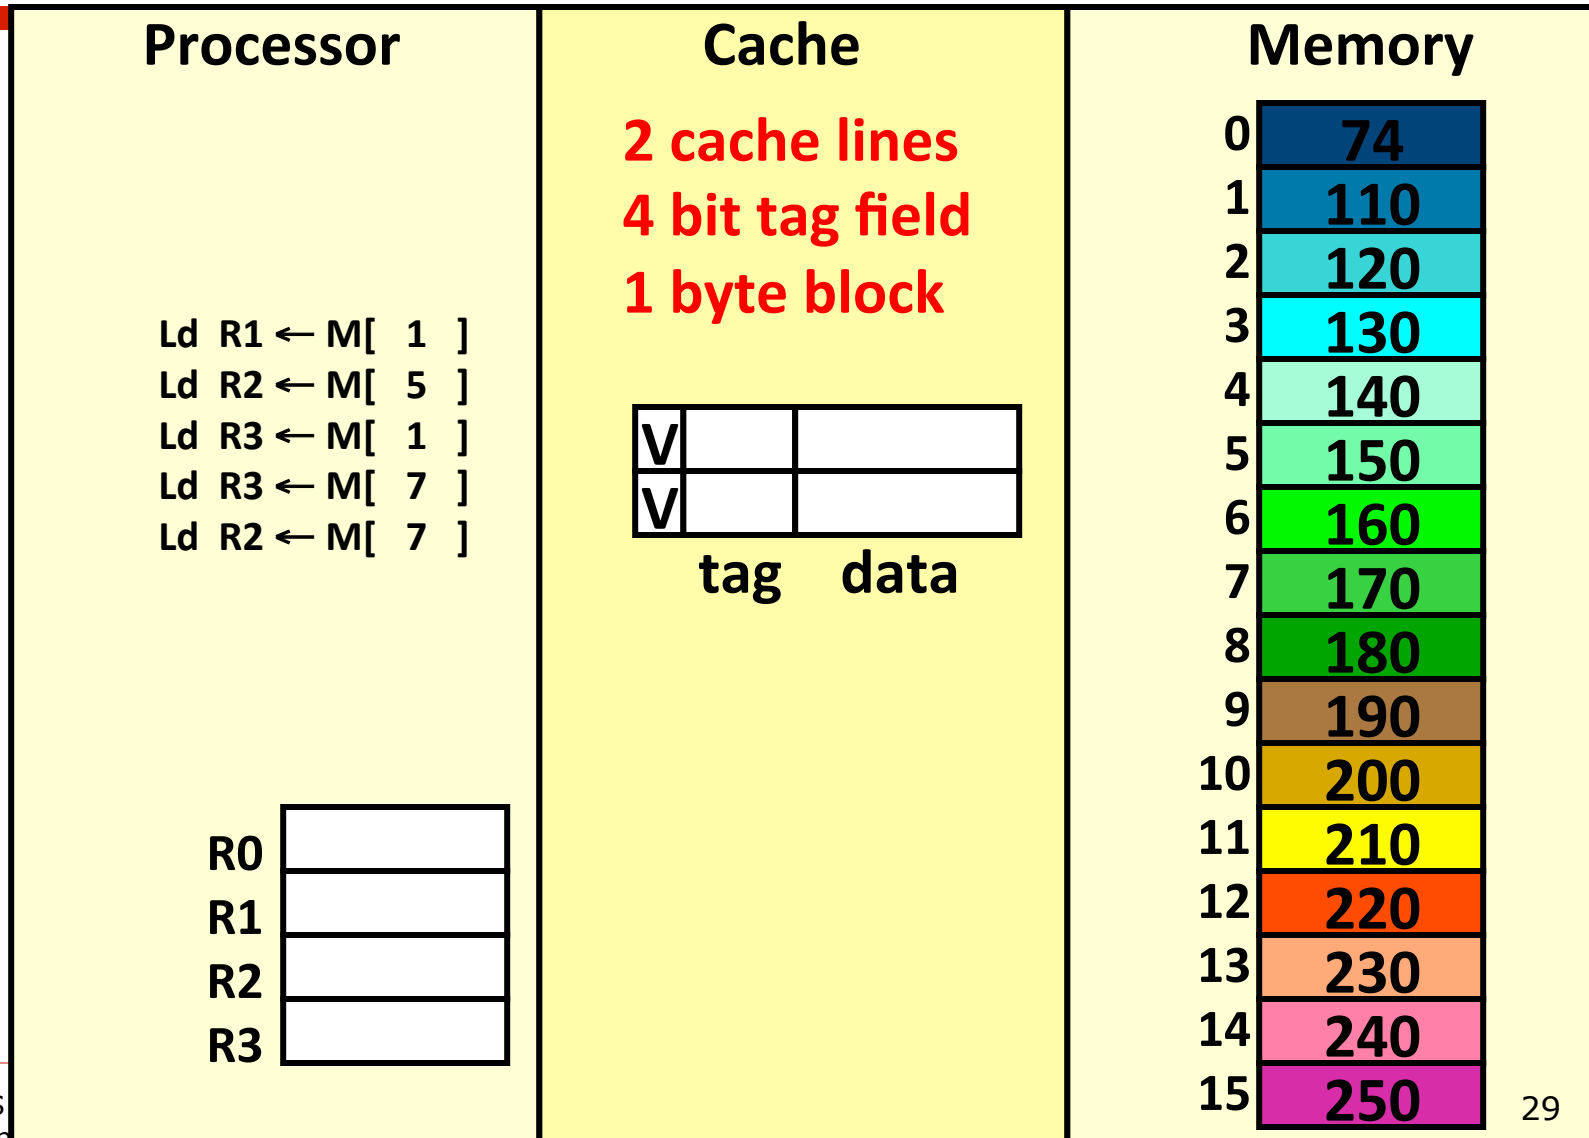
- ❑ What is the probability of accessing a random memory location?
  - With no information, it is just as likely as any other address
- ❑ But programs are not random
  - They tend to use the same memory locations over and over
  - We can use this to pick the most referenced locations to put into the cache

# Temporal Locality

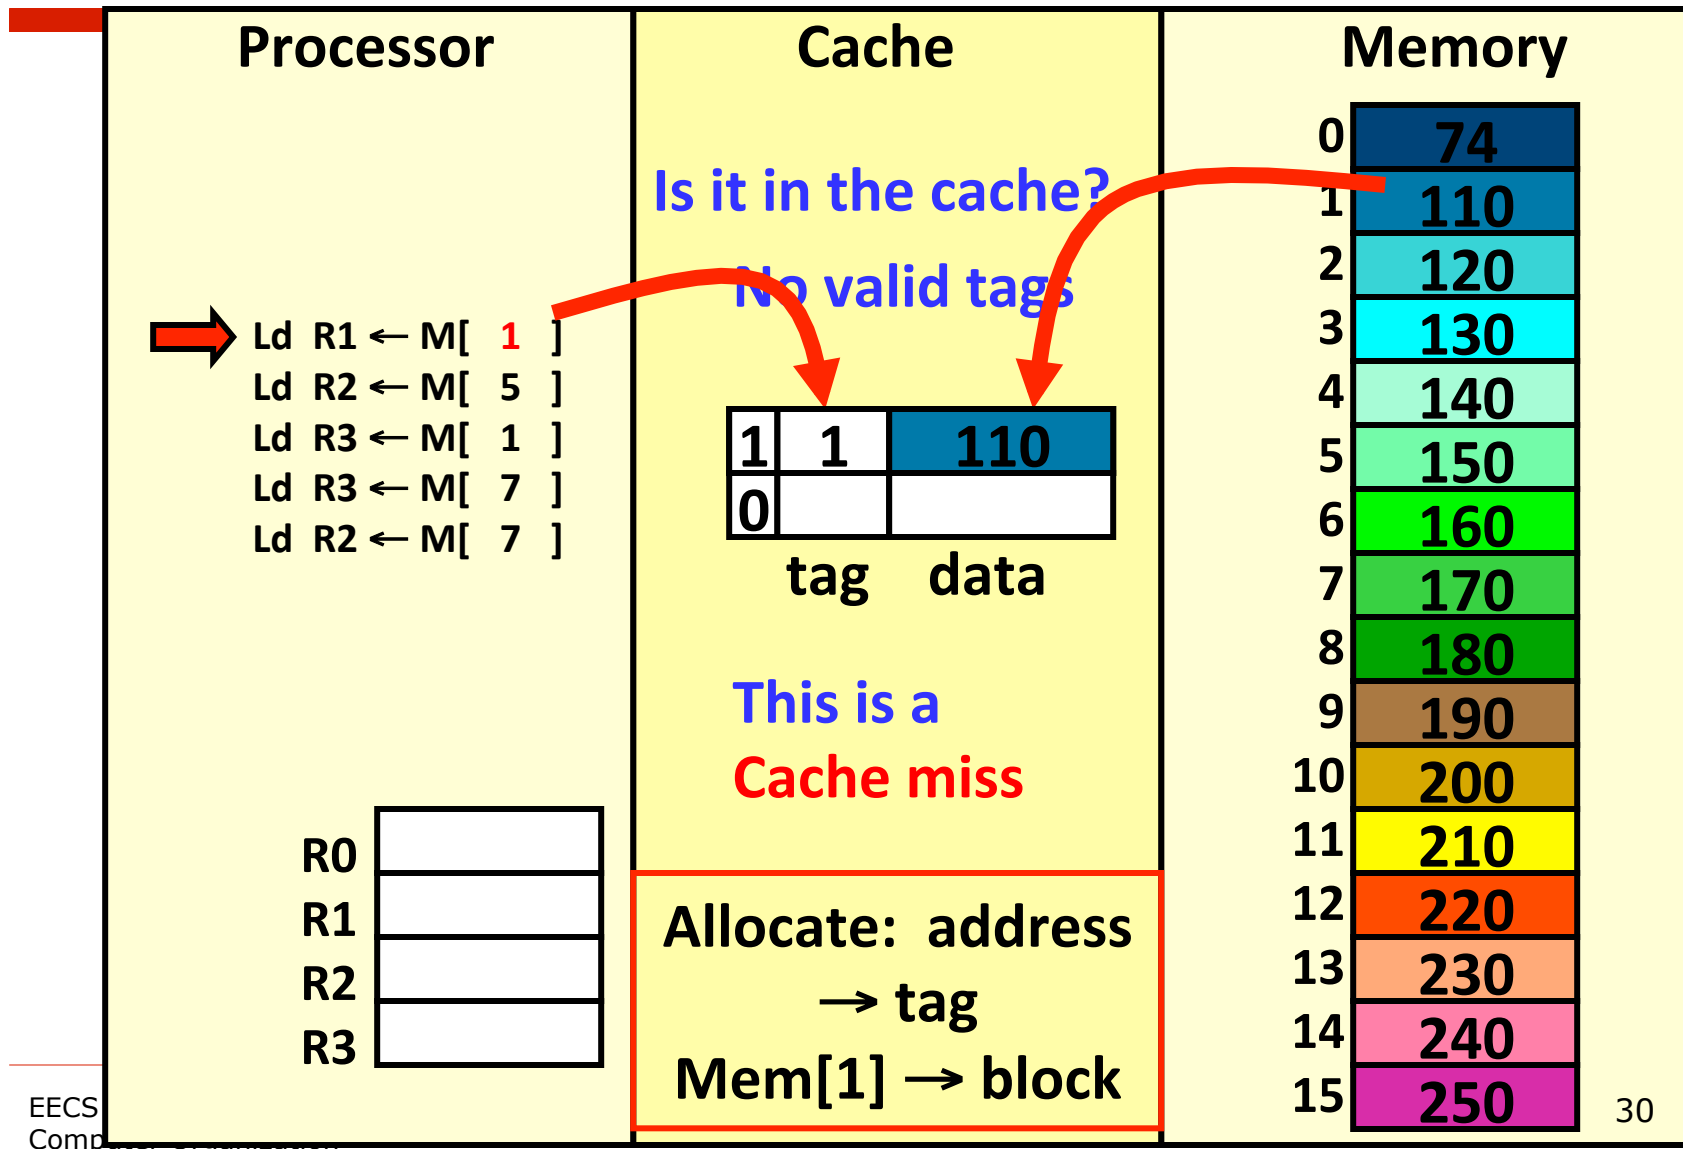
---

- ❑ The principle of **temporal locality** in program references says that if you access a memory location (e.g., 1000) you will be more likely to re-access that location than you will be to reference some other random location
- ❑ Temporal locality says any miss data should be placed into the cache
  - It is the most recent reference location
- ❑ Temporal locality says that the least recently referenced (or least recently used – **LRU** ) cache line should be **evicted** to make room for the new line
  - Because the re-access probability falls over time as a cache line isn't referenced, the LRU line is least likely to be re-referenced

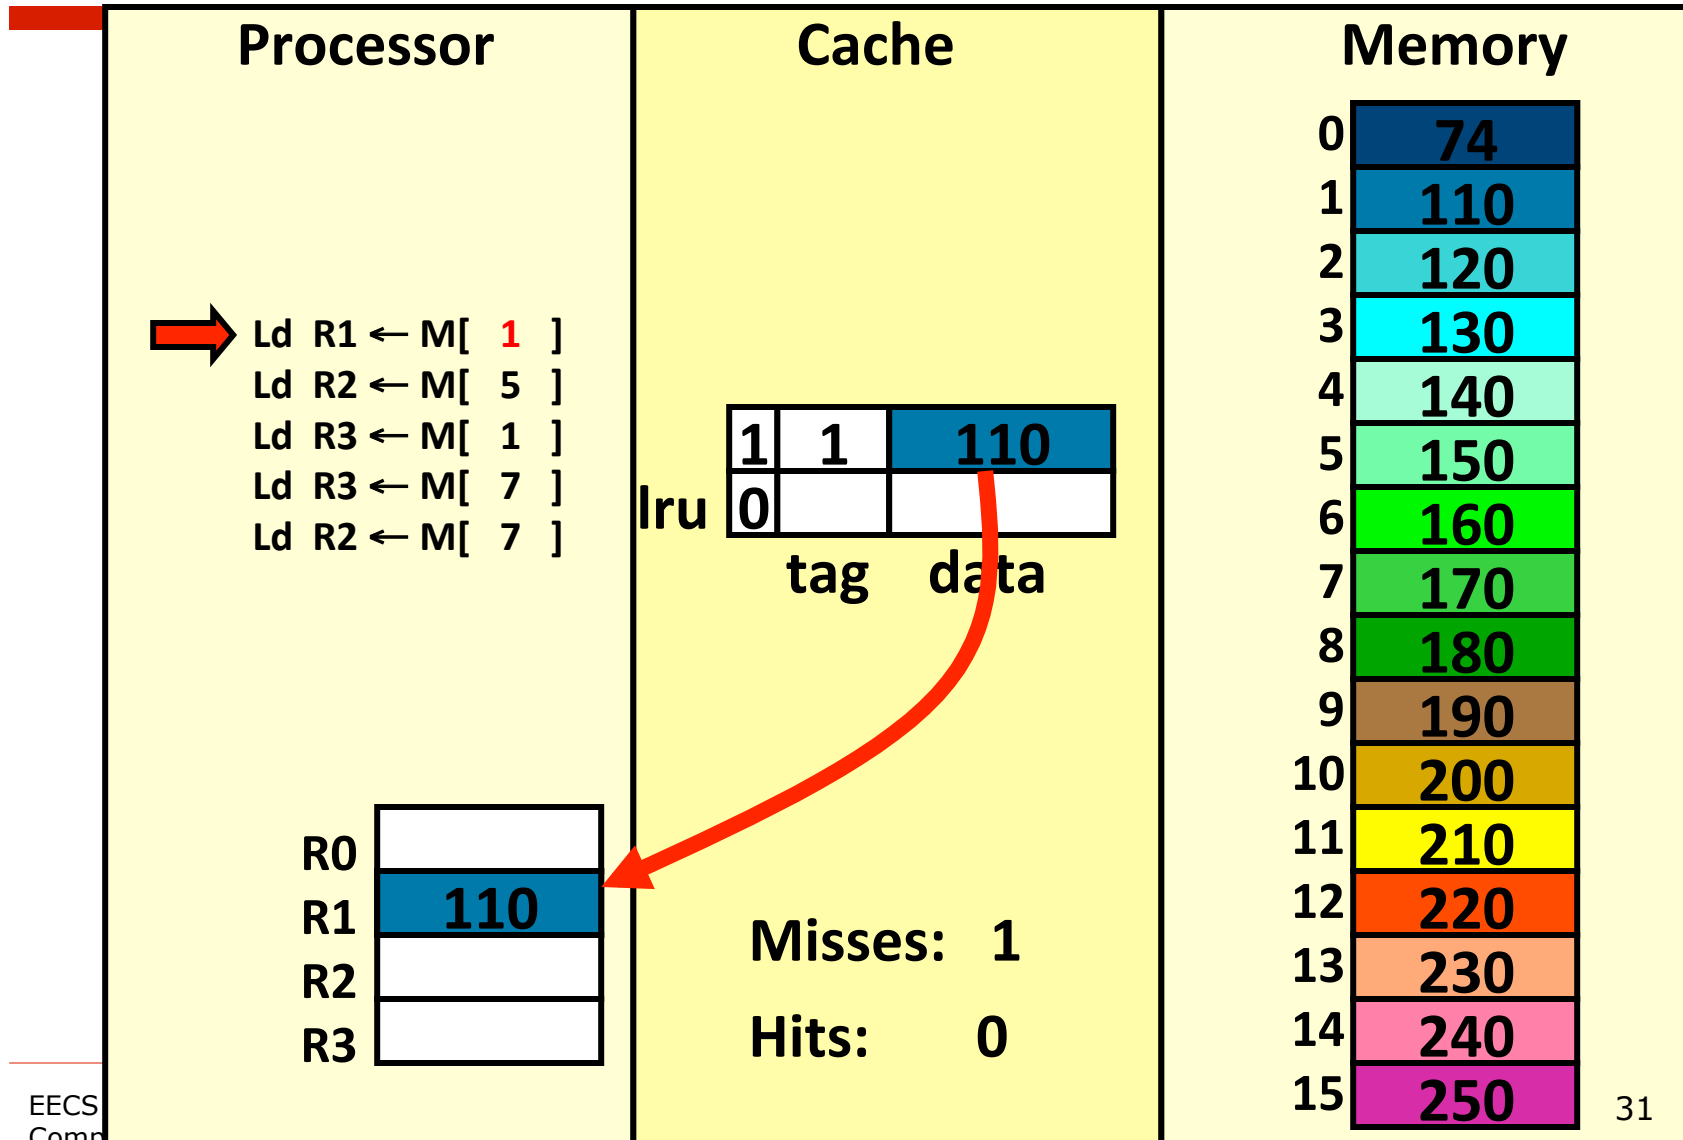
# A Very Simple Memory System



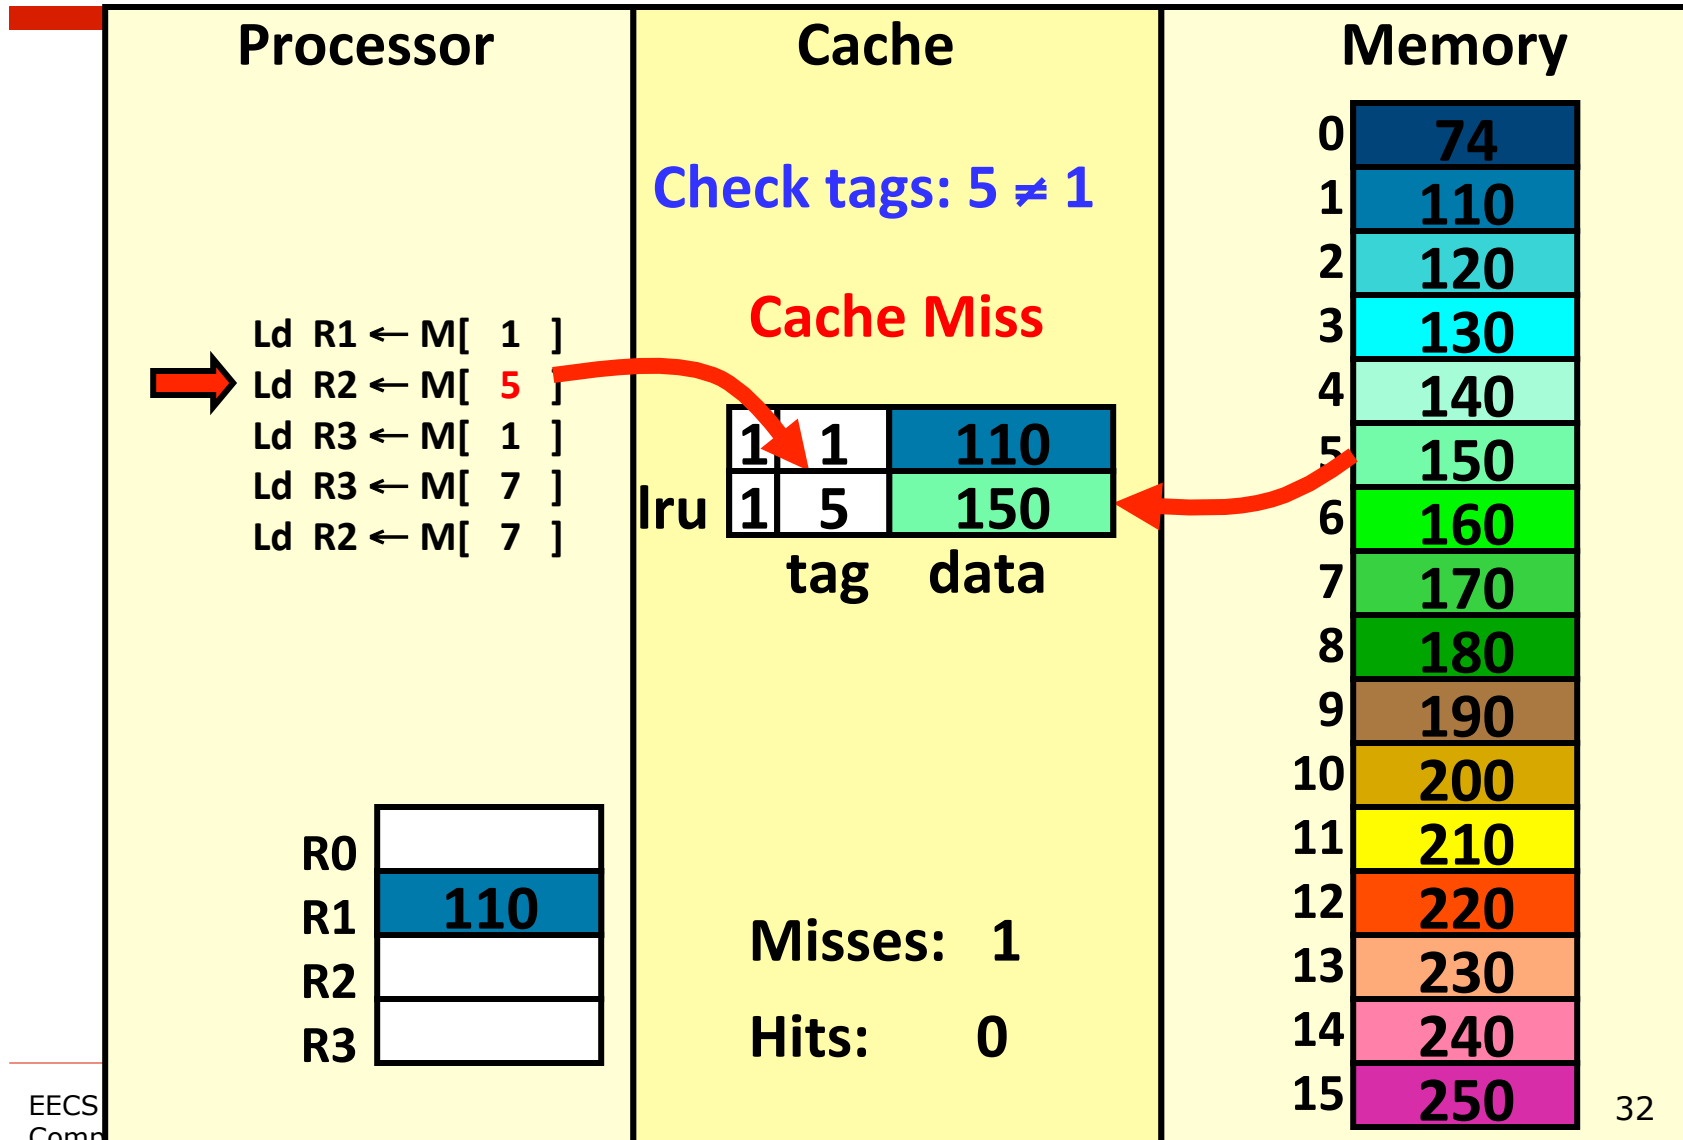
# A Very Simple Memory System



# A Very Simple Memory System

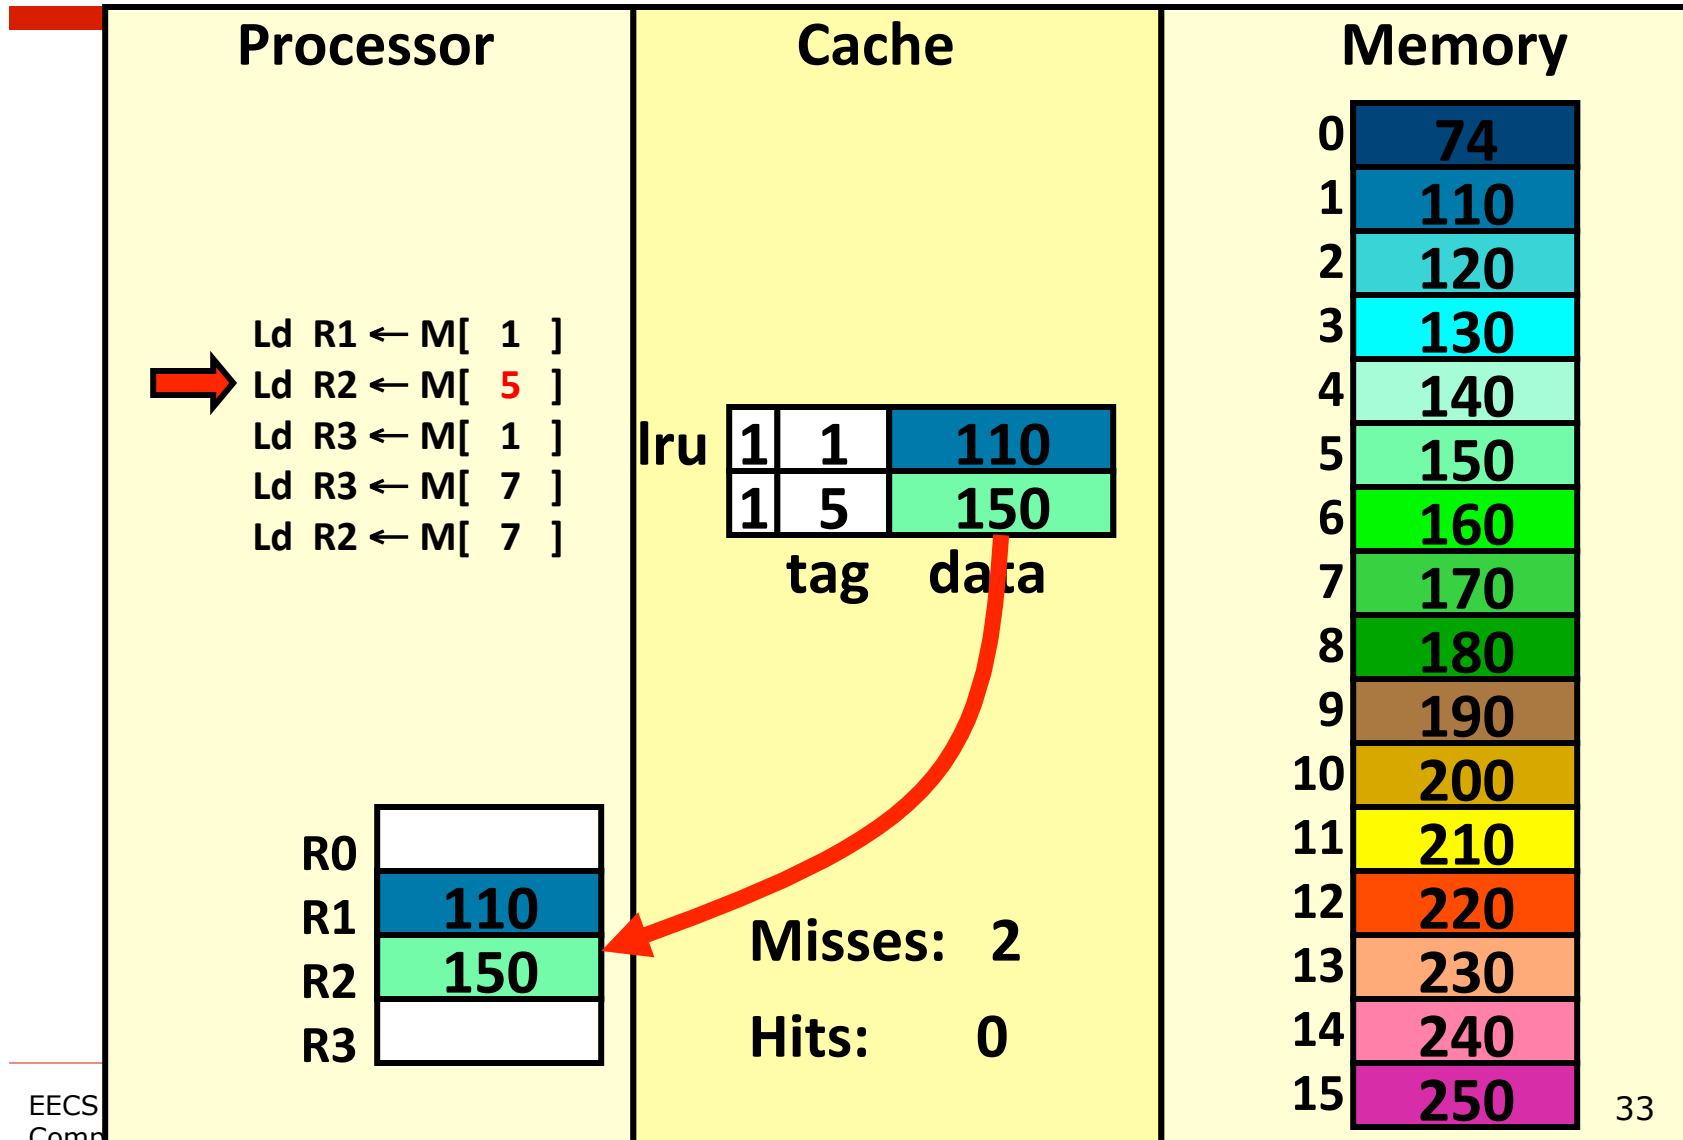


# A Very Simple Memory System

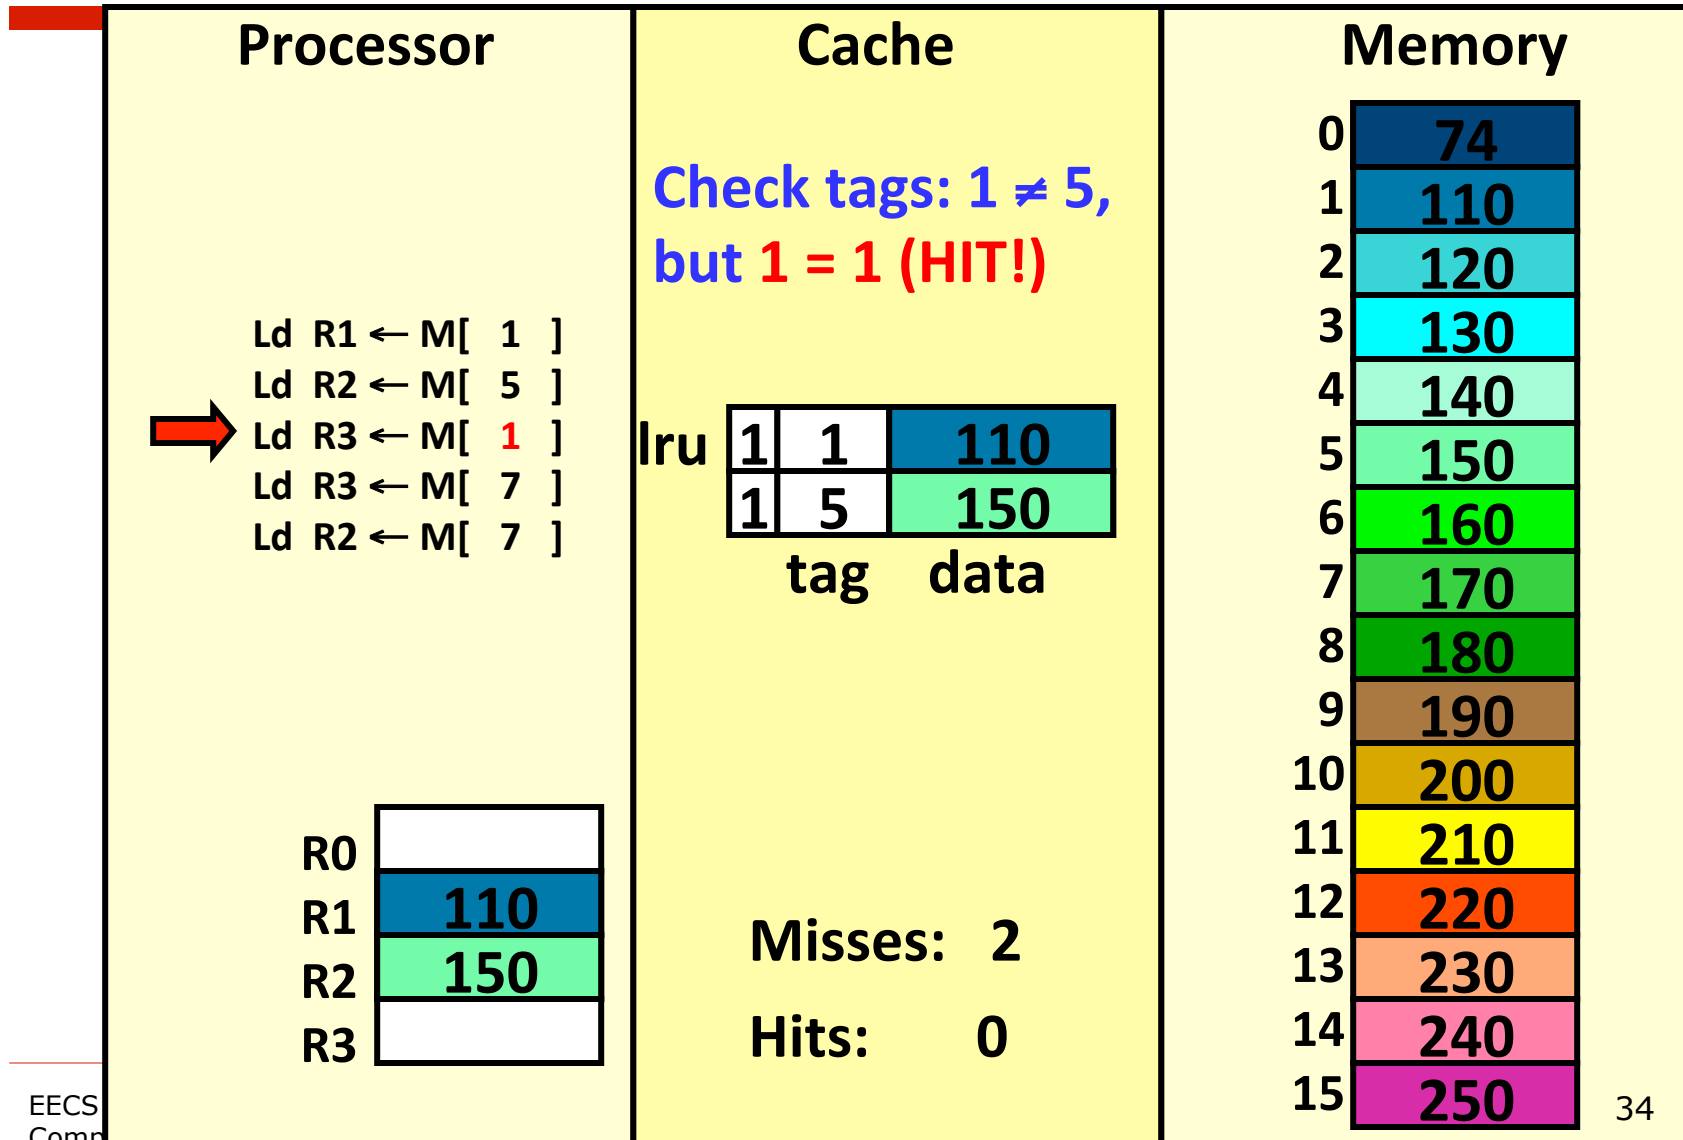




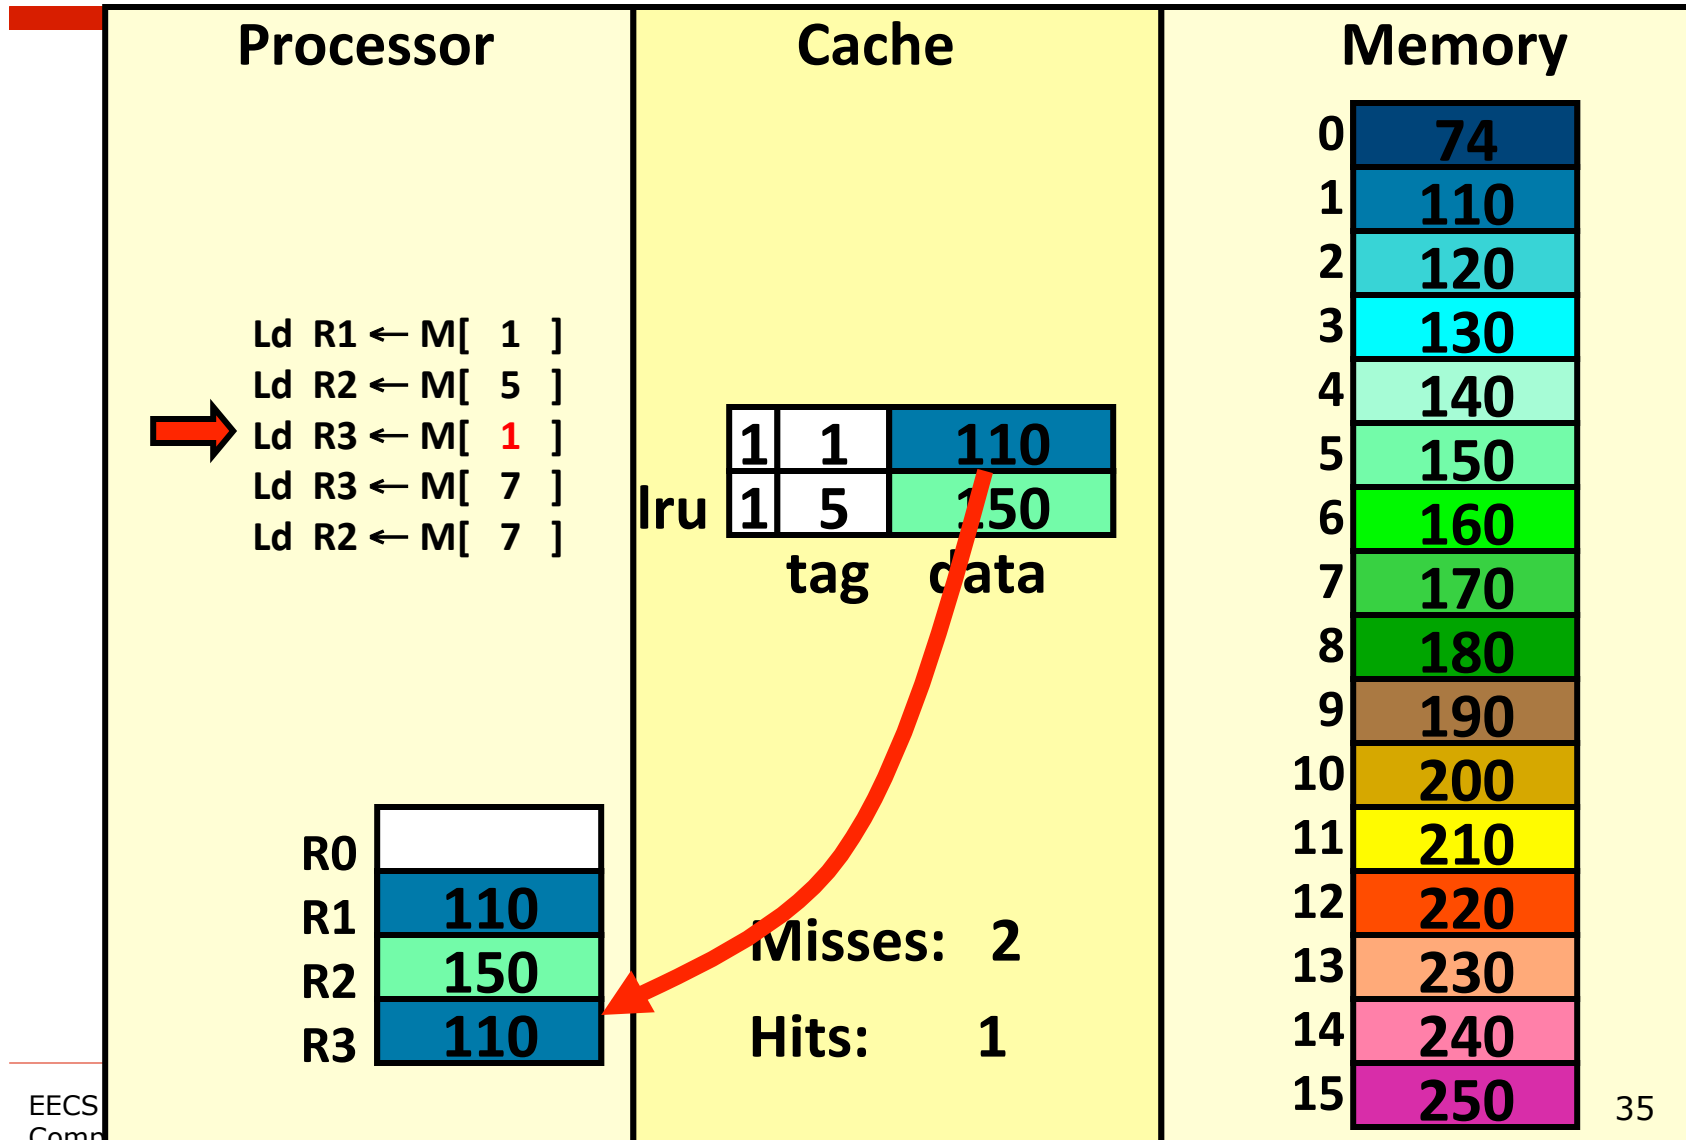
# A Very Simple Memory System



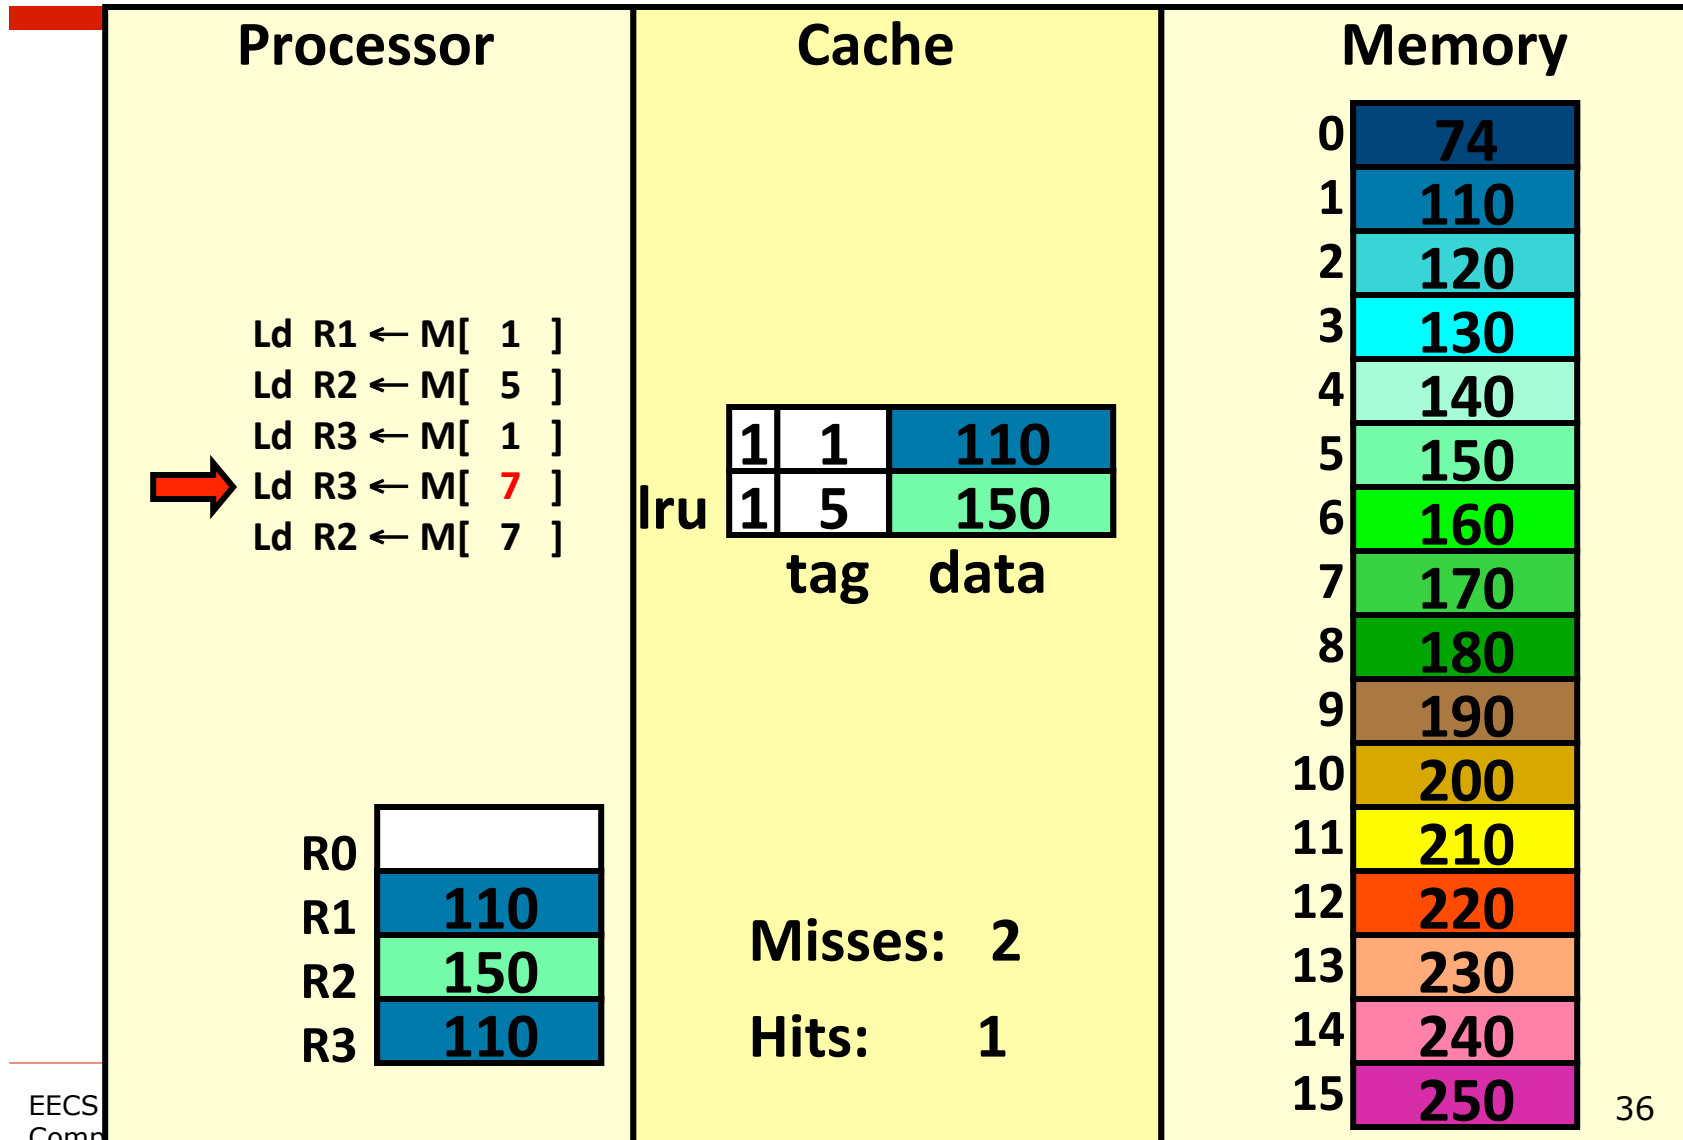
# A Very Simple Memory System



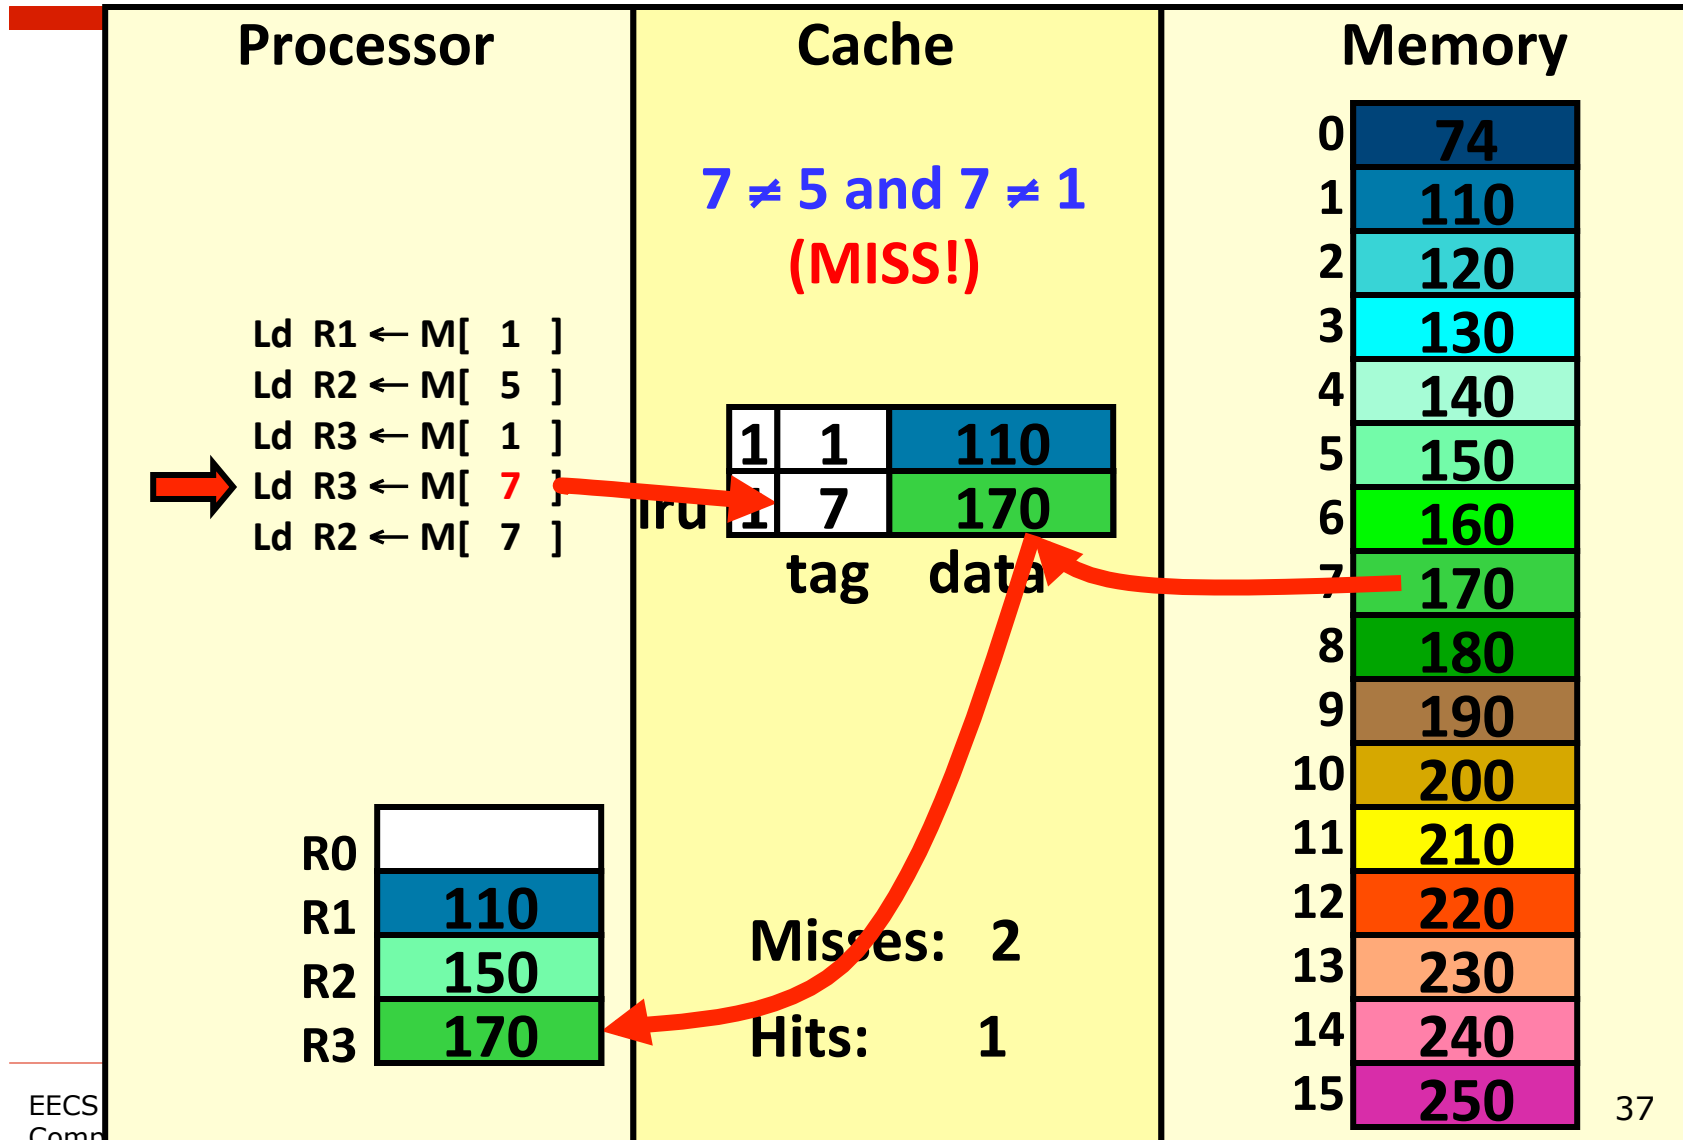
# A Very Simple Memory System



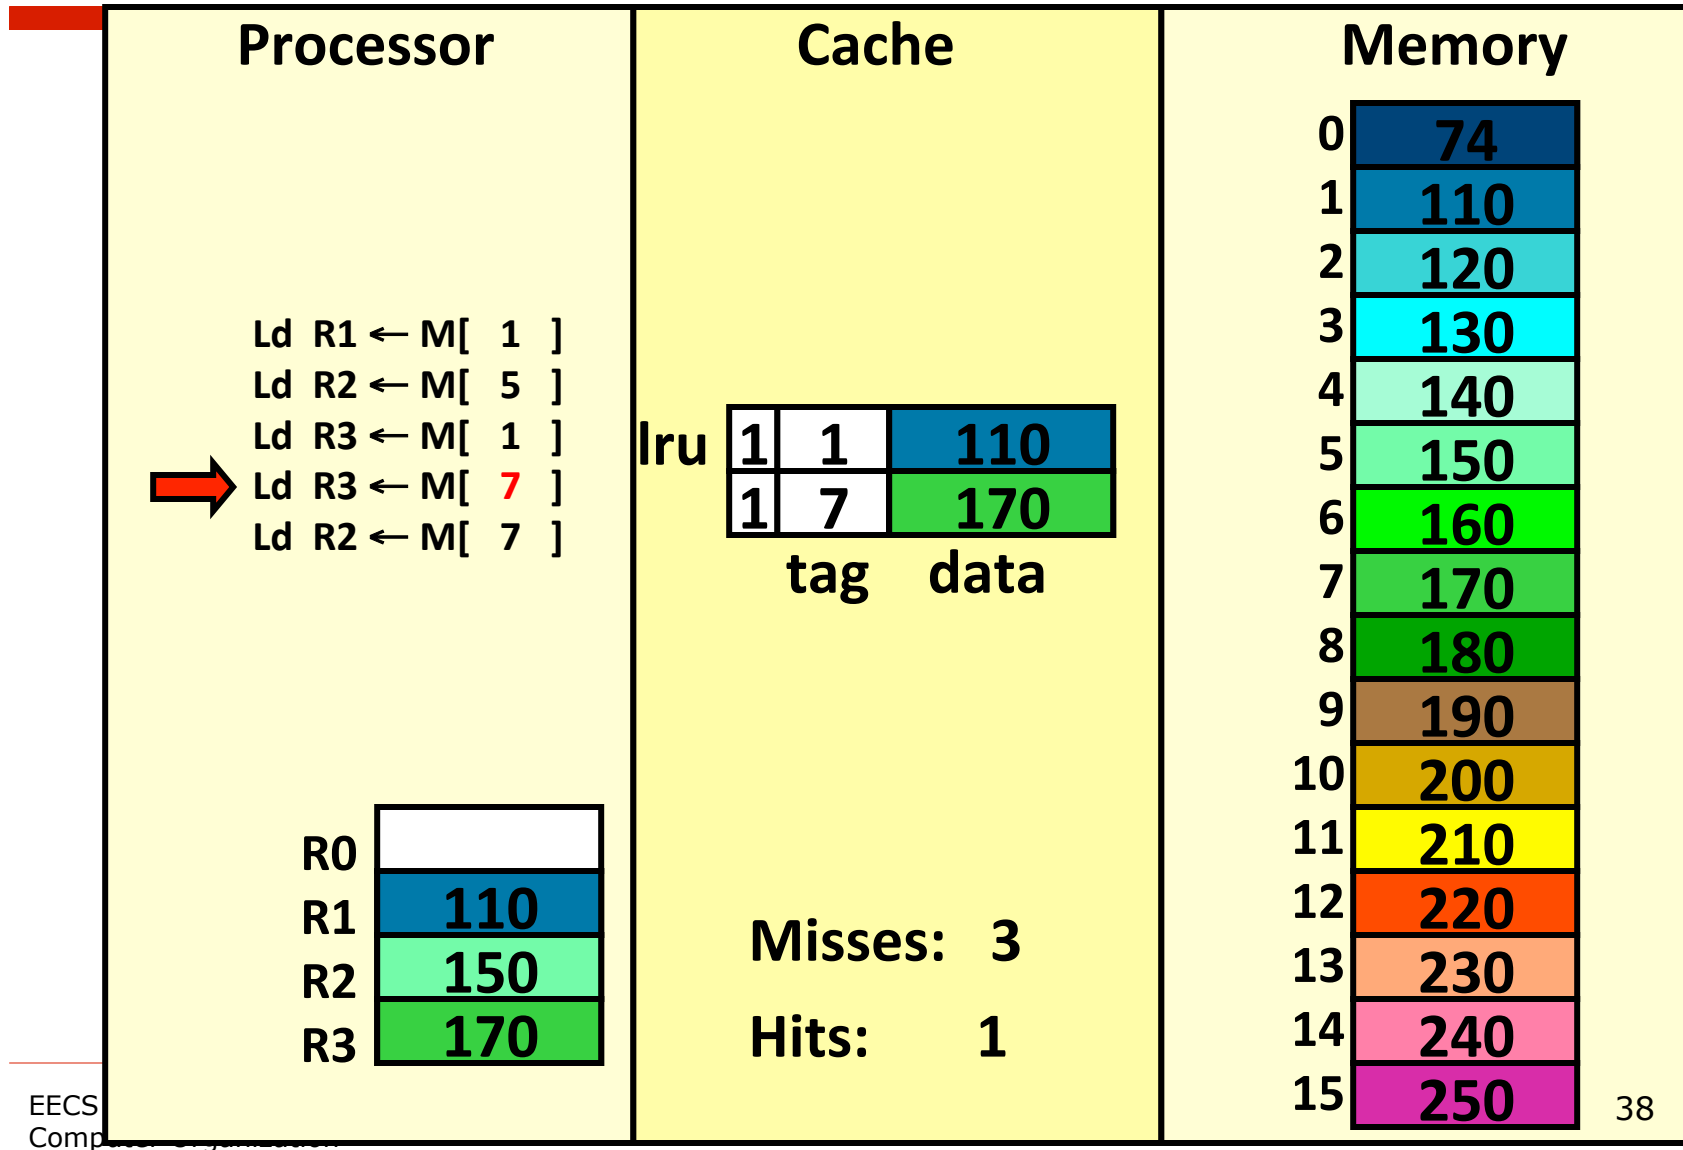
# A Very Simple Memory System



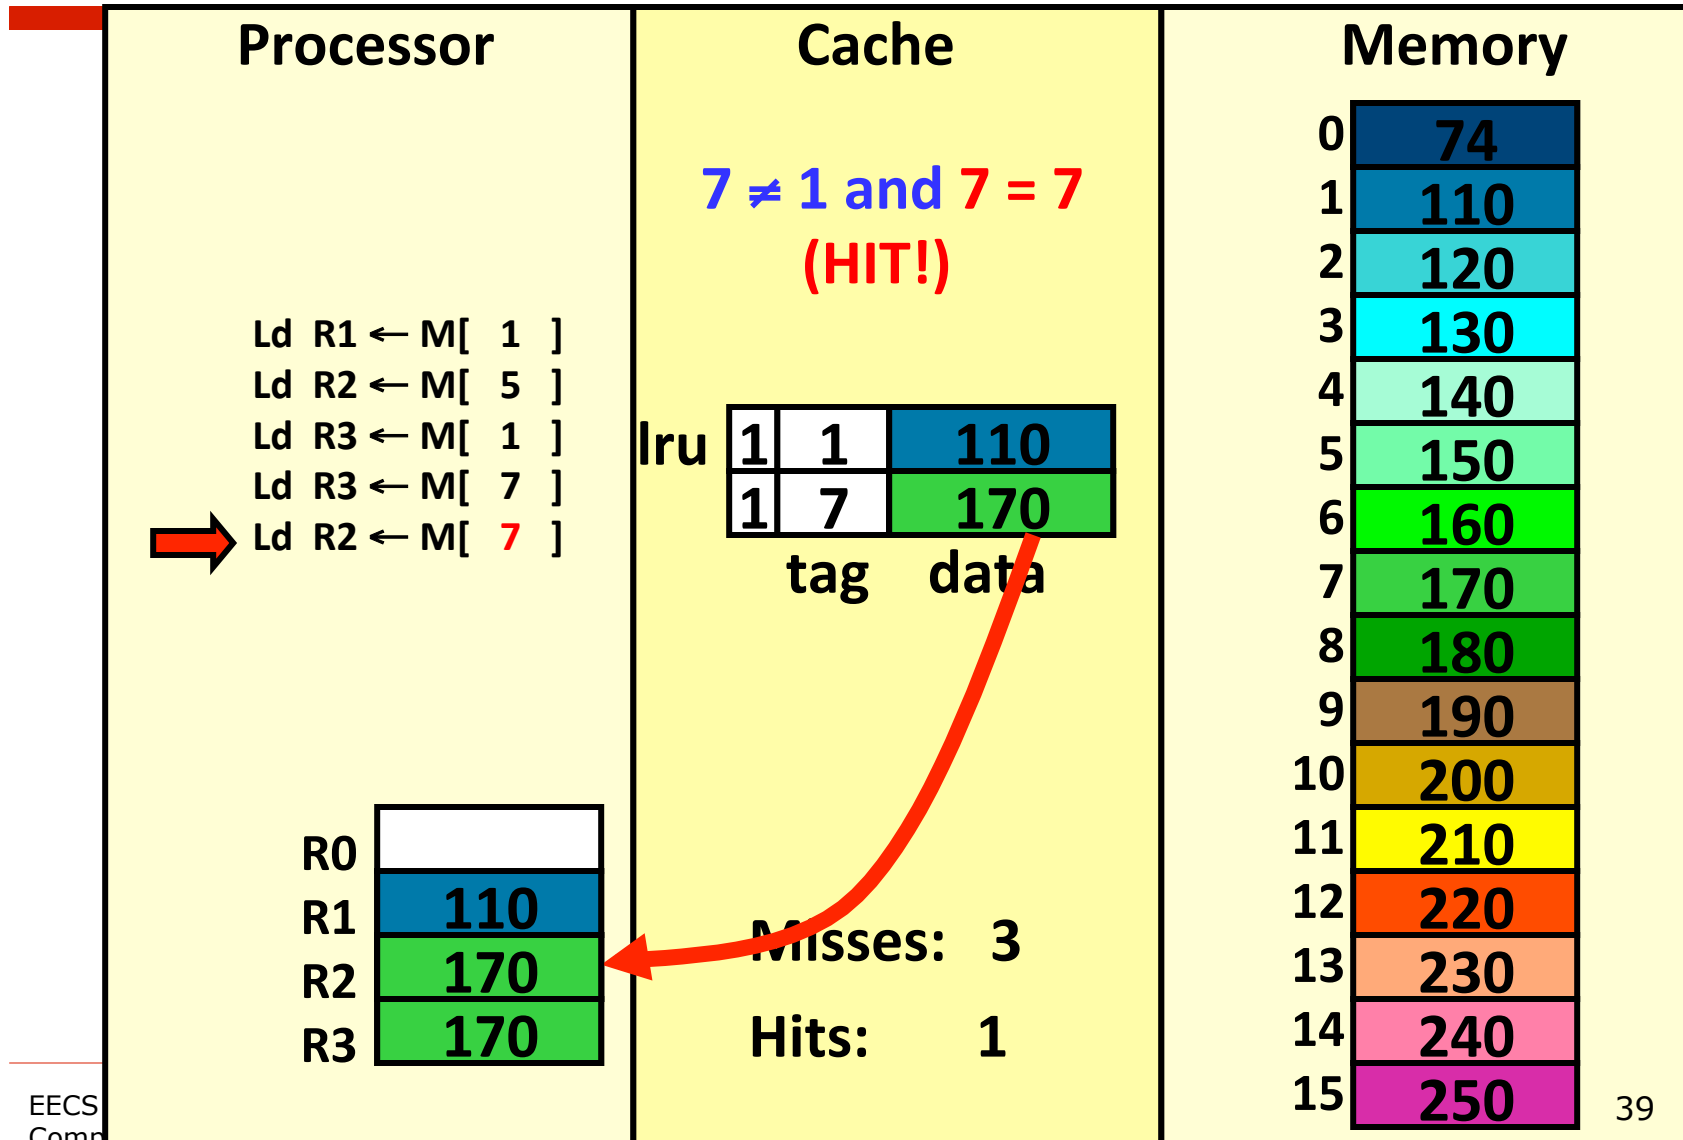
# A Very Simple Memory System



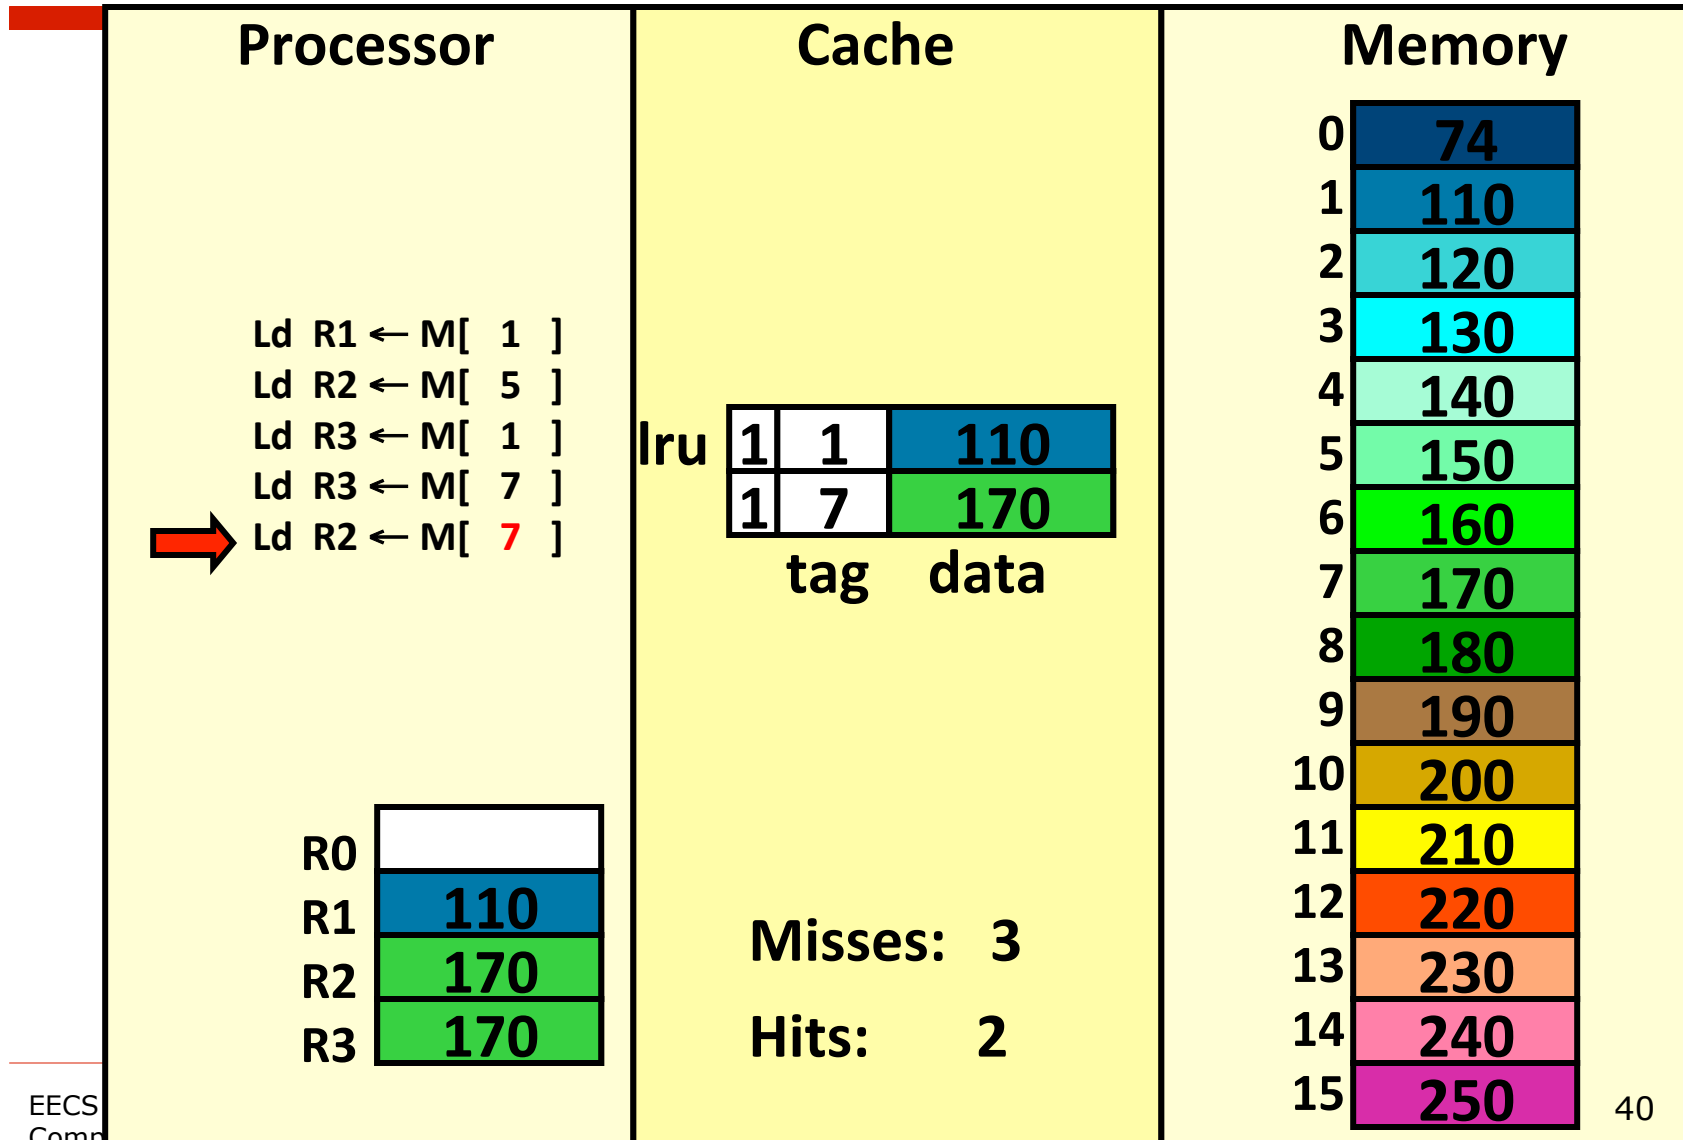
# A Very Simple Memory System



# A Very Simple Memory System



# A Very Simple Memory System





# Calculating Average Access Latency

---

- ❑ **Avg latency**  
= cache latency  $\times$  hit rate + memory latency  $\times$  miss rate
  
- ❑ Avg latency for our example cache  
= 1 cycle  $\times$  (2/5) + 15  $\times$  (3/5)  
= 9.4 cycles per reference
  
- ❑ To improve average latency:
  - Improve memory access latency, or
  - Improve cache access latency, or
  - Improve cache hit rate

# Calculating Cost

---

- ❑ How much does our example cache cost (in bits)?
  - Calculate storage requirements
    - 2 bytes of SRAM
  - Calculate overhead to support access (tags)
    - 2 4-bit tags
    - The cost of the tags is often forgotten for caches, but this cost drives the design of real caches
    - 2 valid bits
- ❑ What is the cost if a 32 bit address is used?

# How can we reduce the overhead?

---

- ❑ Have a small address.
  - Impractical, and caches are supposed to be micro-architectural
  
- ❑ Cache bigger units than bytes
  - Each block has a single tag, and blocks can be whatever size we choose.
  
- ❑ To Be Continued...