

# 21. Putting it all together: Virtual memory and Caches

---

EECS 370 – Introduction to Computer Organization - Winter 2016

**Profs. Valeria Bertacco & Reetu Das**

EECS Department  
University of Michigan in Ann Arbor, USA

© Bertacco-Das, 2016

The material in this presentation cannot be  
copied in any form without our written permission


# Announcements

---

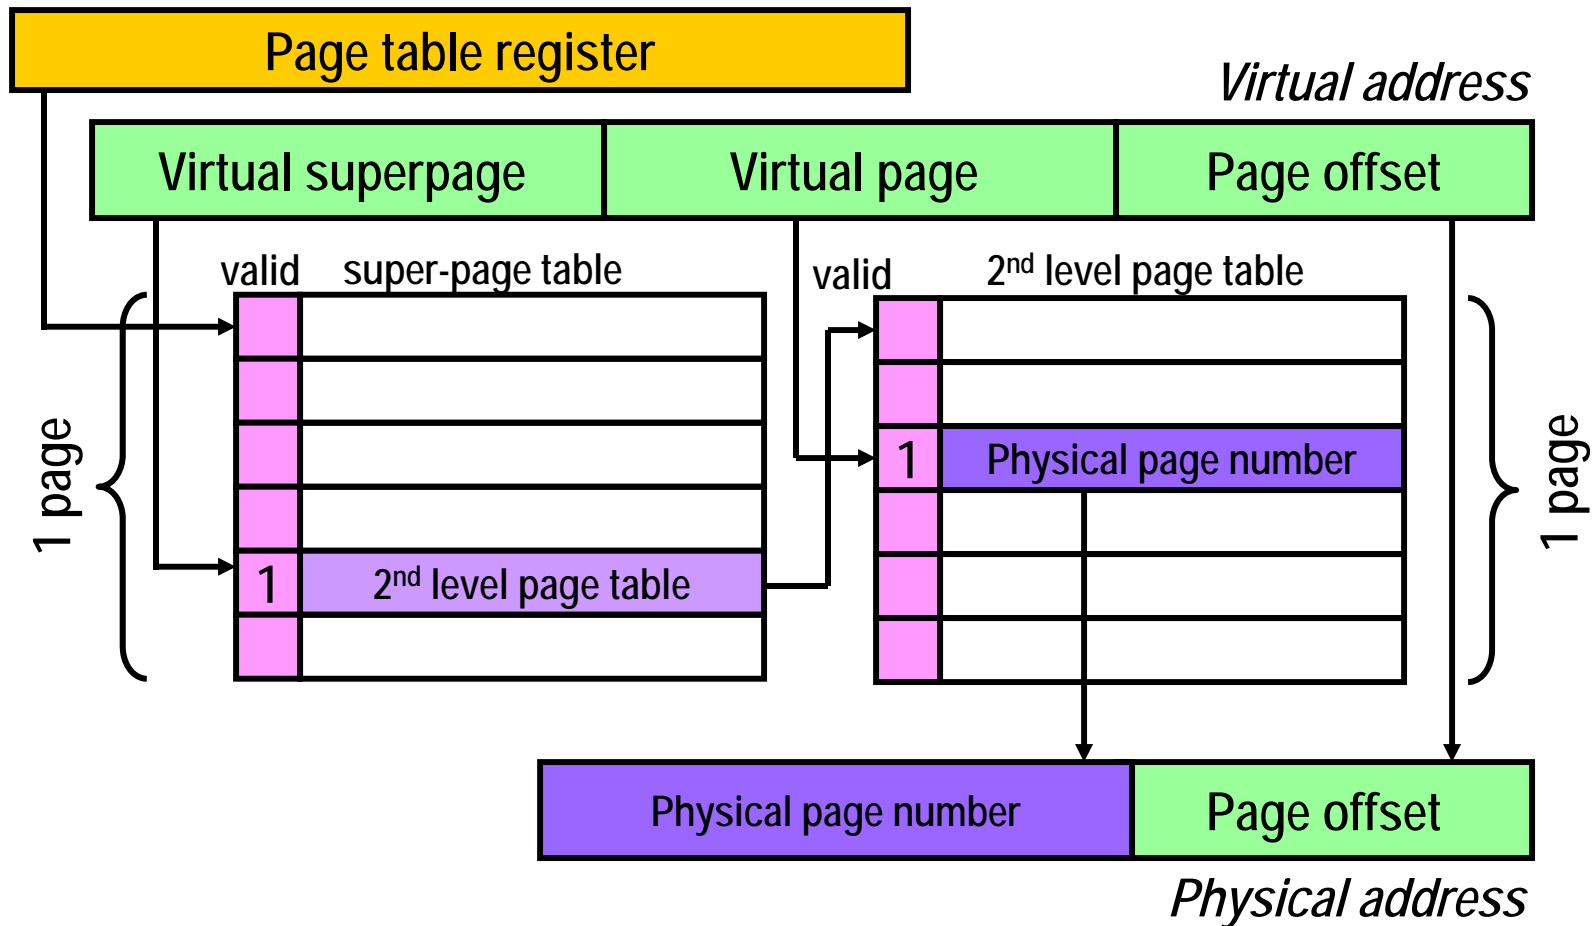
- ❑ Project 4 – due today
- ❑ Homework 7 – due Tuesday
  
- ❑ FINAL EXAM preview:
  - Cumulative
  - Cheat sheet – 3 sides (2 sheets)
  - 120 minutes long (longer than midterms!)
  - Make sure to not get sick
  
- ❑ Midterm2 grades will be republished this afternoon

# REVIEW: key pieces in the life of VM

The address translation information of the program is contained in the *Page Table*



# REVIEW: Hierarchical page tables




# REVIEW: Translation Look-aside Buffer

---


The TLB is a **CAM** = content-addressable memory

Virtual page	Pg offset
--------------	-----------



All tags are compared concurrently against the virtual page #. If anyone hits, we gather the physical page #. Otherwise, we go through the usual process (consult page table).

v	tag	Physical page



# Class Problem 1 – VM architecture

---

A new generation U-tanium processor has a 32-bit virtual address space and it can support up to 512 MB of byte addressable RAM. The machine has single level page tables and the page size is set to 4KB. Each entry of the page table contains an additional byte with control information (dirty, valid, etc.).

- ❑ What is the size of one page table entry in bytes? Remember that page table entries are byte aligned
- ❑ How many pages does the page table occupy?
- ❑ How many programs can run simultaneously on this machine? Each program has its own page table which is kept in memory at all times. Assume at least 128MB of RAM must be devoted to data (not page tables).
- ❑ How will increasing the page size to 32KB affect the page table? (same size, larger, smaller)

# Caches in systems with VM

---

- ❑ VM systems give us two different addresses: virtual and physical.
- ❑ Which address should we use to access the data cache?
  - Physical address (after VM translations).
    - Delayed access.
  - Virtual address (before VM translation).
    - Faster access.
    - More complex.


# Cache & VM Organization: option 1

---

## Physically-addressed \$

✗ Slower

✓ Low complexity



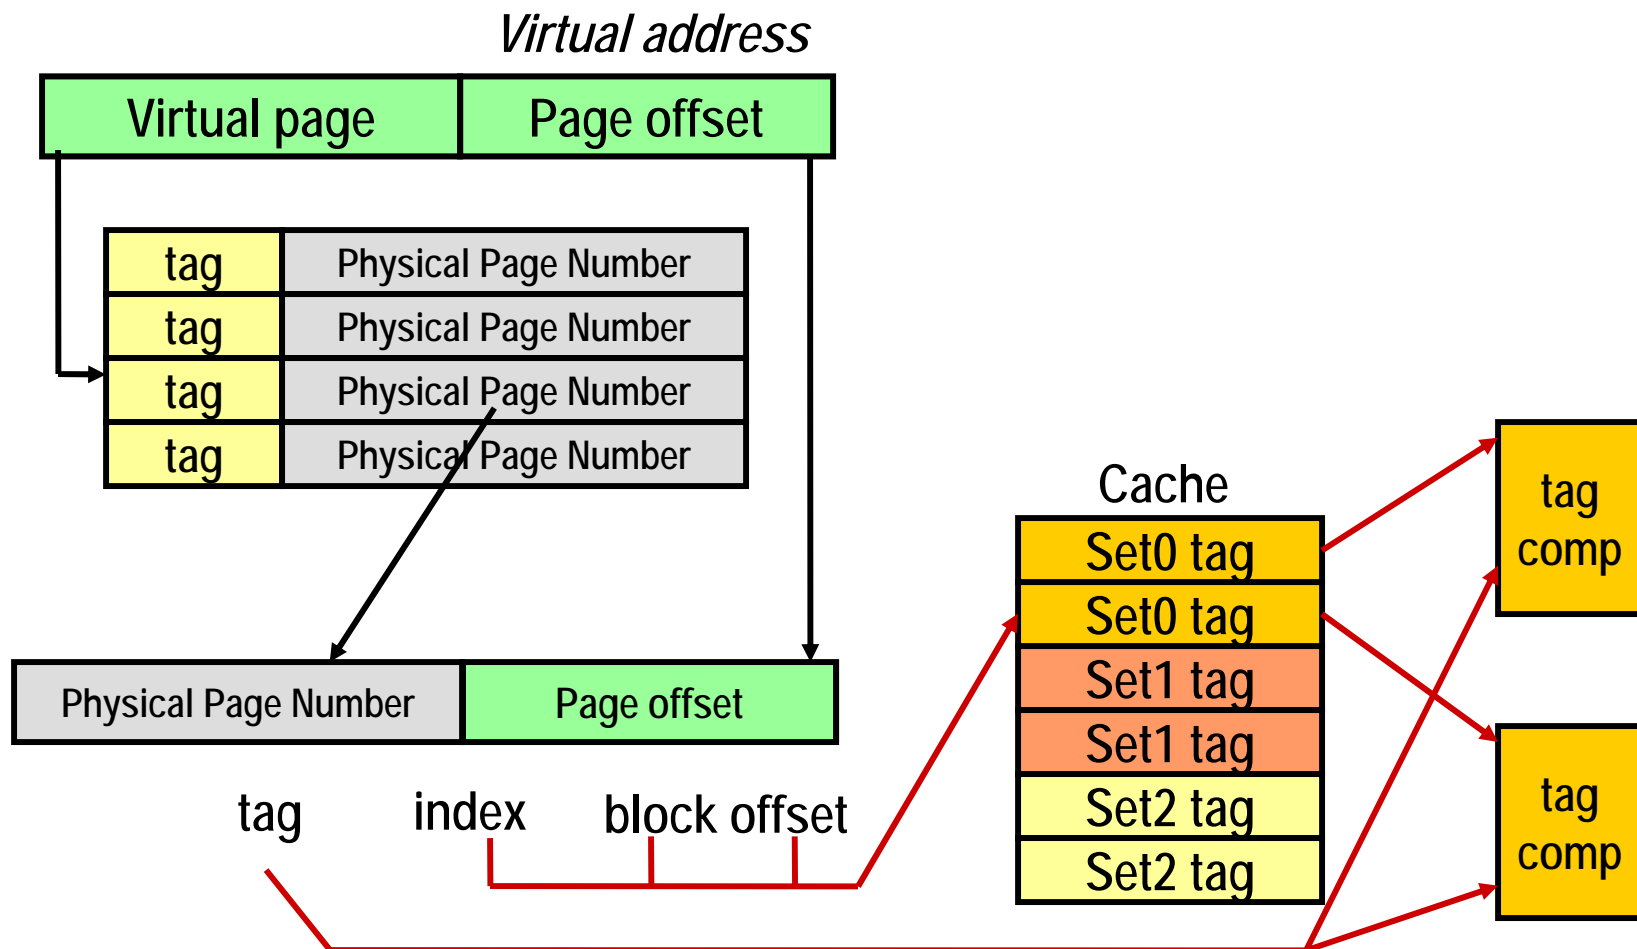


# Physically addressed caches

---

- ❑ Perform TLB lookup *before* cache tag comparison.
  - Use bits from physical address to index set.
  - Use bits from physical address to compare tag.
  
- ❑ Slower access?
  - Tag lookup takes place *after* the TLB lookup.
  
- ❑ Simplifies some VM management.
  - When switching processes, TLB must be invalidated, but cache OK to stay as is.
  - Implications? Might result in fewer cache misses if context switches very common (but they generally are not).

# Physically addressed caches




# Cache & VM Organization: option 2

## Virtually-addressed \$

✗ High complexity (aliasing)

✓ Faster



# Virtually addressed caches

---


- ❑ Perform the TLB lookup at the same time as the cache tag compare.
  - Uses bits from the virtual address to index the cache set
  - Uses bits from the virtual address for tag match.
- ❑ Problems:
  - Aliasing: Two processes may refer to the same physical location with different virtual addresses.
  - When switching processes, TLB must be invalidated, dirty cache blocks must be written back to memory, and cache must be invalidated.

# Multitasking with VM

---

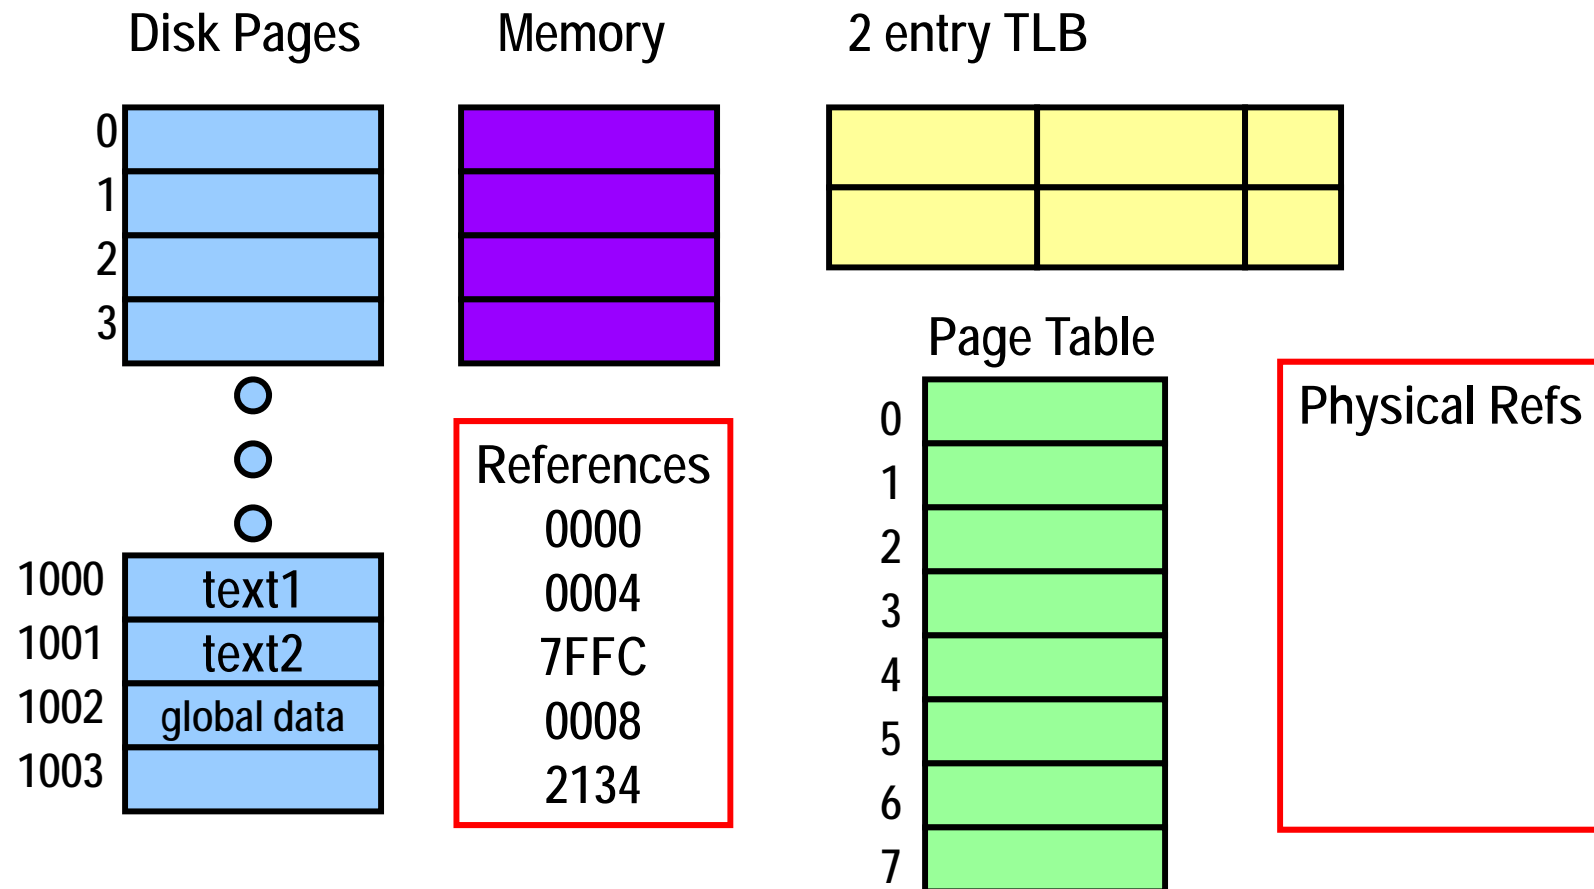
- ❑ If virtually-addressed caches:
  - Flush the cache between each **context switch**.
  - OR
  - Use **processID** (a unique number for each processes given by the operating system) as part of the tag

# Virtually addressed caches



- TLB is accessed in parallel with cache lookup.
- Physical address is used to access main memory in case of a cache miss.

# Loading a program into memory



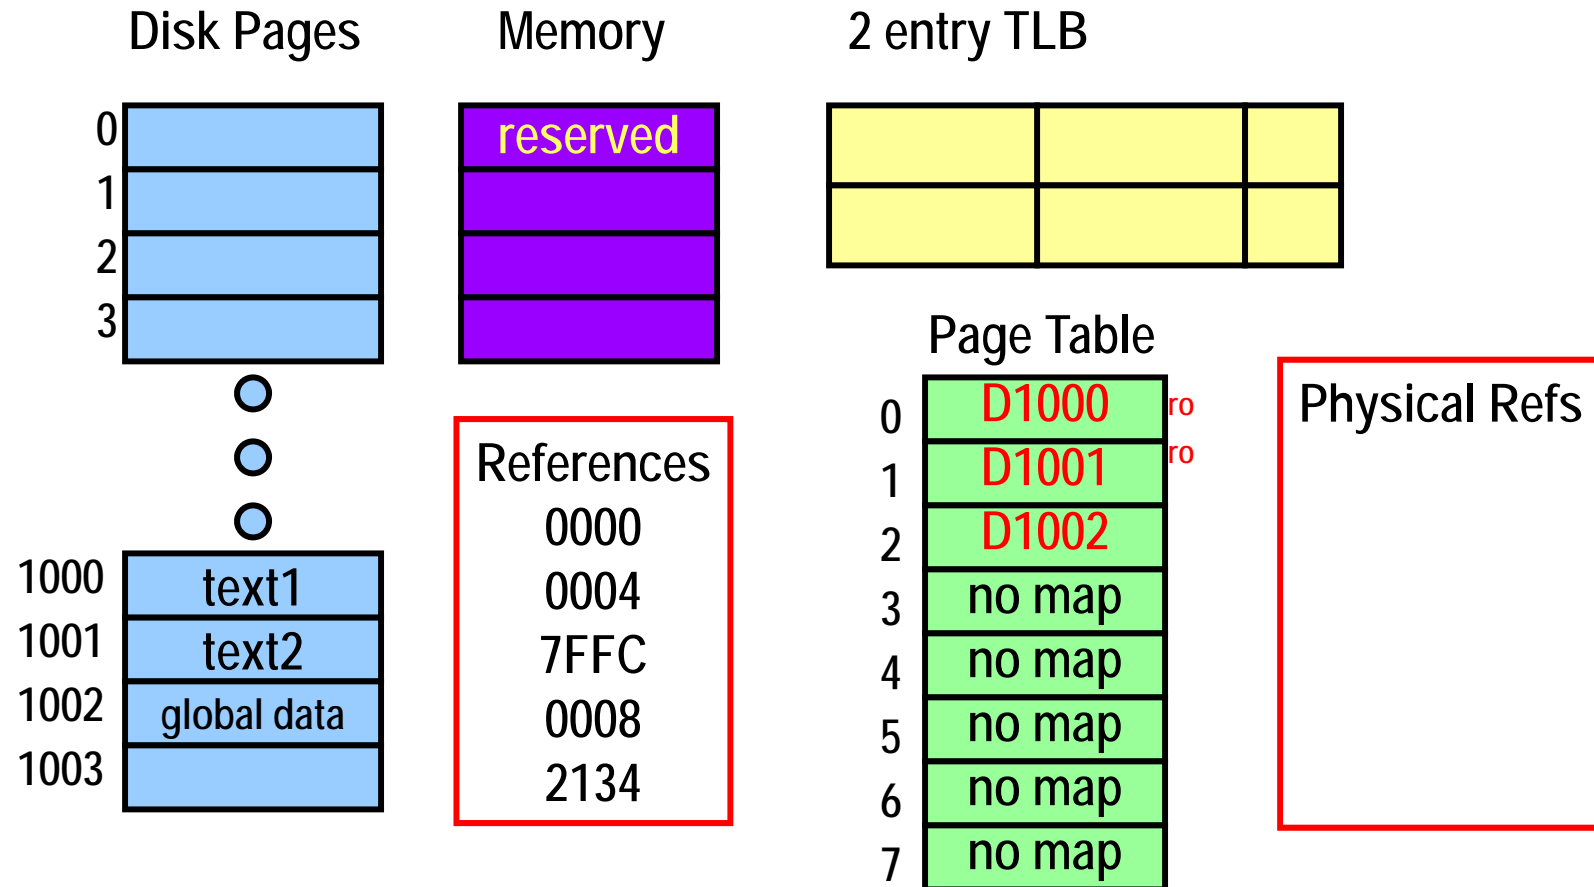
# Additional information

---

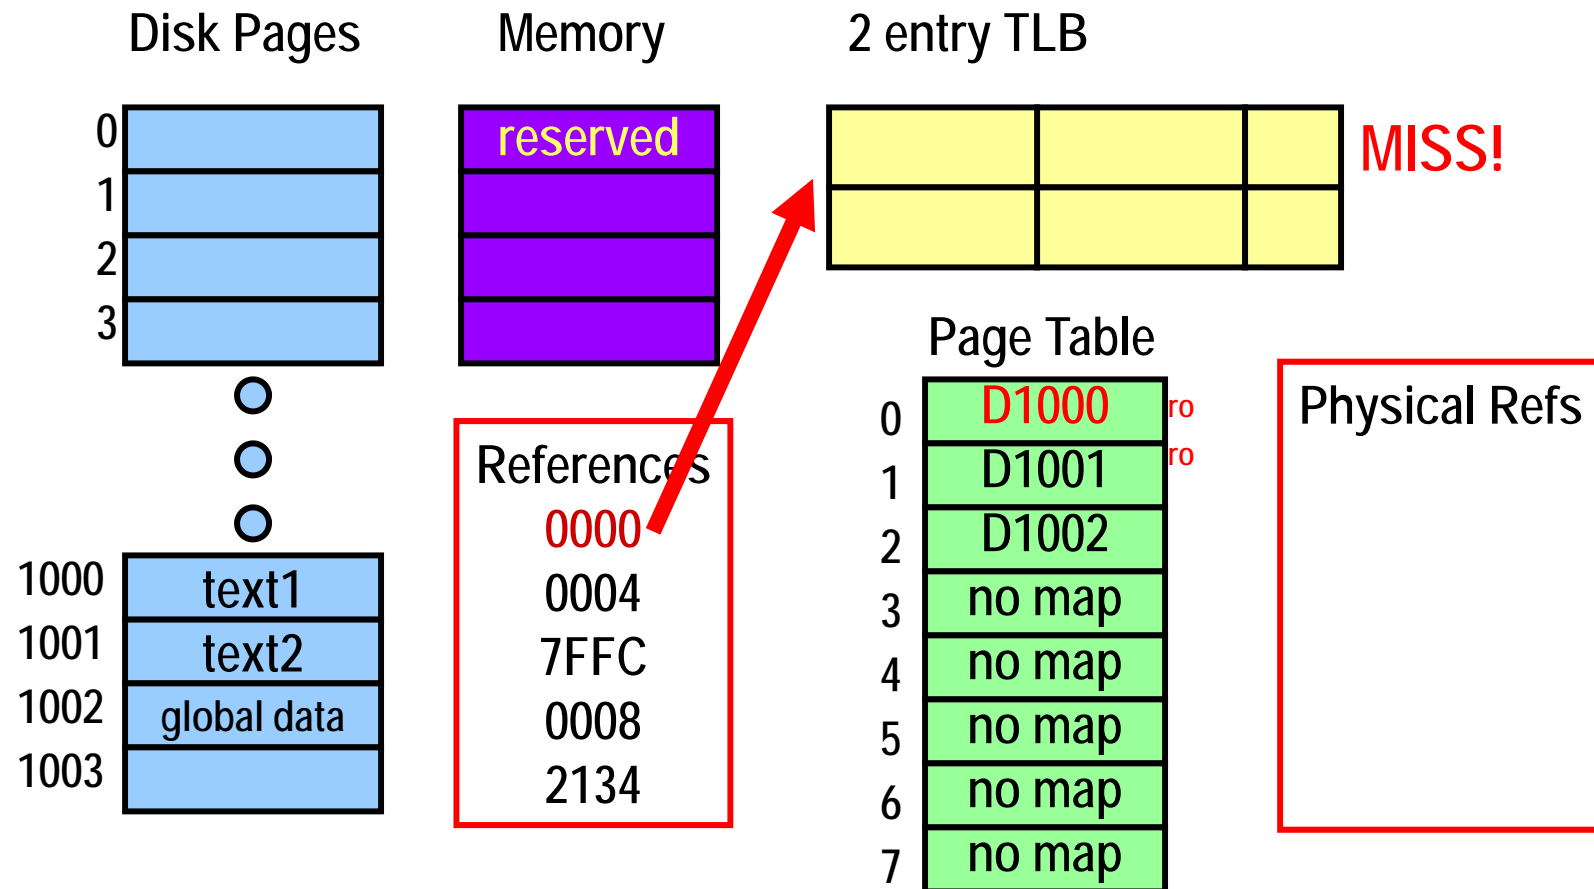
- ❑ Page size = 4 KB.
- ❑ Page table entry size = 4 B.
- ❑ Page table register points to physical address 0000.



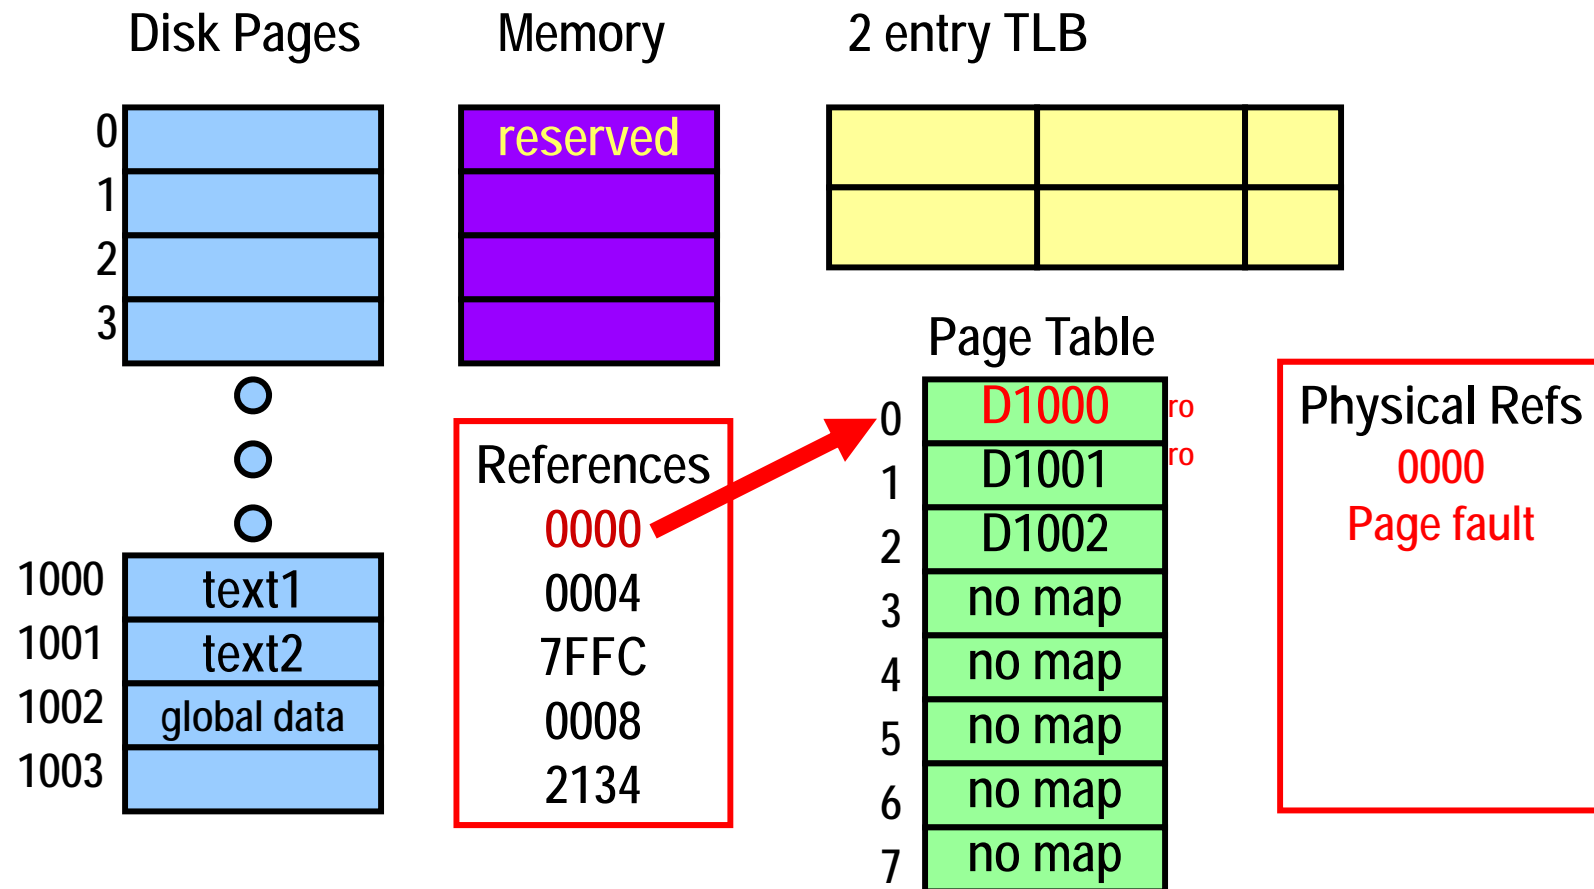
# Step 1: read executable header and initialize page table



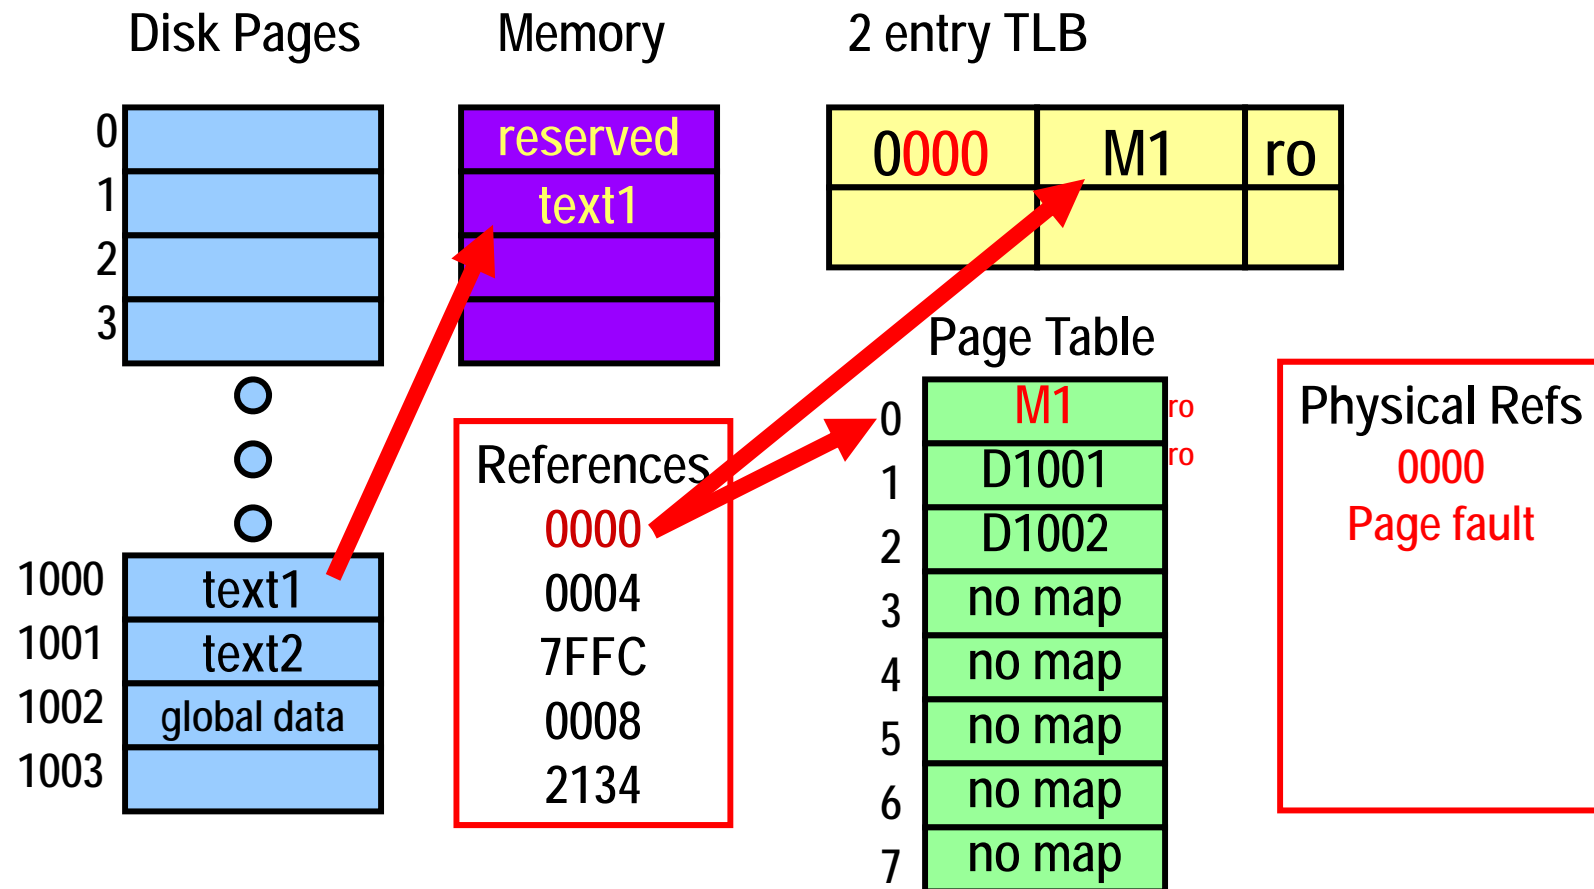
# Step 2: load PC from header and start execution



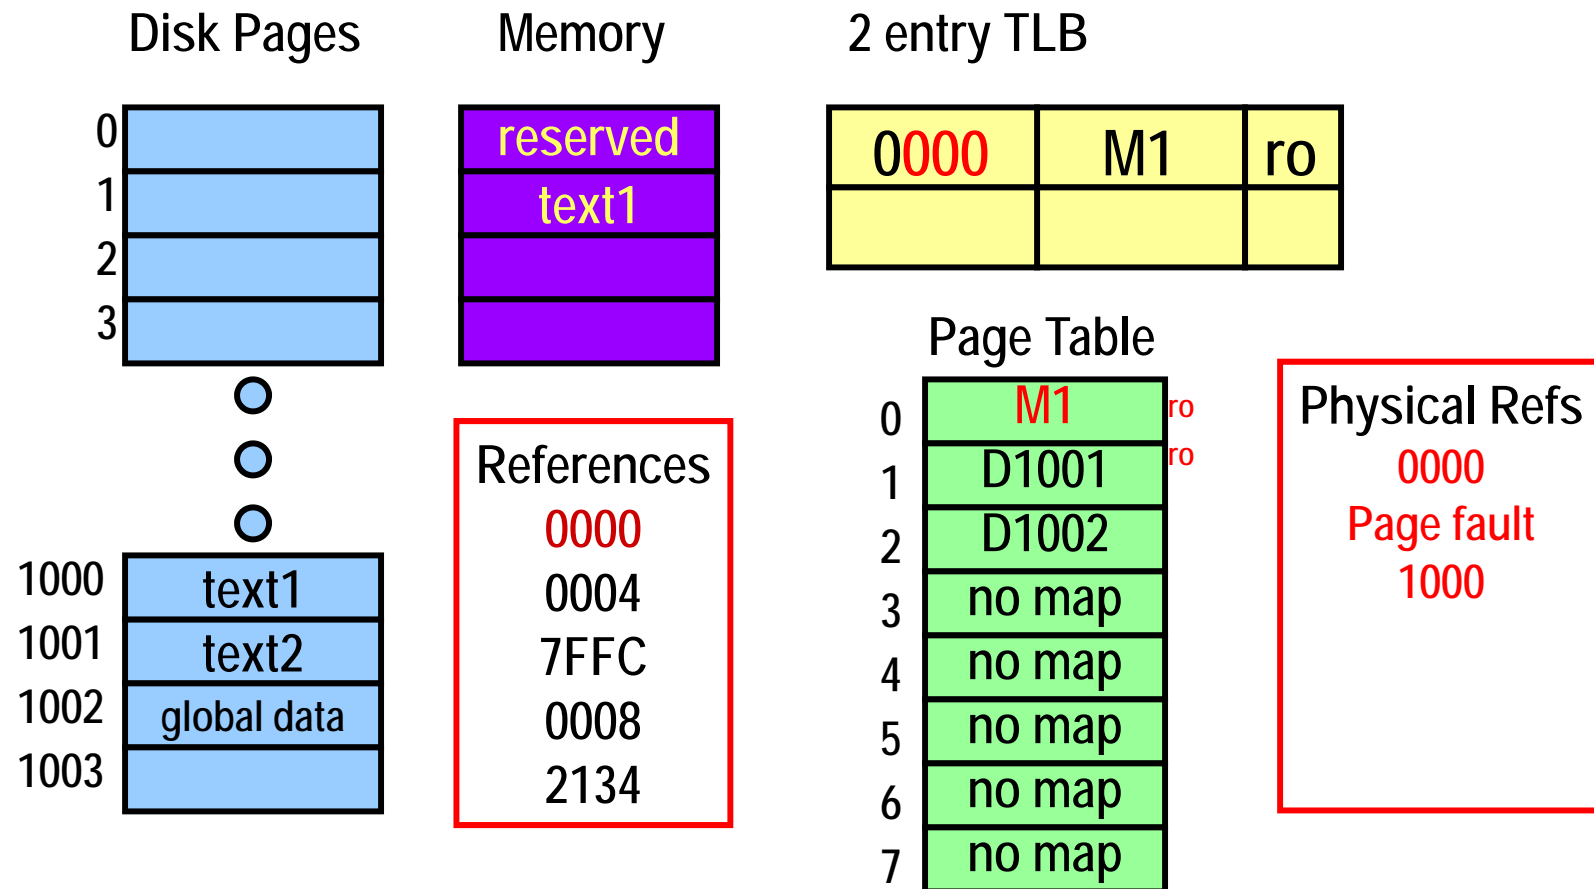
# Fetching instruction 0000



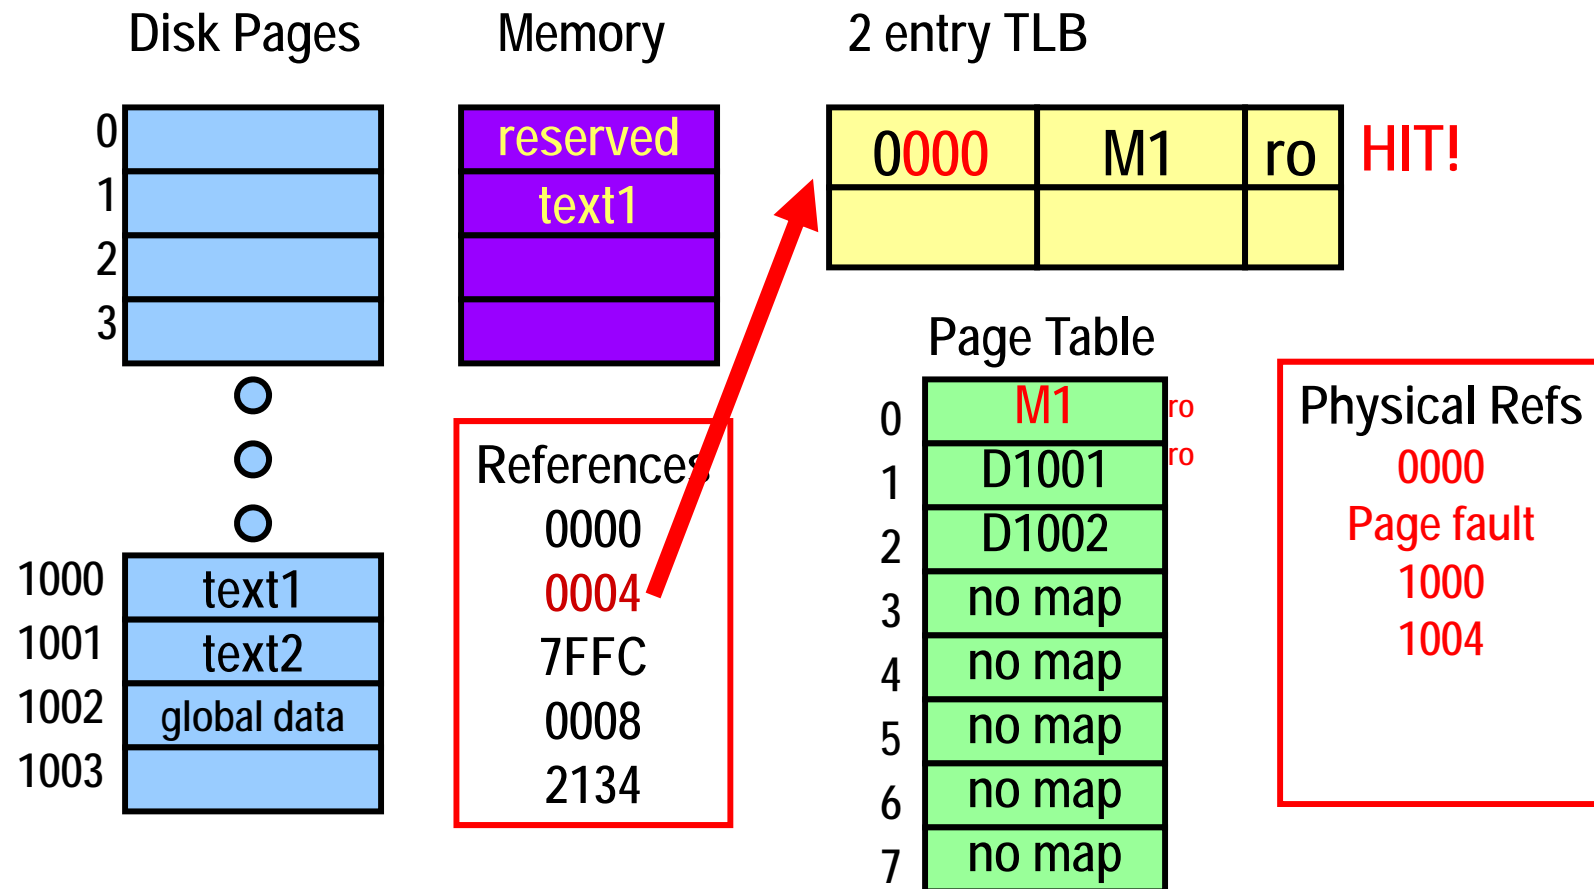
# Fetching instruction 0000



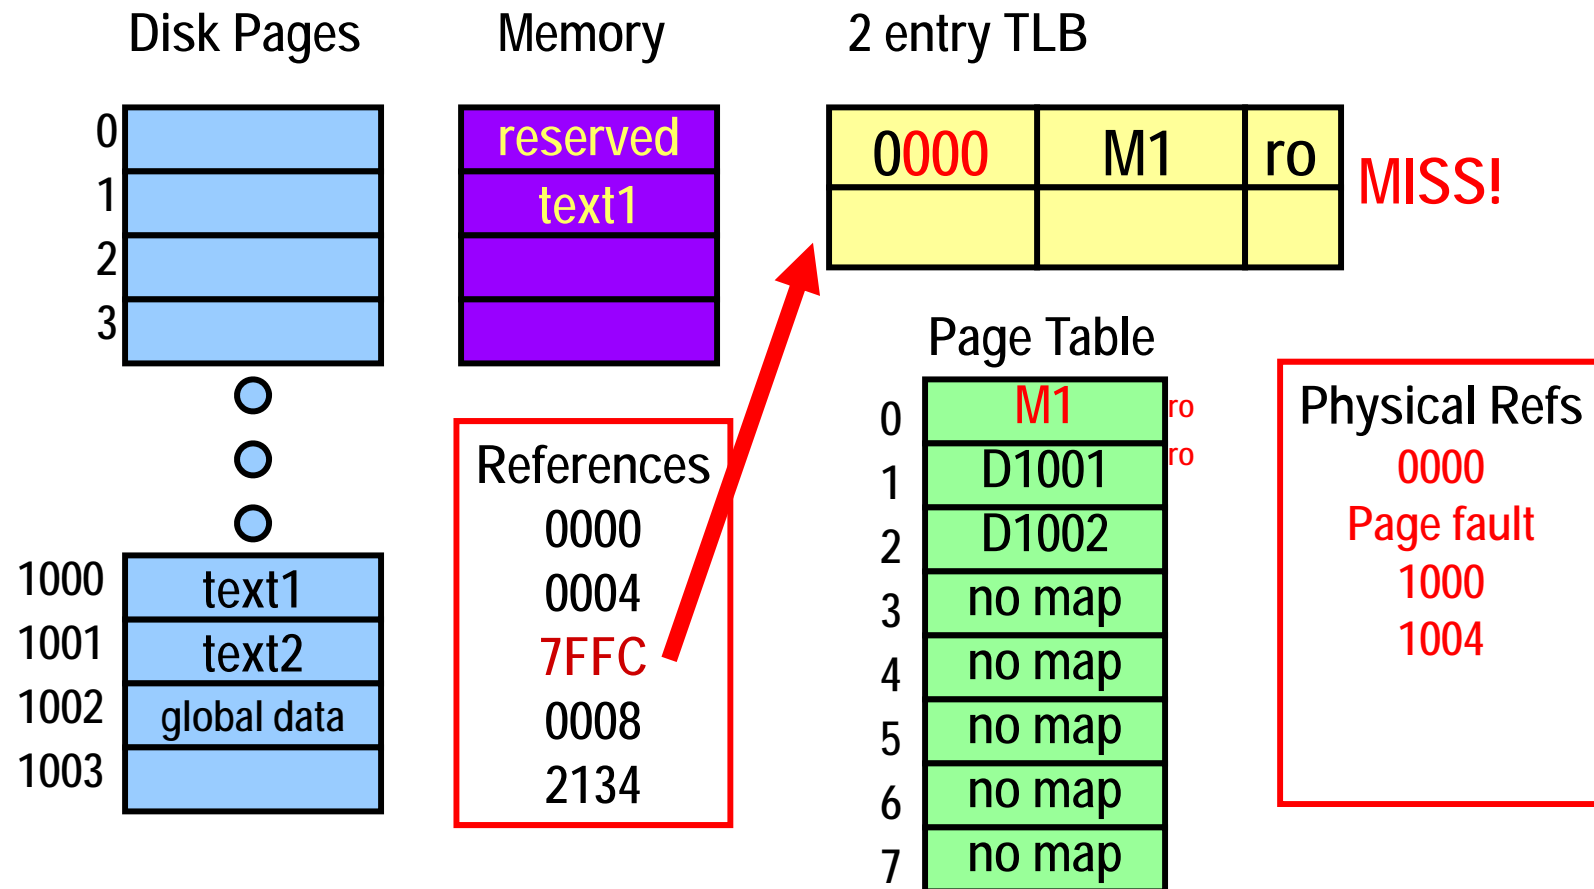
# Fetching instruction 0000



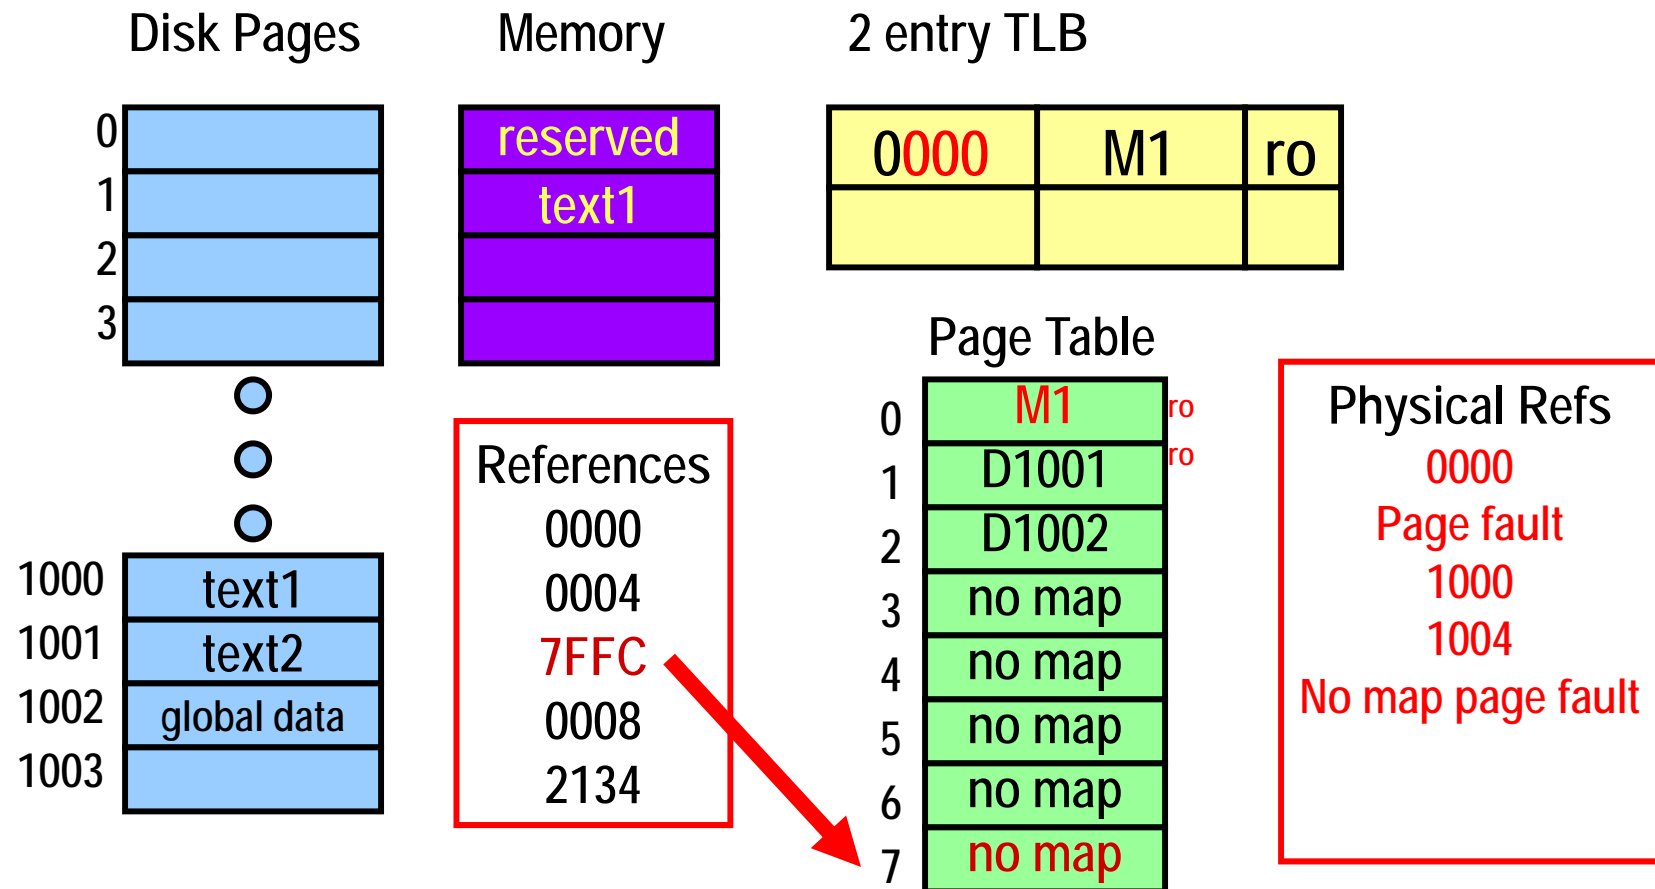
# Fetching instruction 0004



# Reference 7FFC

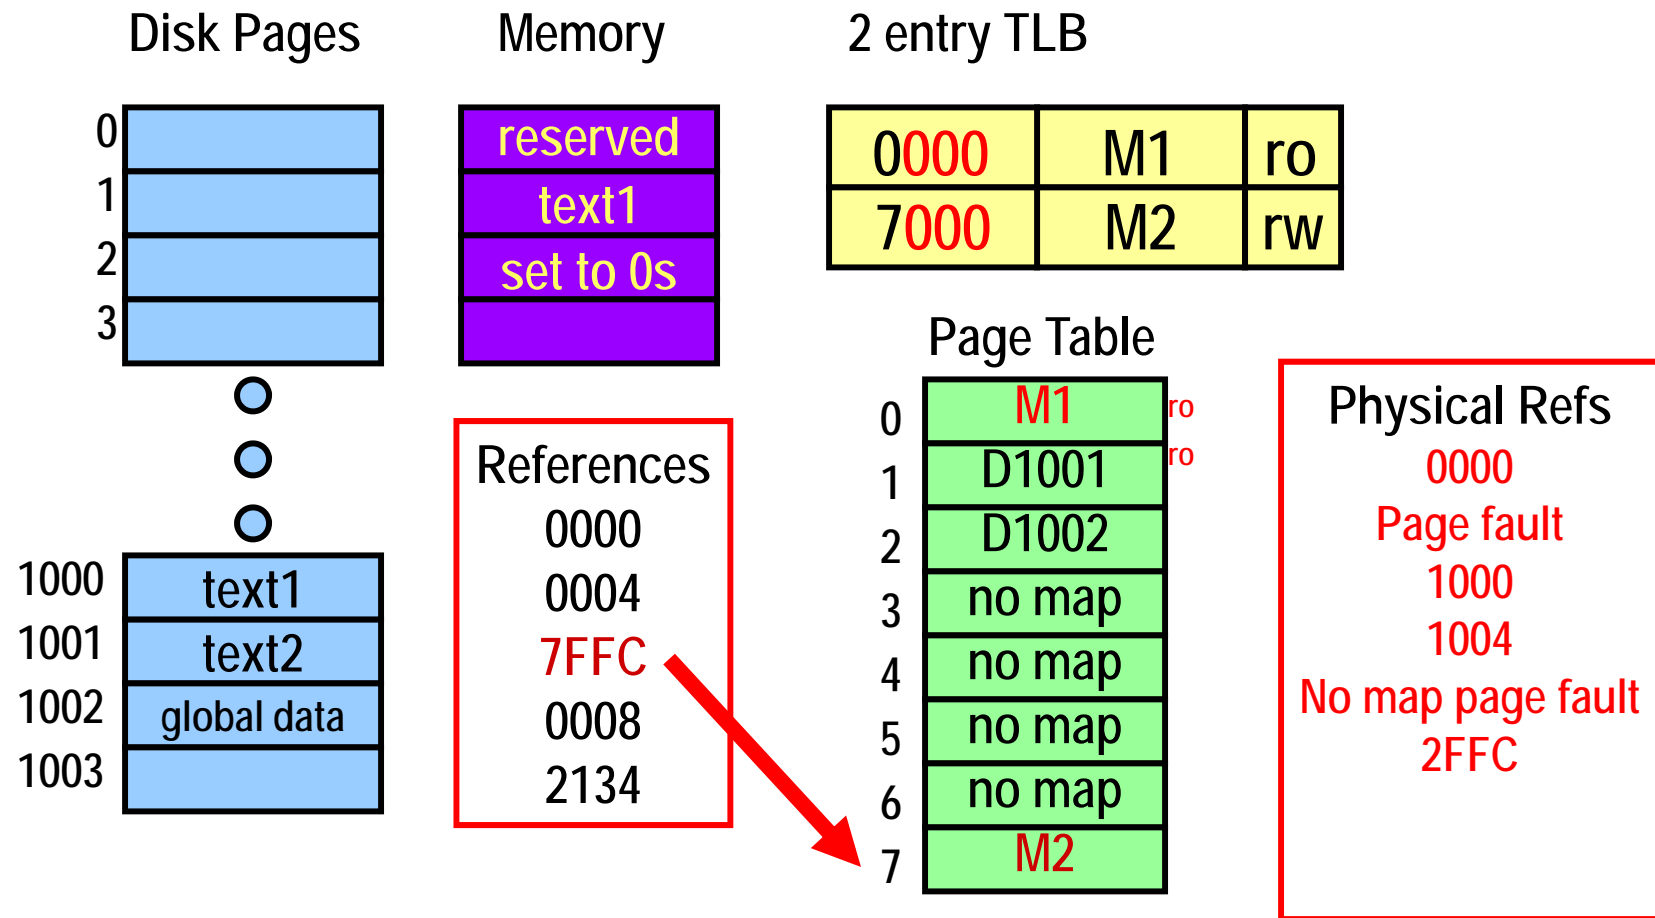


# Reference 7FFC

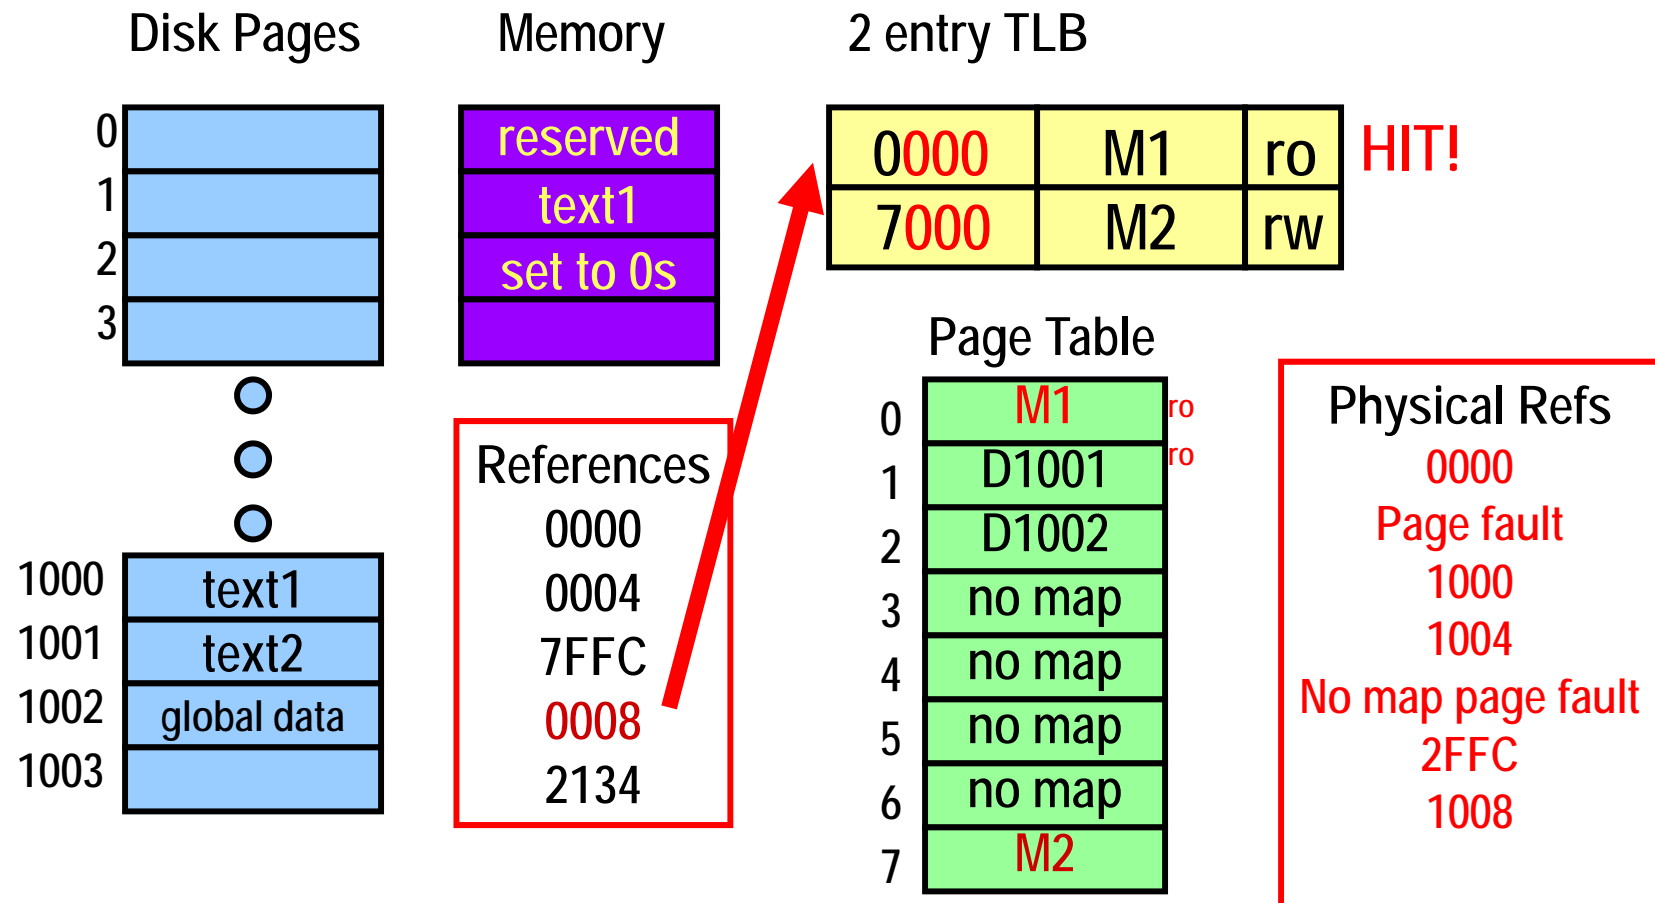




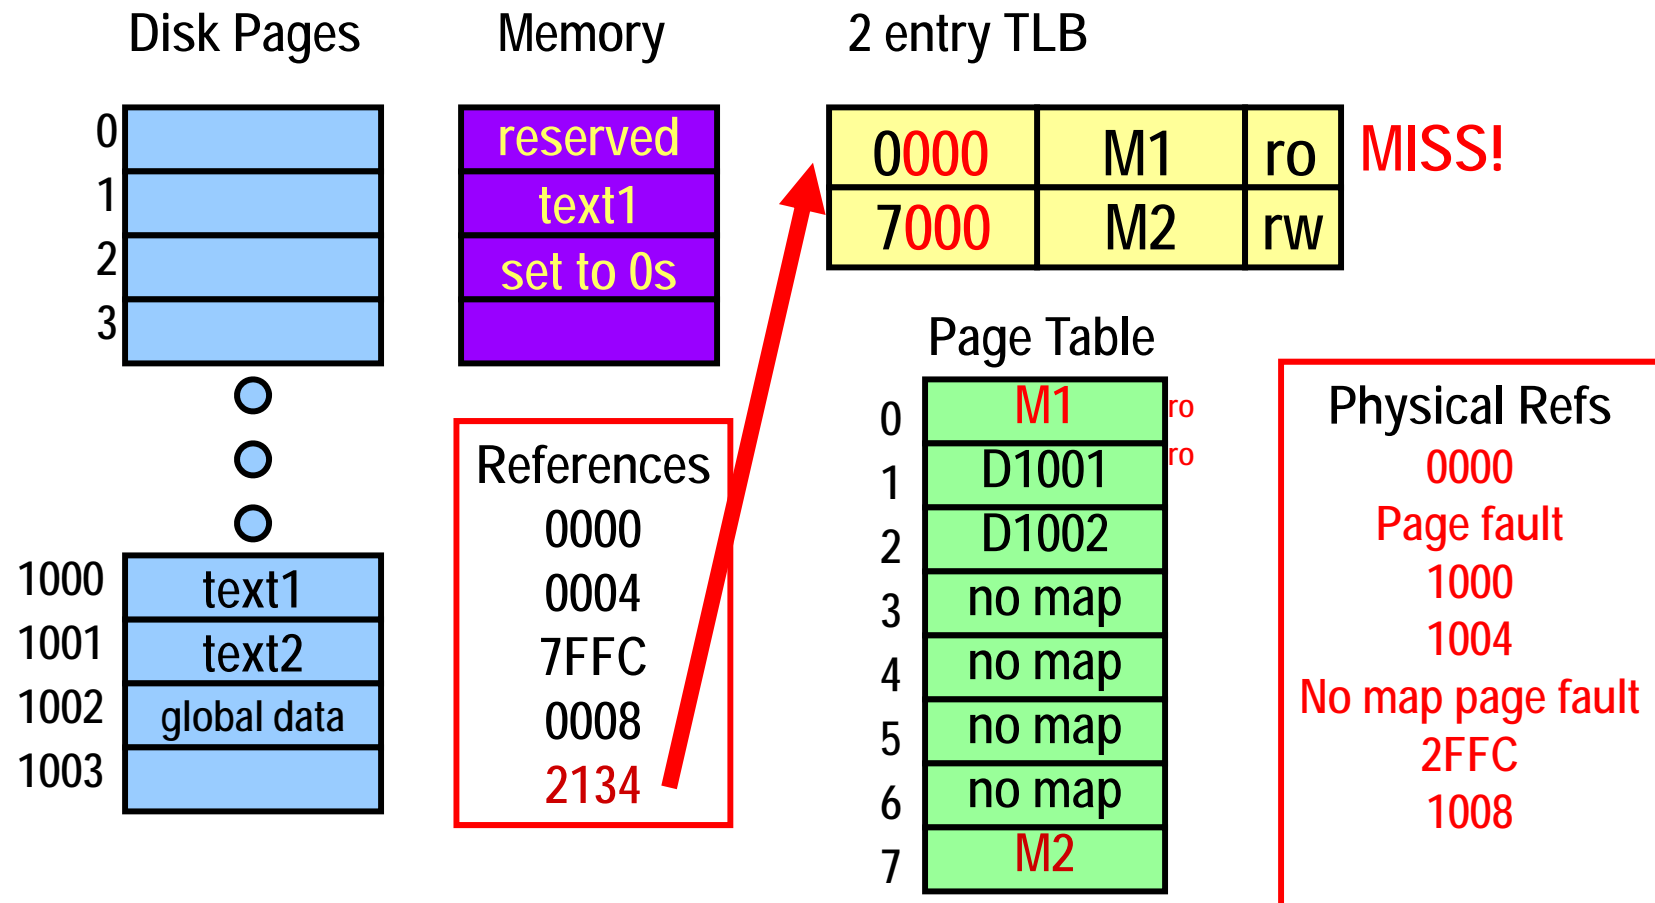
# Reference 7FFC



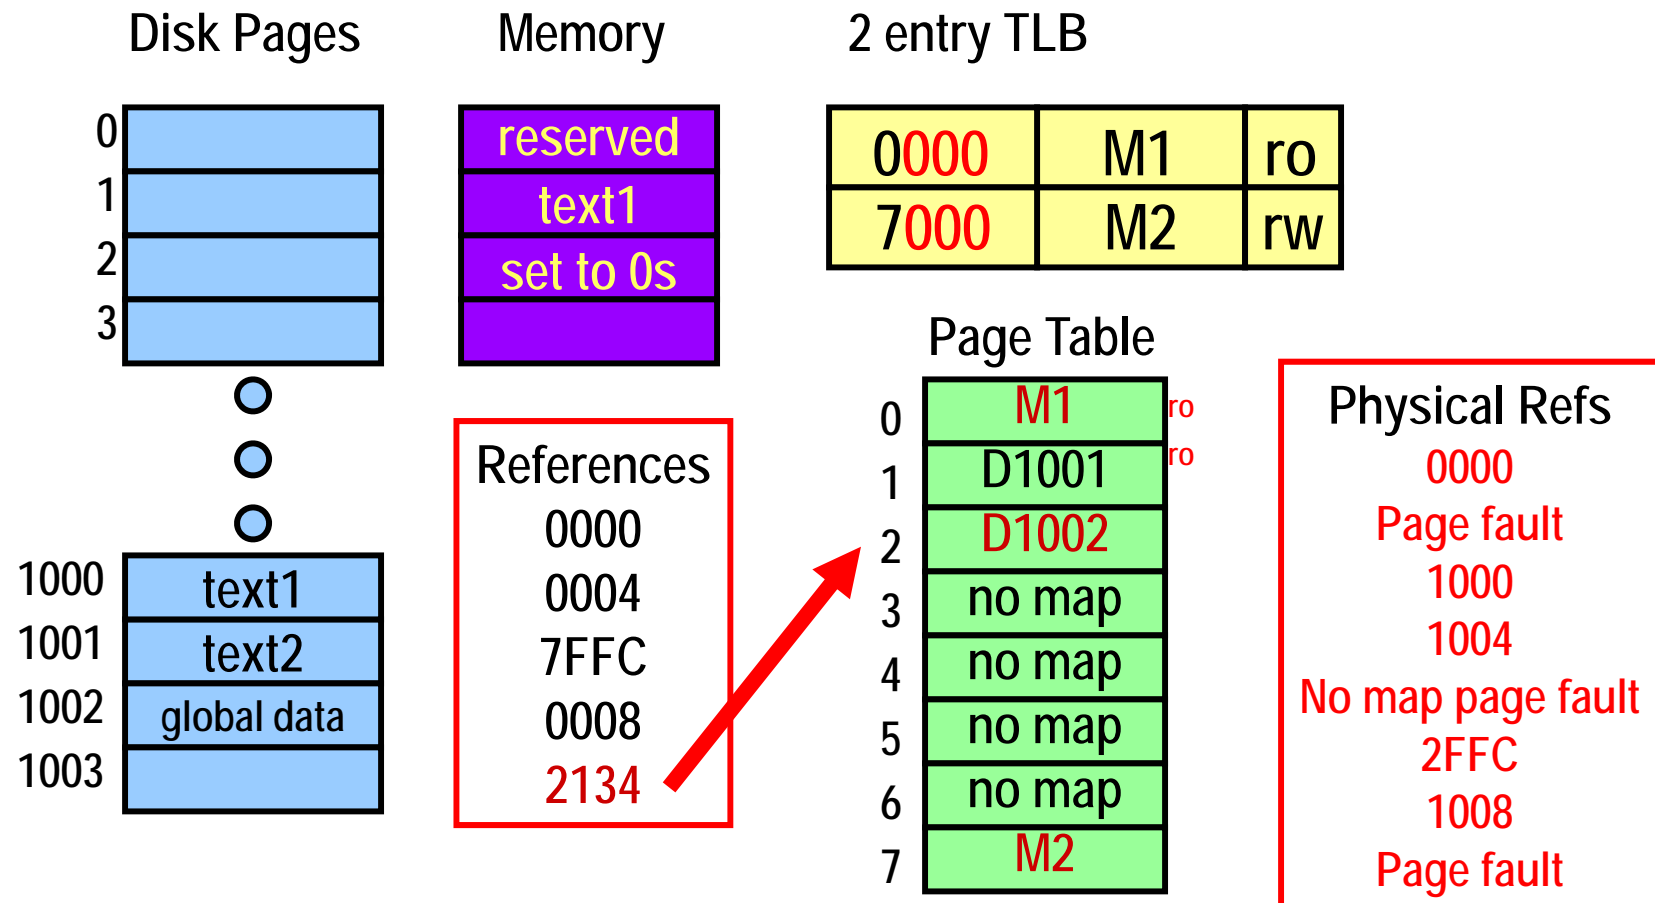
# Fetching instruction 0008



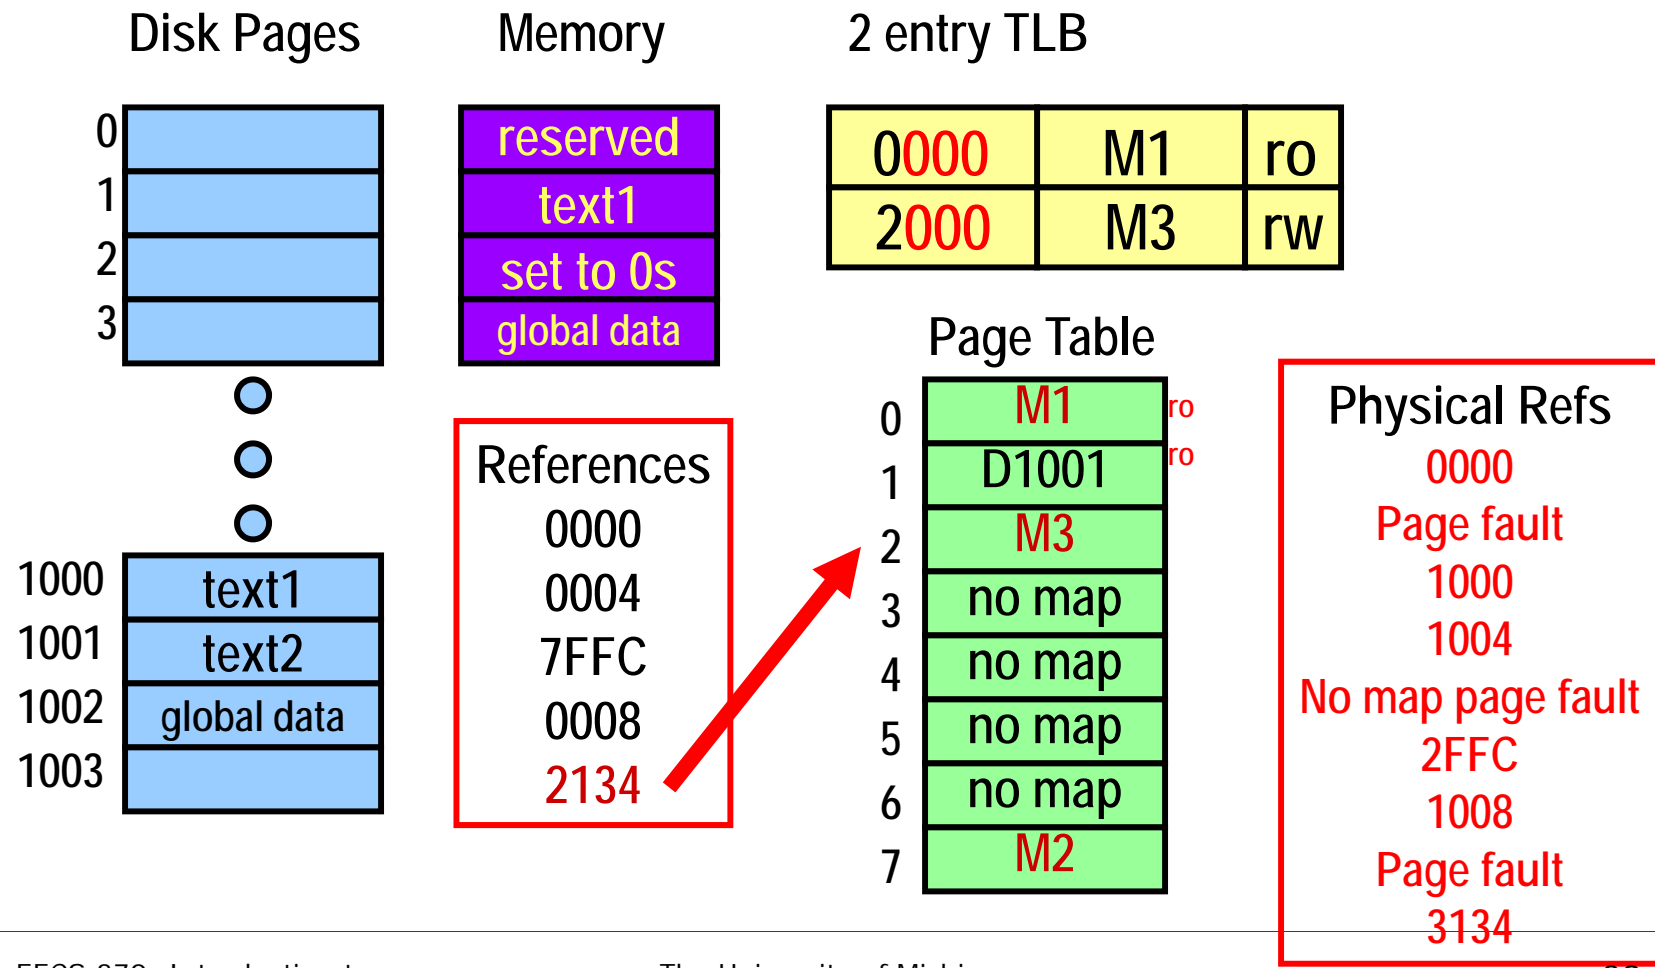
# Reference 2134



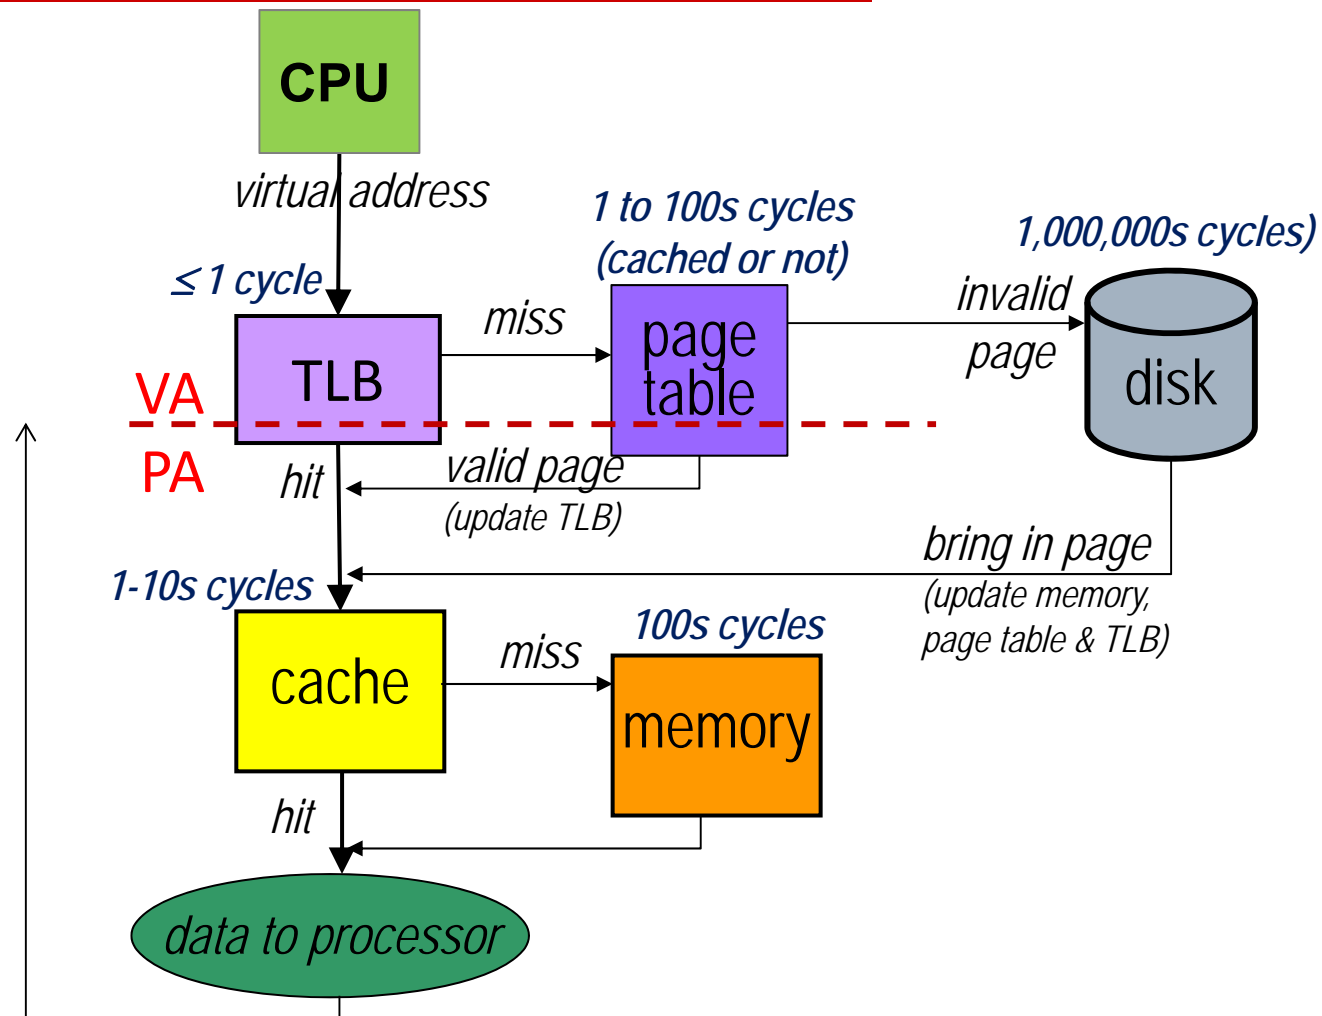
# Reference 2134



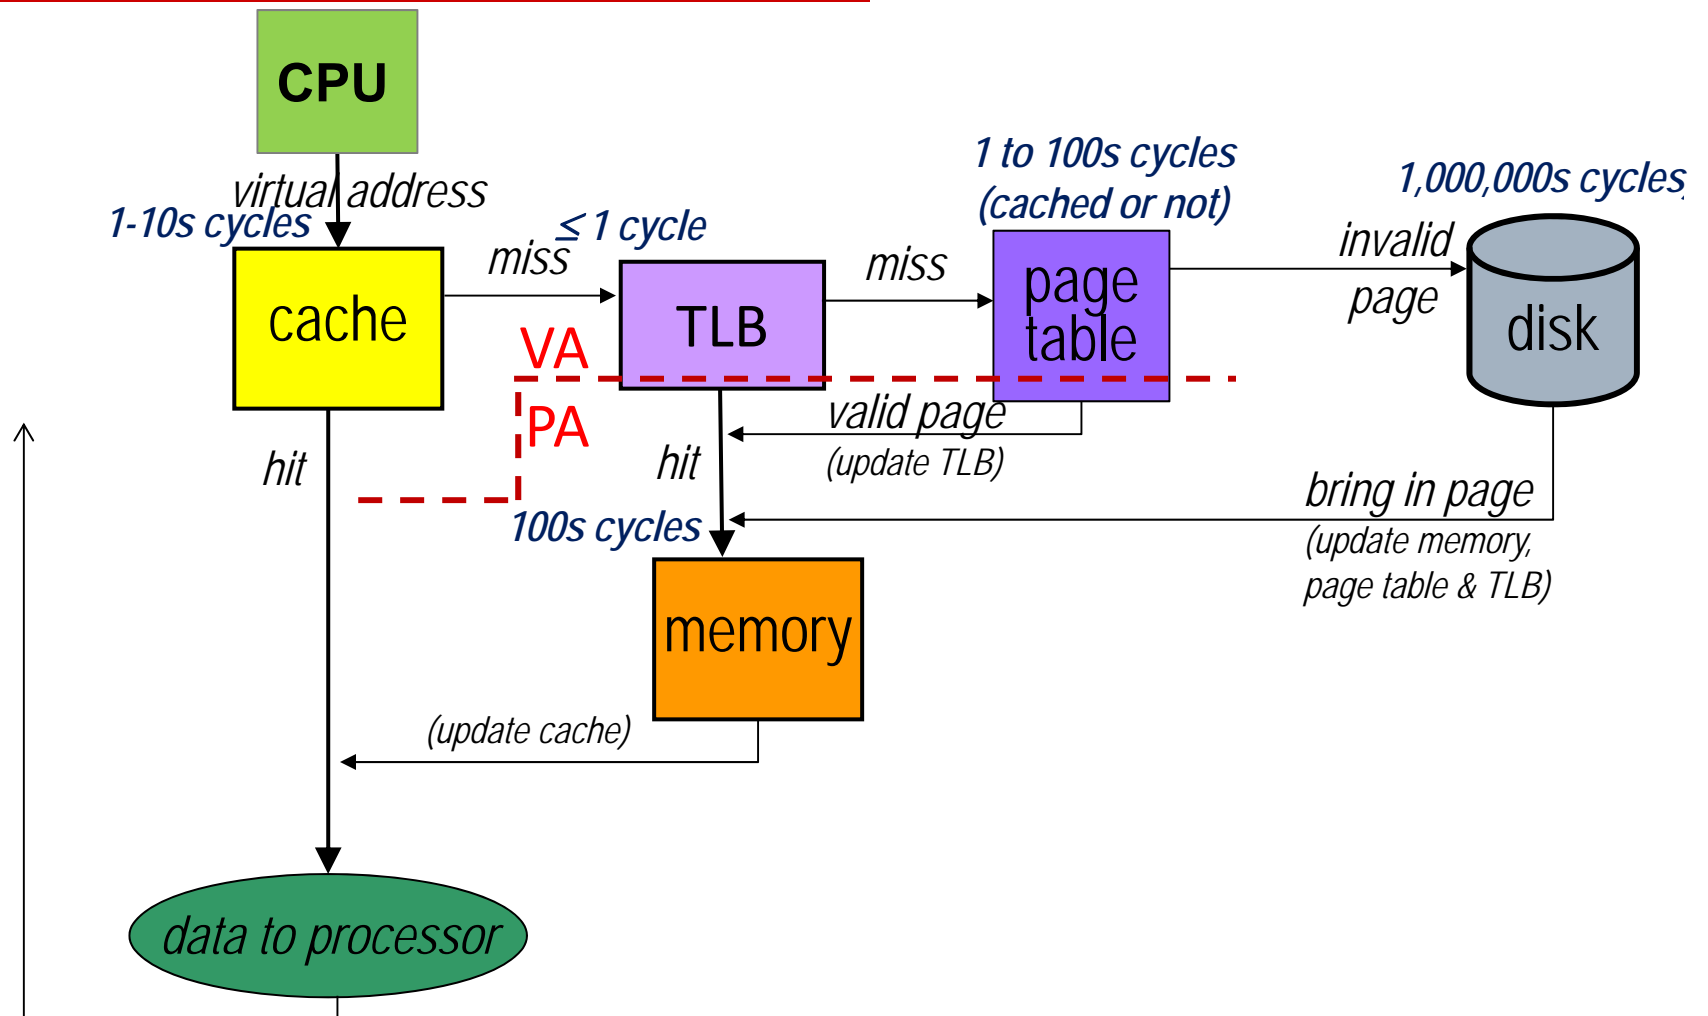
# Reference 2134



# Physically addressed caches: detailed flow



# Virtually addressed caches: detailed flow



# Class problem 2 – VM performance

---

- ❑ Consider a system with unified data and instruction cache. The virtual memory system is a one-level page table system.
  - TLB hit rate: 99 %, TLB access time is 1 cycle
  - Cache hit rate: 95 %, cache access time is 1 cycle
  - 'Page not loaded in memory rate: 0.1 % (that's 1 in one thousand memory accesses)
  - Accesses to main memory require 30 cycles and hard drive require 100,000 cycles.

## Notes:

- The TLB access and cache access are sequential.
- If data is swapped out to hard drive, the TLB will always miss.
- TLB hits will not cause a page fault.
- The page table is uncacheable and always available in main memory.
- Upon retrieval from cache, main memory or hard drive, the data is sent immediately to the CPU, while other updates occur in parallel.

Compute the average latency of a memory access for this machine.



# Problems solutions

---

# Lecture 20 – problem 2, part 2

---

Virtual Address	Virtual Super Page	Virtual Page	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C	0x01	0x07	0x10C	Y	0x002	0x0050C

Virtual superpage = 7b	Virtual page = 8b	Page offset = 9b
------------------------	-------------------	------------------

*Virtual address = 24b*

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

*Physical address = 18b*

Assume memory for pages is “somewhere else” in memory

---

# Class Problem 1 – VM architecture

---

Page Offset: 12 bits (4KB page size)

Physical page number = 17b	Page offset = 12b
----------------------------	-------------------

*Physical address = 29b (512MB Mem size)*

Virtual page number = 20b	Page offset = 12b
---------------------------	-------------------

*Virtual address = 32b*

Page table entry = 17b + 8b ~ 4B

Page table size =  $2^{20} \times 4B = 4MB$

How many pages does the page table occupy?  $4MB / 4KB = 1K$  pages

Memory left for page tables =  $512MB - 128MB$  (Data) = 384MB

Max # processes =  $384MB / 4MB$  (1 page table per process) = 96

smaller

# Class problem 2 – VM performance

---

Compute the average latency of a memory access for this machine.

Physically tagged:

Address translation:  $1 + 0.01 \cdot 30$

Data access:  $1 + 0.05 \cdot (30 + 0.001 \cdot 100k)$

Virtually tagged:

Address translation:  $0.05 \cdot (1 + 0.01 \cdot 30)$

Data access:  $1 + 0.05 \cdot (30 + 0.001 \cdot 100k)$

**Avg Mem Access = Addr translation + Data access**